



Ground up Introduction to In-memory Data (Grids)

Tweet about this
`#hazelcast #imdg #devnexus`



Why you here?

A photograph of five business professionals in a modern office setting. In the foreground, a man with dark hair and a light beard, wearing a grey suit, stands with his arms crossed. Behind him, four other individuals are visible: two women on the left and two men on the right, all dressed in professional attire and smiling. The background shows large windows overlooking a city skyline.

**Java Developer on a quest for scalability
frameworks**

Architect on low-latency projects

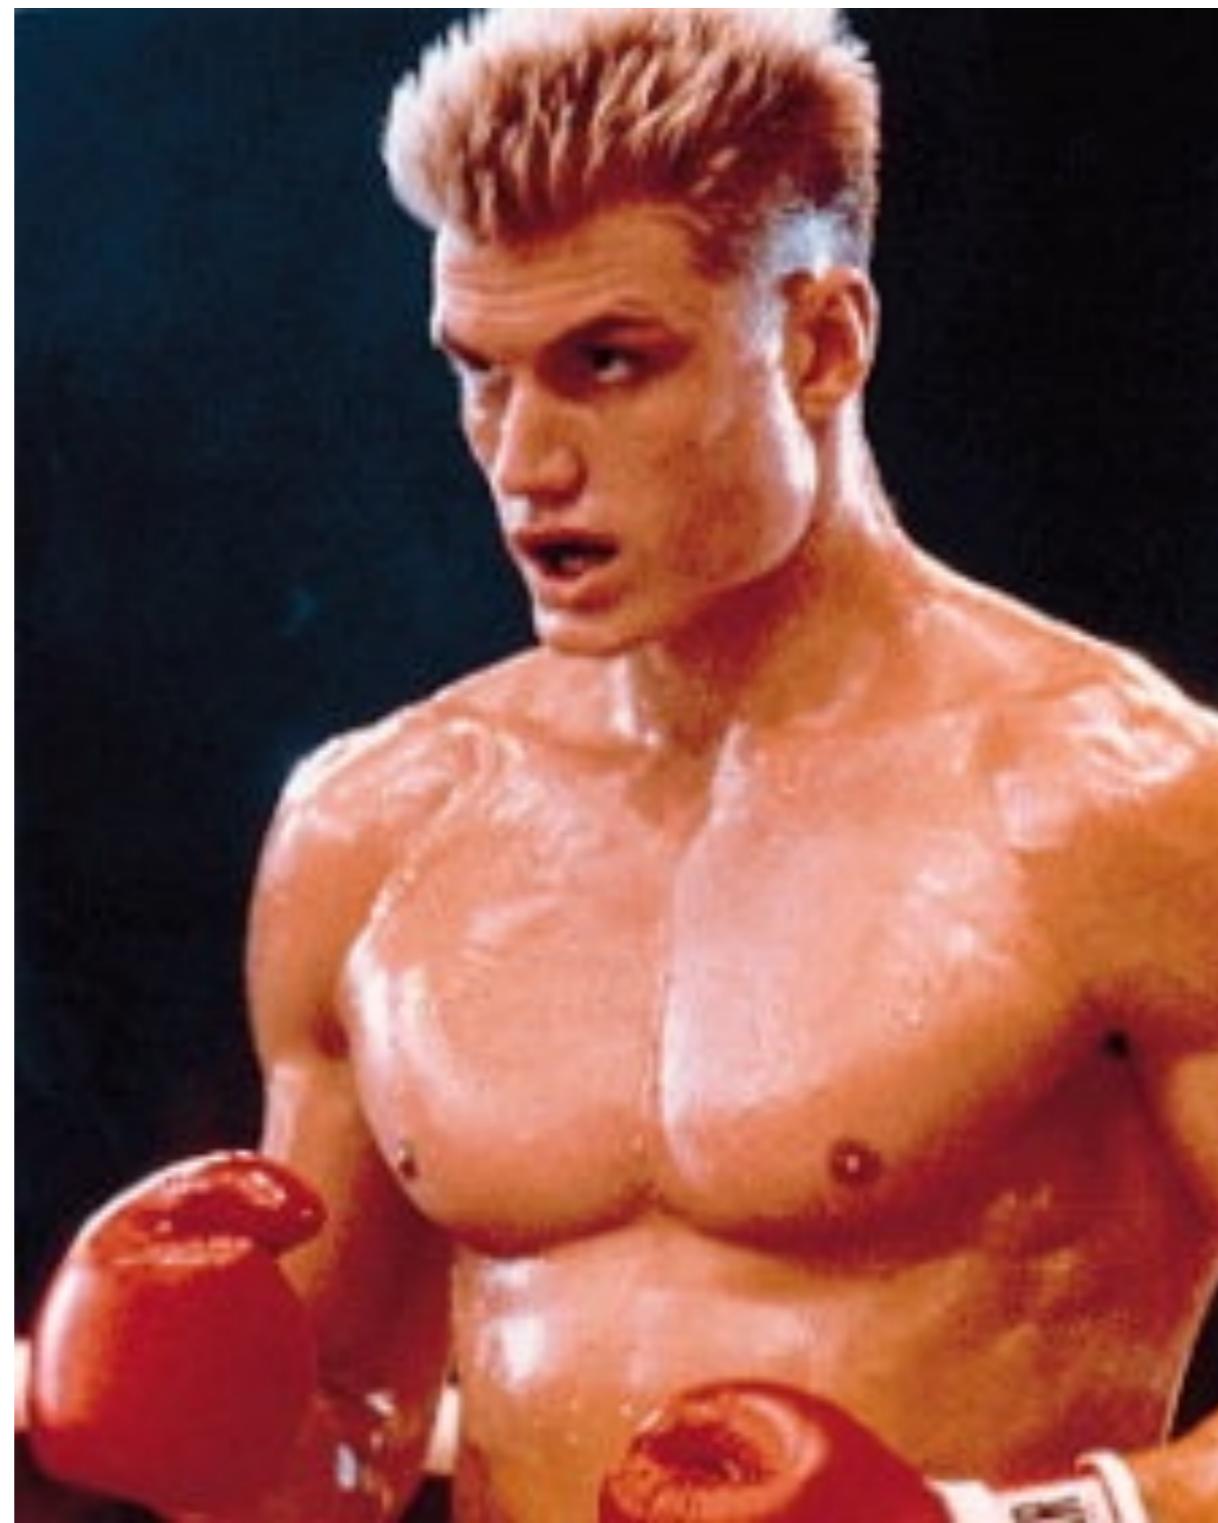
NoSQL practitioner

IMDG newbie

Data enthusiast



Who am I ?





Viktor Gamov
@gamussa



Building HTML5 Applications: From Desktop to Mobile



*Yakov Fain, Victor Rasputnis,
Anatole Tartakovsky & Viktor Gamov*

- Co-author of O'Reilly's
«Enterprise Web Development»



**hitect at Hazelcast
nal Conferences
peaker**





Agenda

- The In-Memory landscape you need to know
- Distributed Data Applications
- Making In-memory Reliable, Scalable and Durable
- QA





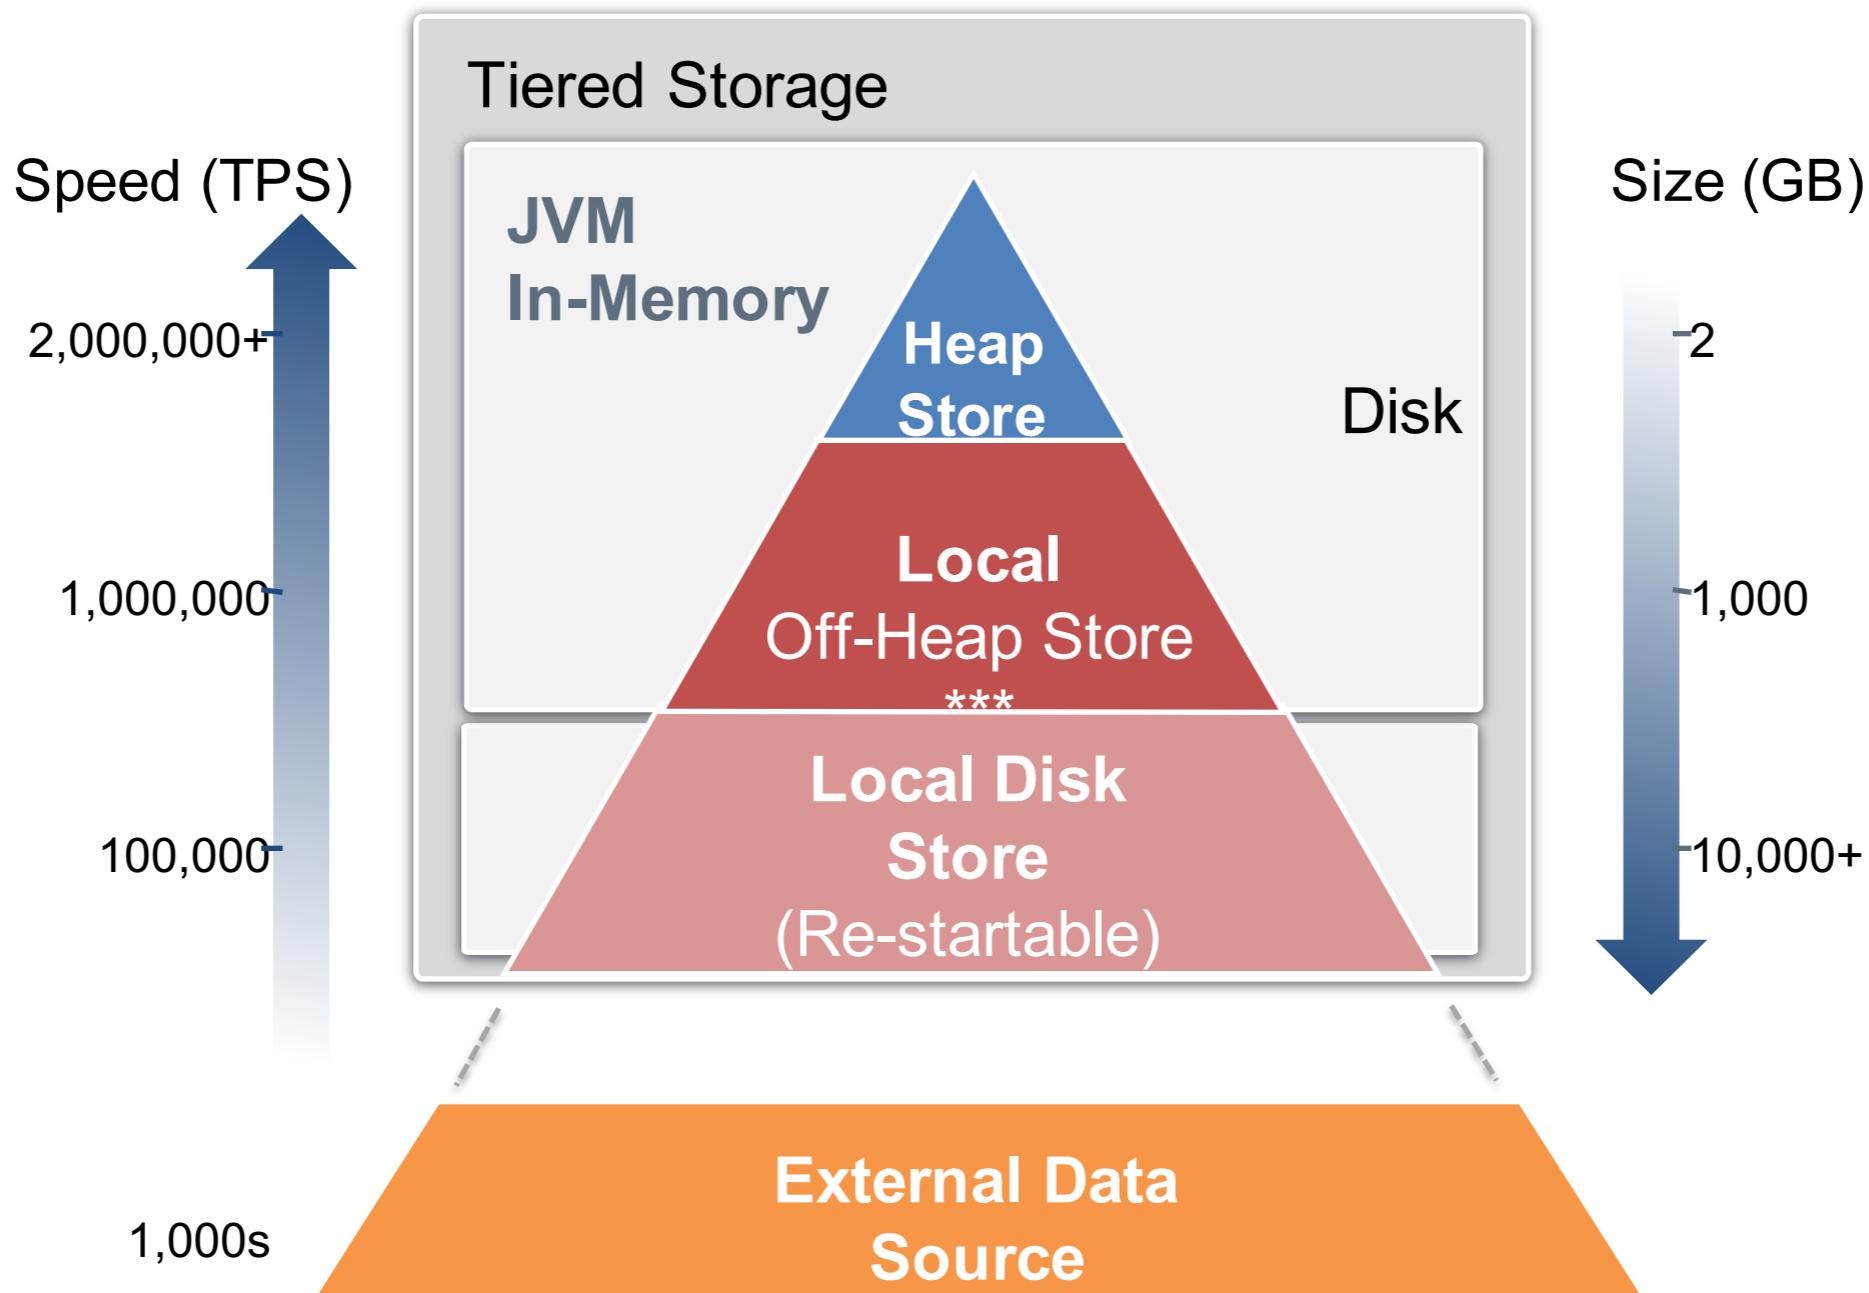
Why In-Memory Data?



Problem

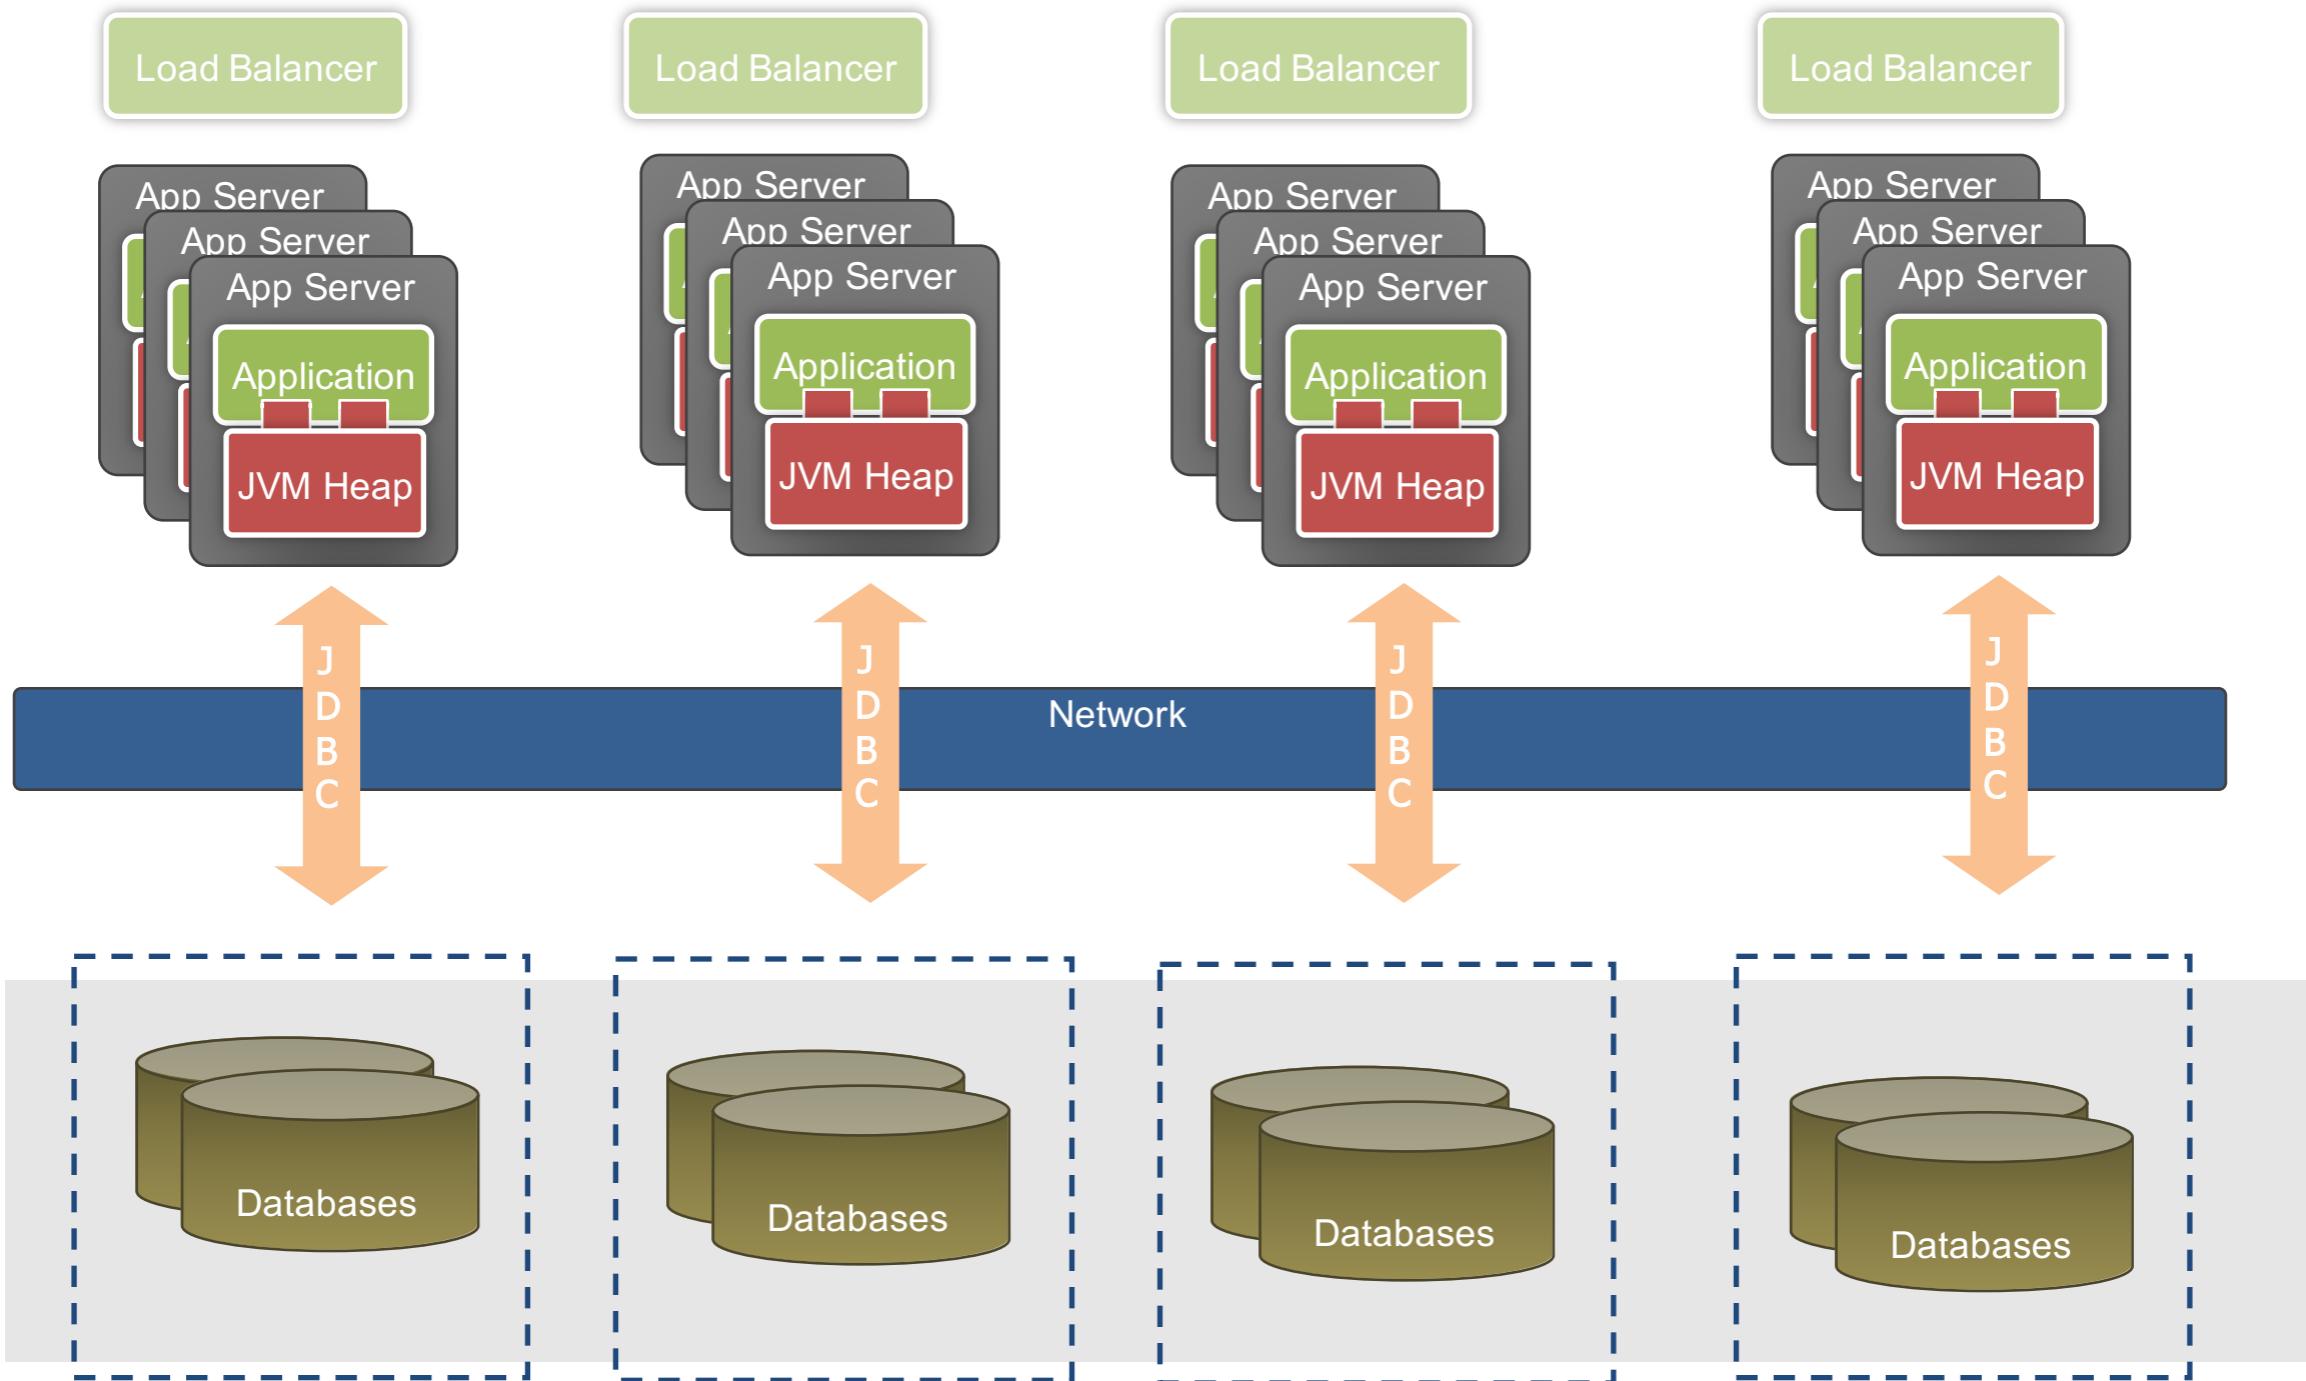


Anatomy of an Application: Tiered Data Storage





Common Enterprise Application Landscape





Performance & Scalability Challenges

- **JVM memory constraints**
 - Heap size limitations
 - GC pauses and tuning
 - Handling peak loads
 - Swapping
- **CPU limitations**
 - Limited processor power
 - Many applications sharing CPU cycles
 - CPU intensive processes

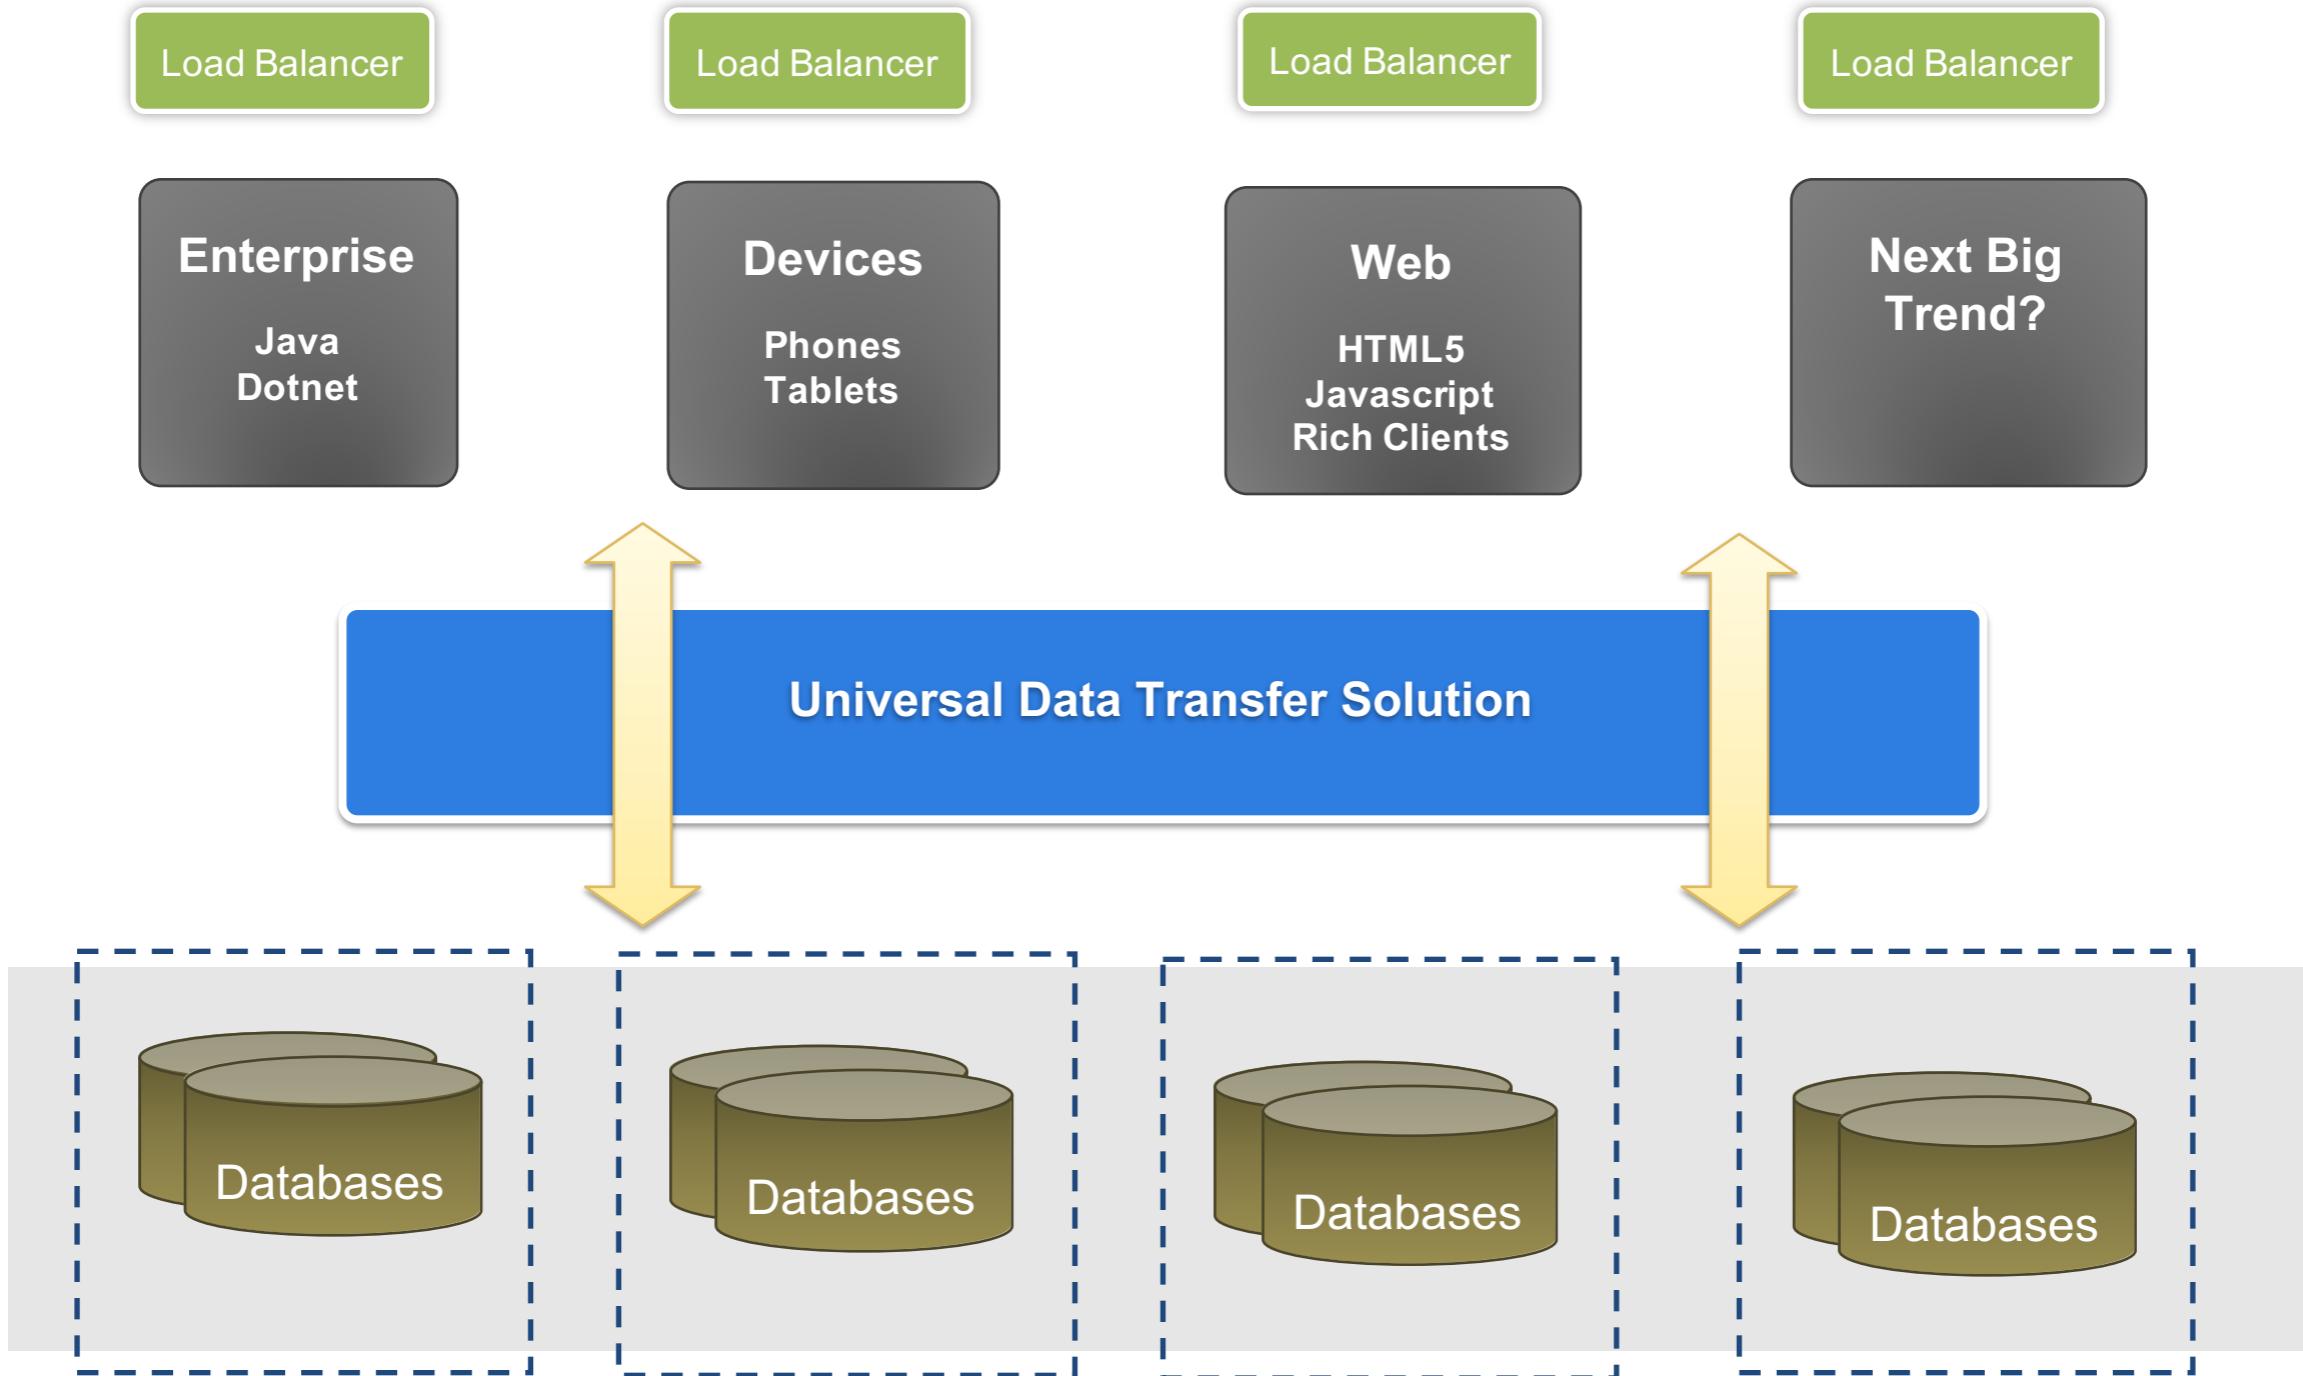


Performance & Scalability Challenges

- **Reliance on the database**
 - All transactions go through the database
 - Expensive and complex to scale
 - Reliance on DBA(s) to maintain and tune
- **Network I/O**
 - Bandwidth limitations
 - Latency
- **State-aware cluster**
 - Session replication

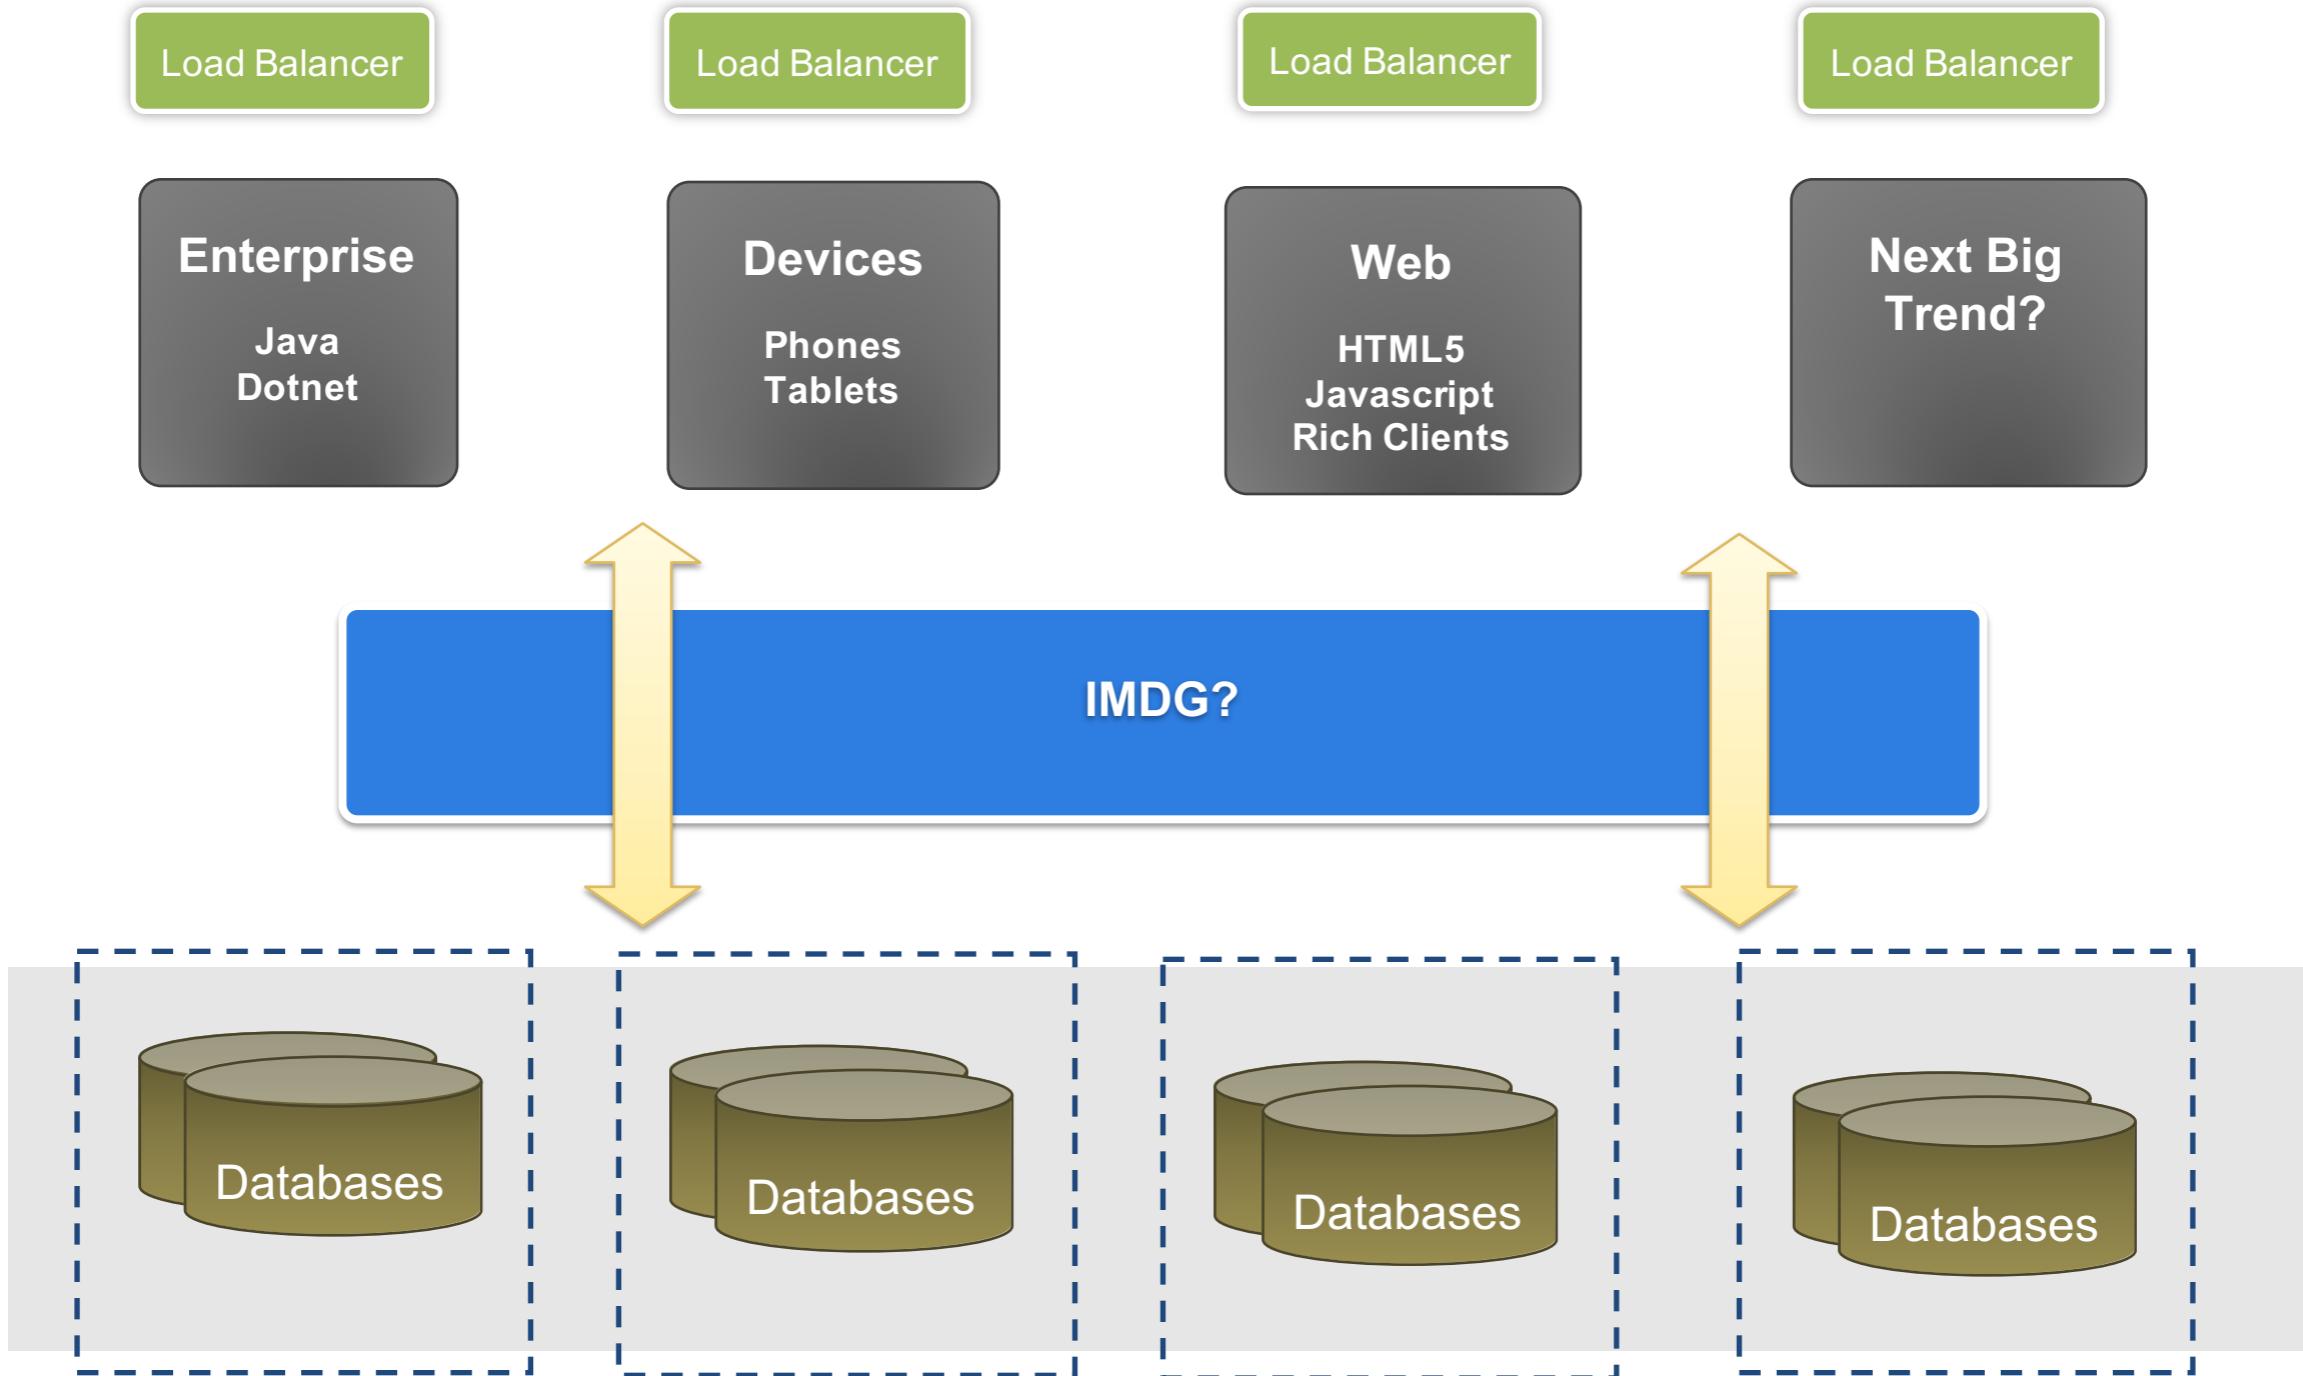


Even More Common Enterprise Application Landscape





Even More Common Enterprise Application Landscape





In-Memory Data Lanscape

Messaging Middleware

IMDG

NoSQL DATABASES

SQL DATABASES

In-Memory Computing



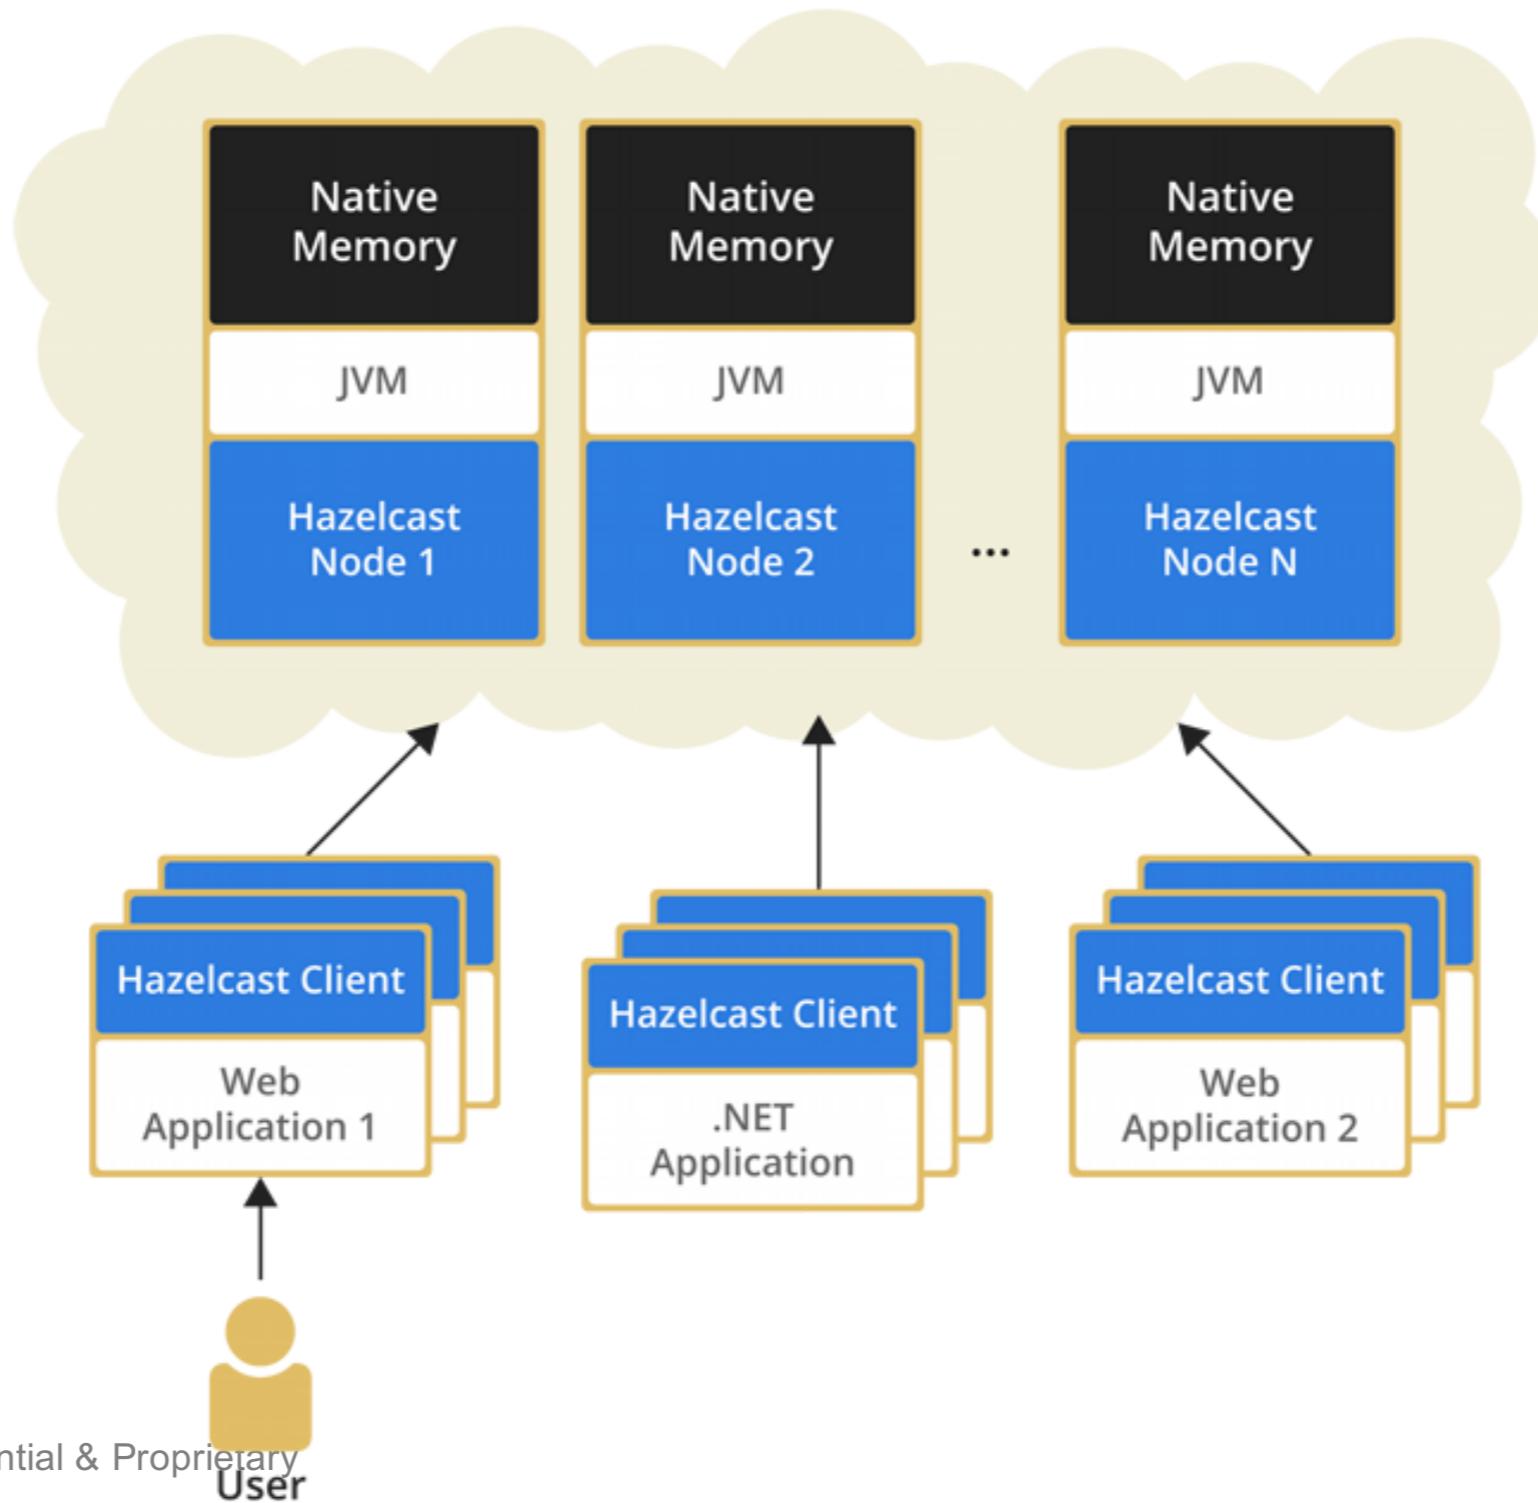
Distributed Data Applications



Distributed Data Applications

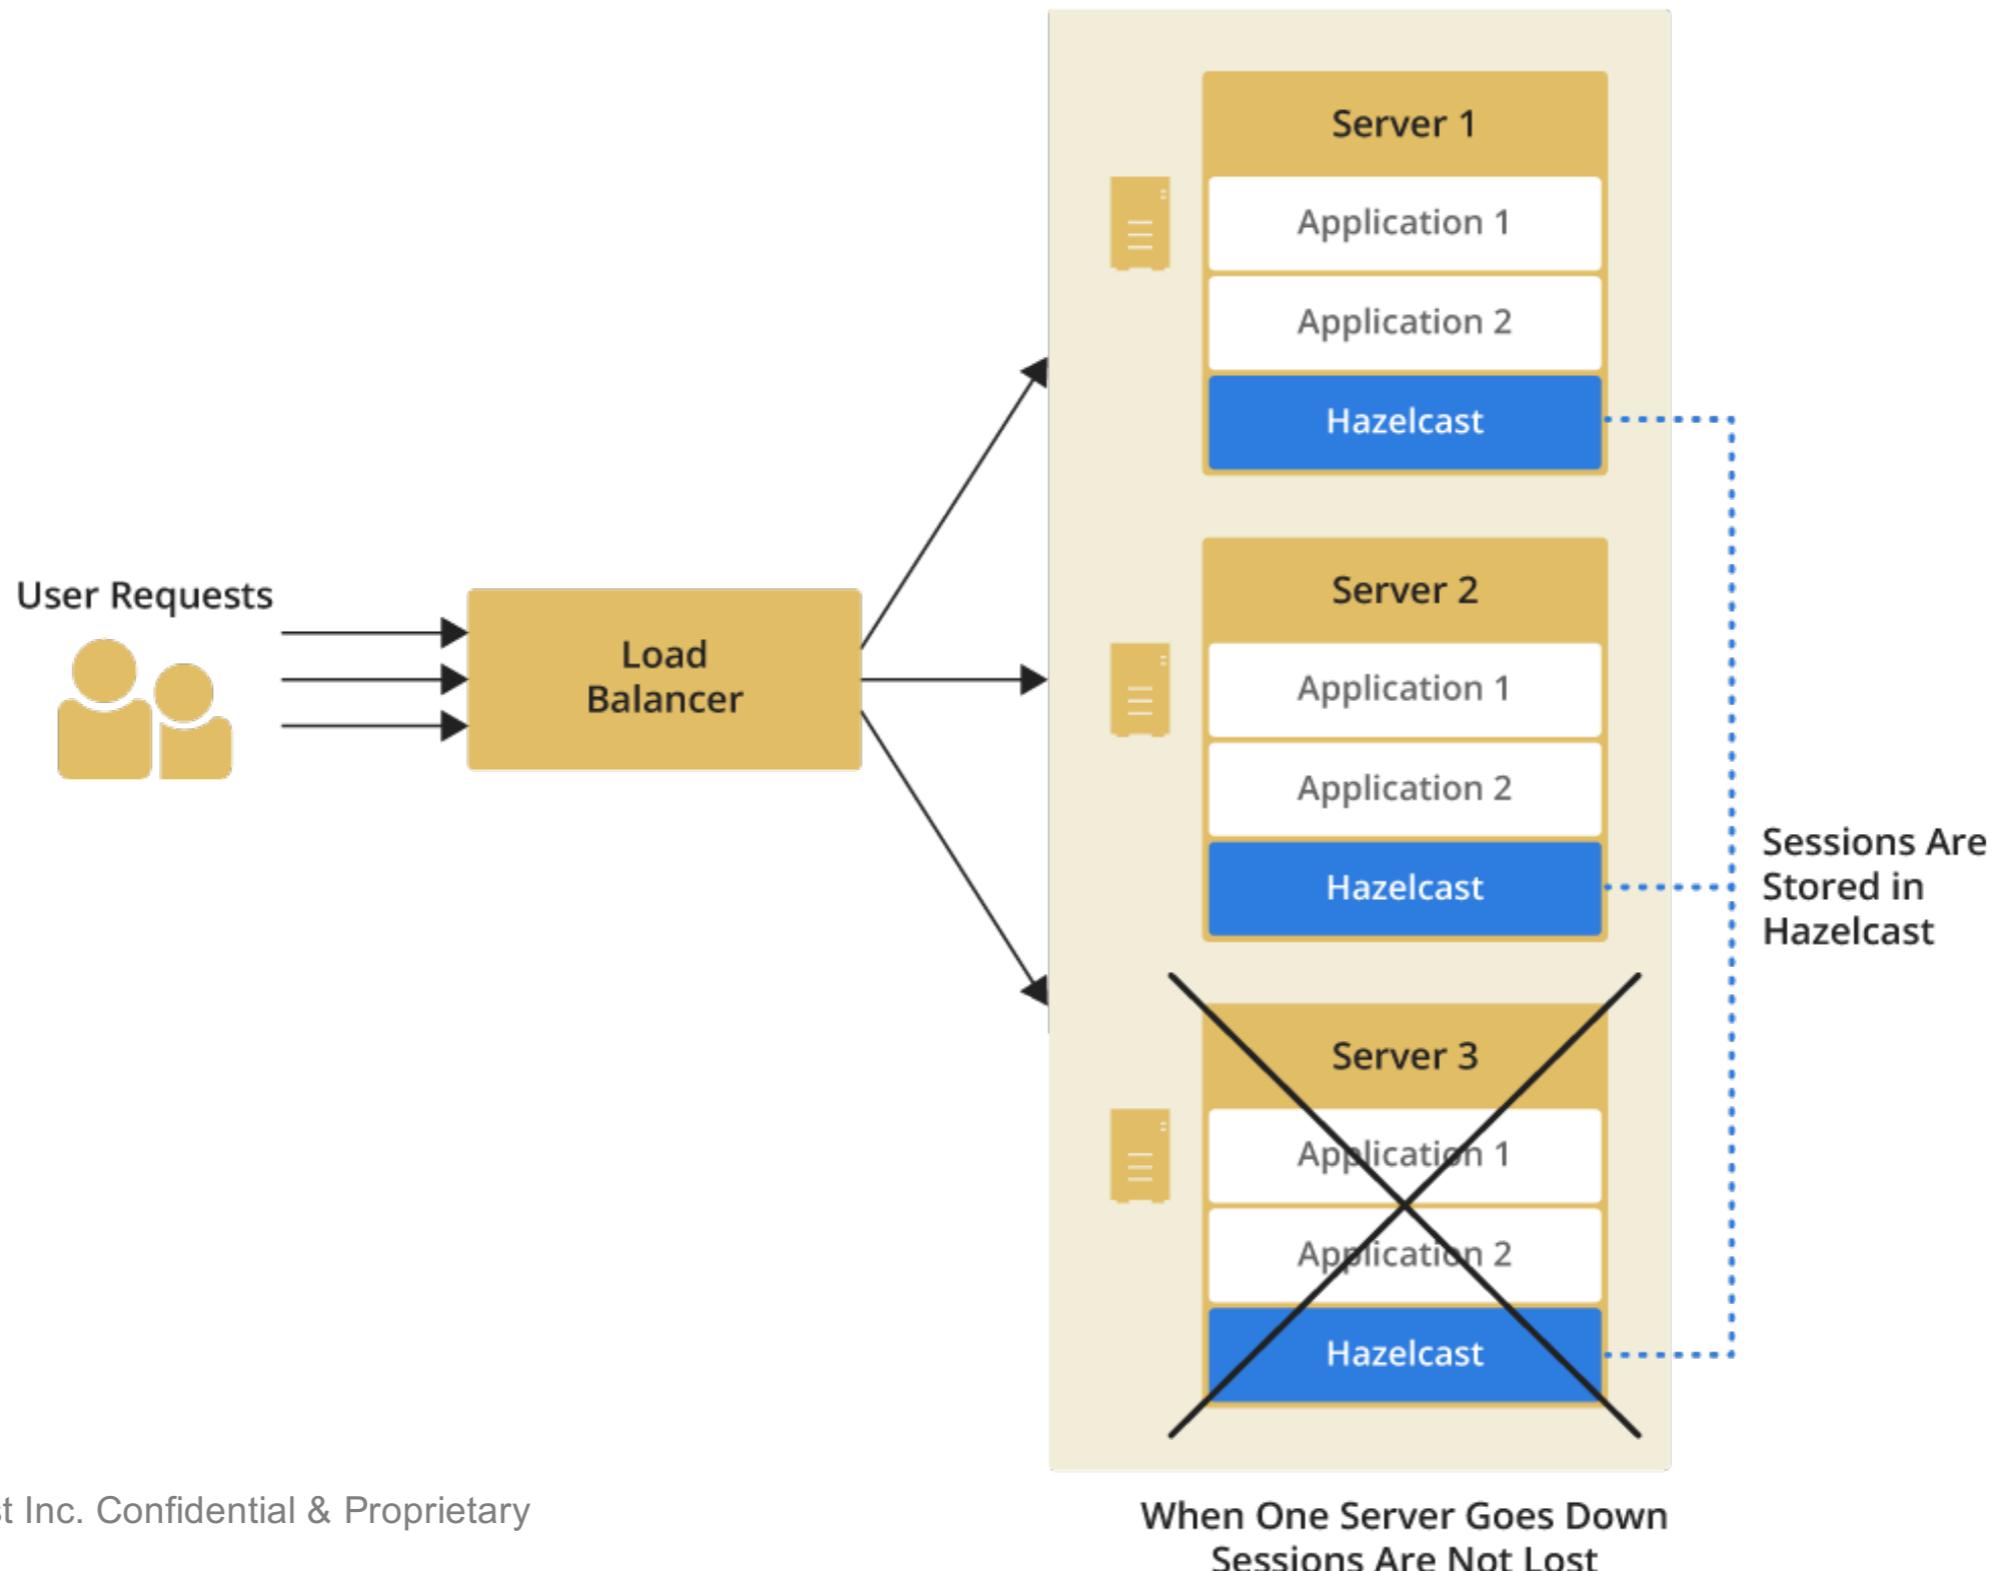
- Cache-as-a-Service
- Application Scaling
- Database Caching
- Distributed Computing
- Reactive / Smart Clients

Cache-as-a-Service





Durability of application data



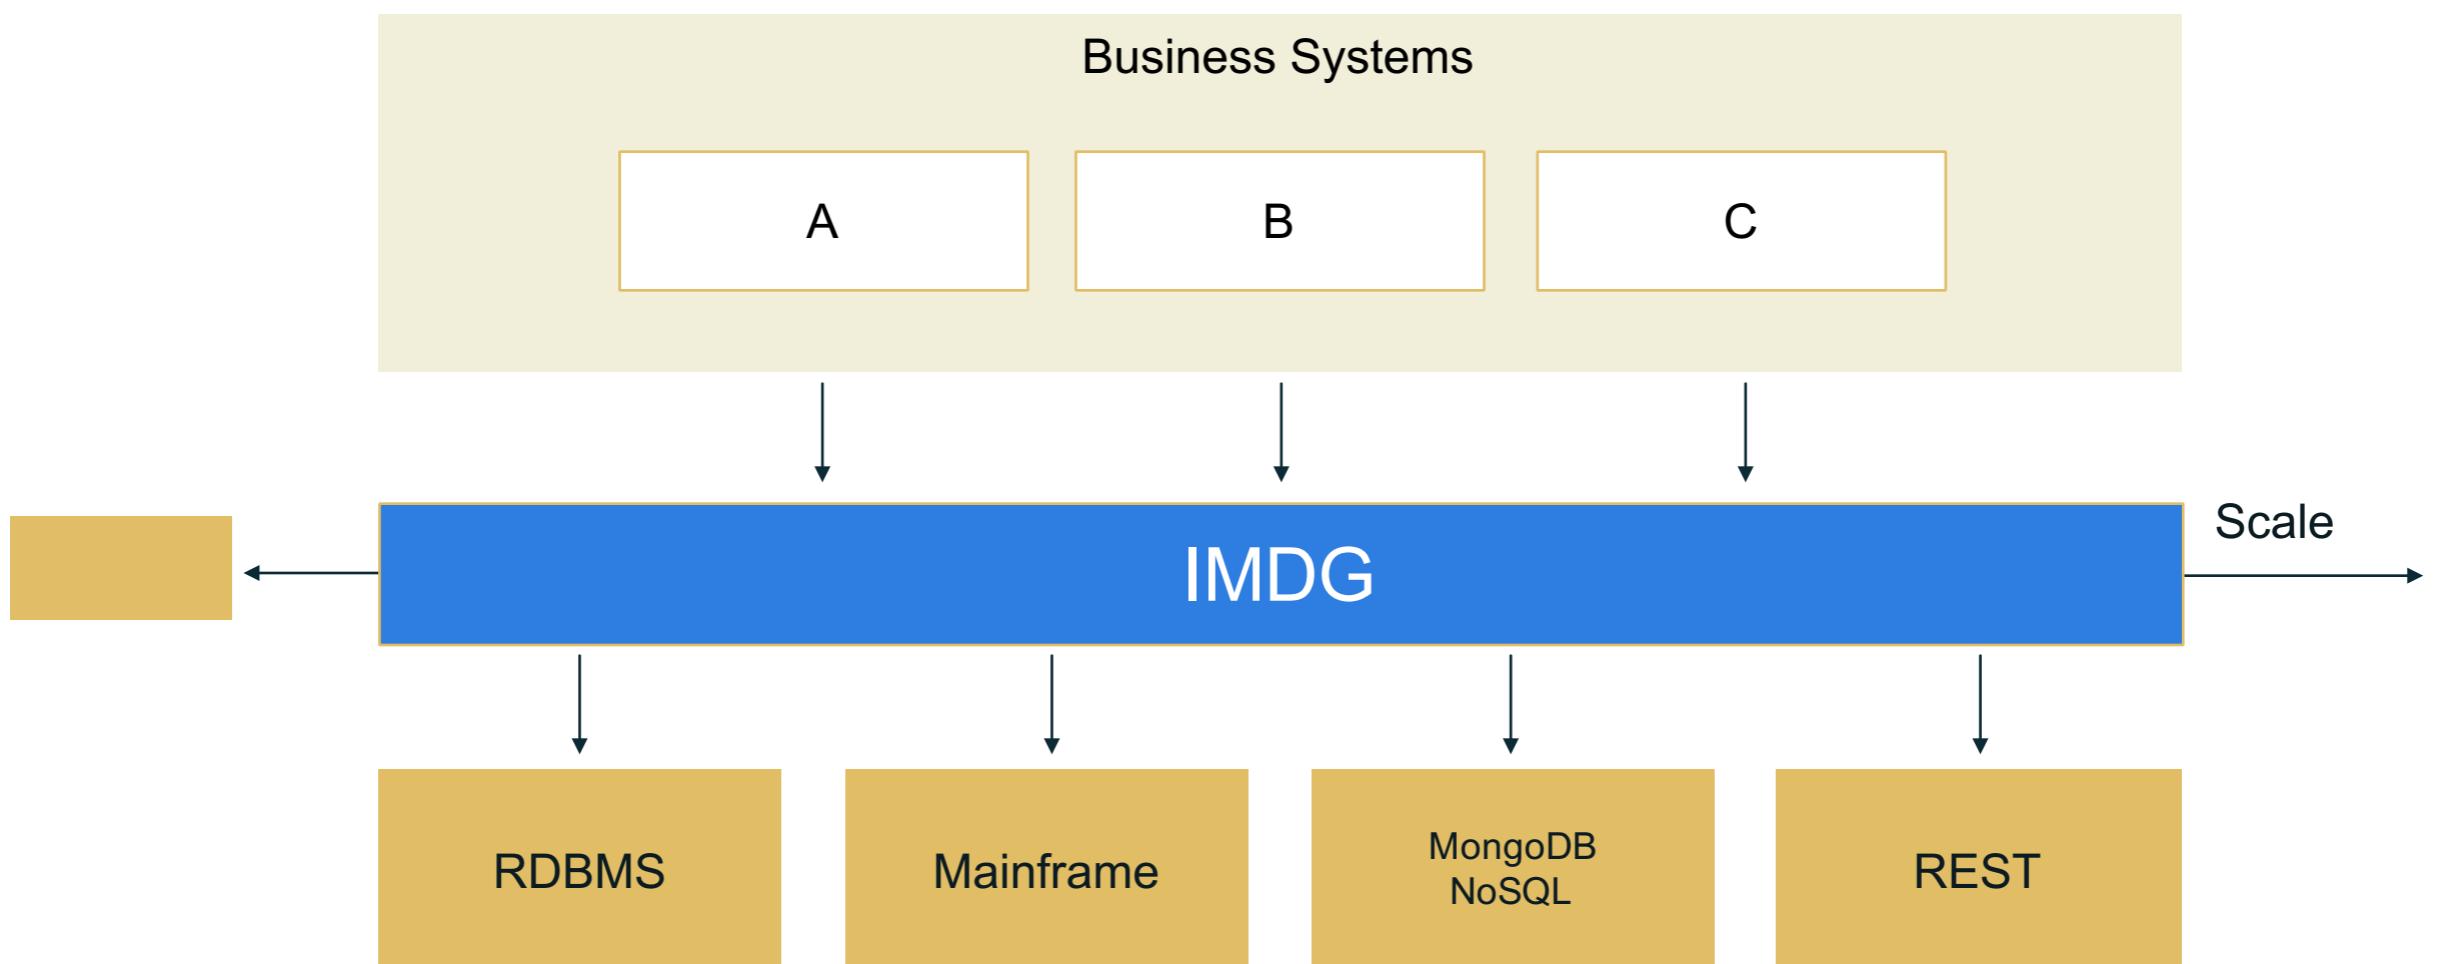


Application Scaling

- Elastic Scalability
- Super Speeds
- High Availability
- Fault Tolerance
- Cloud Readiness
 - EC2, GCE, Docker deployment

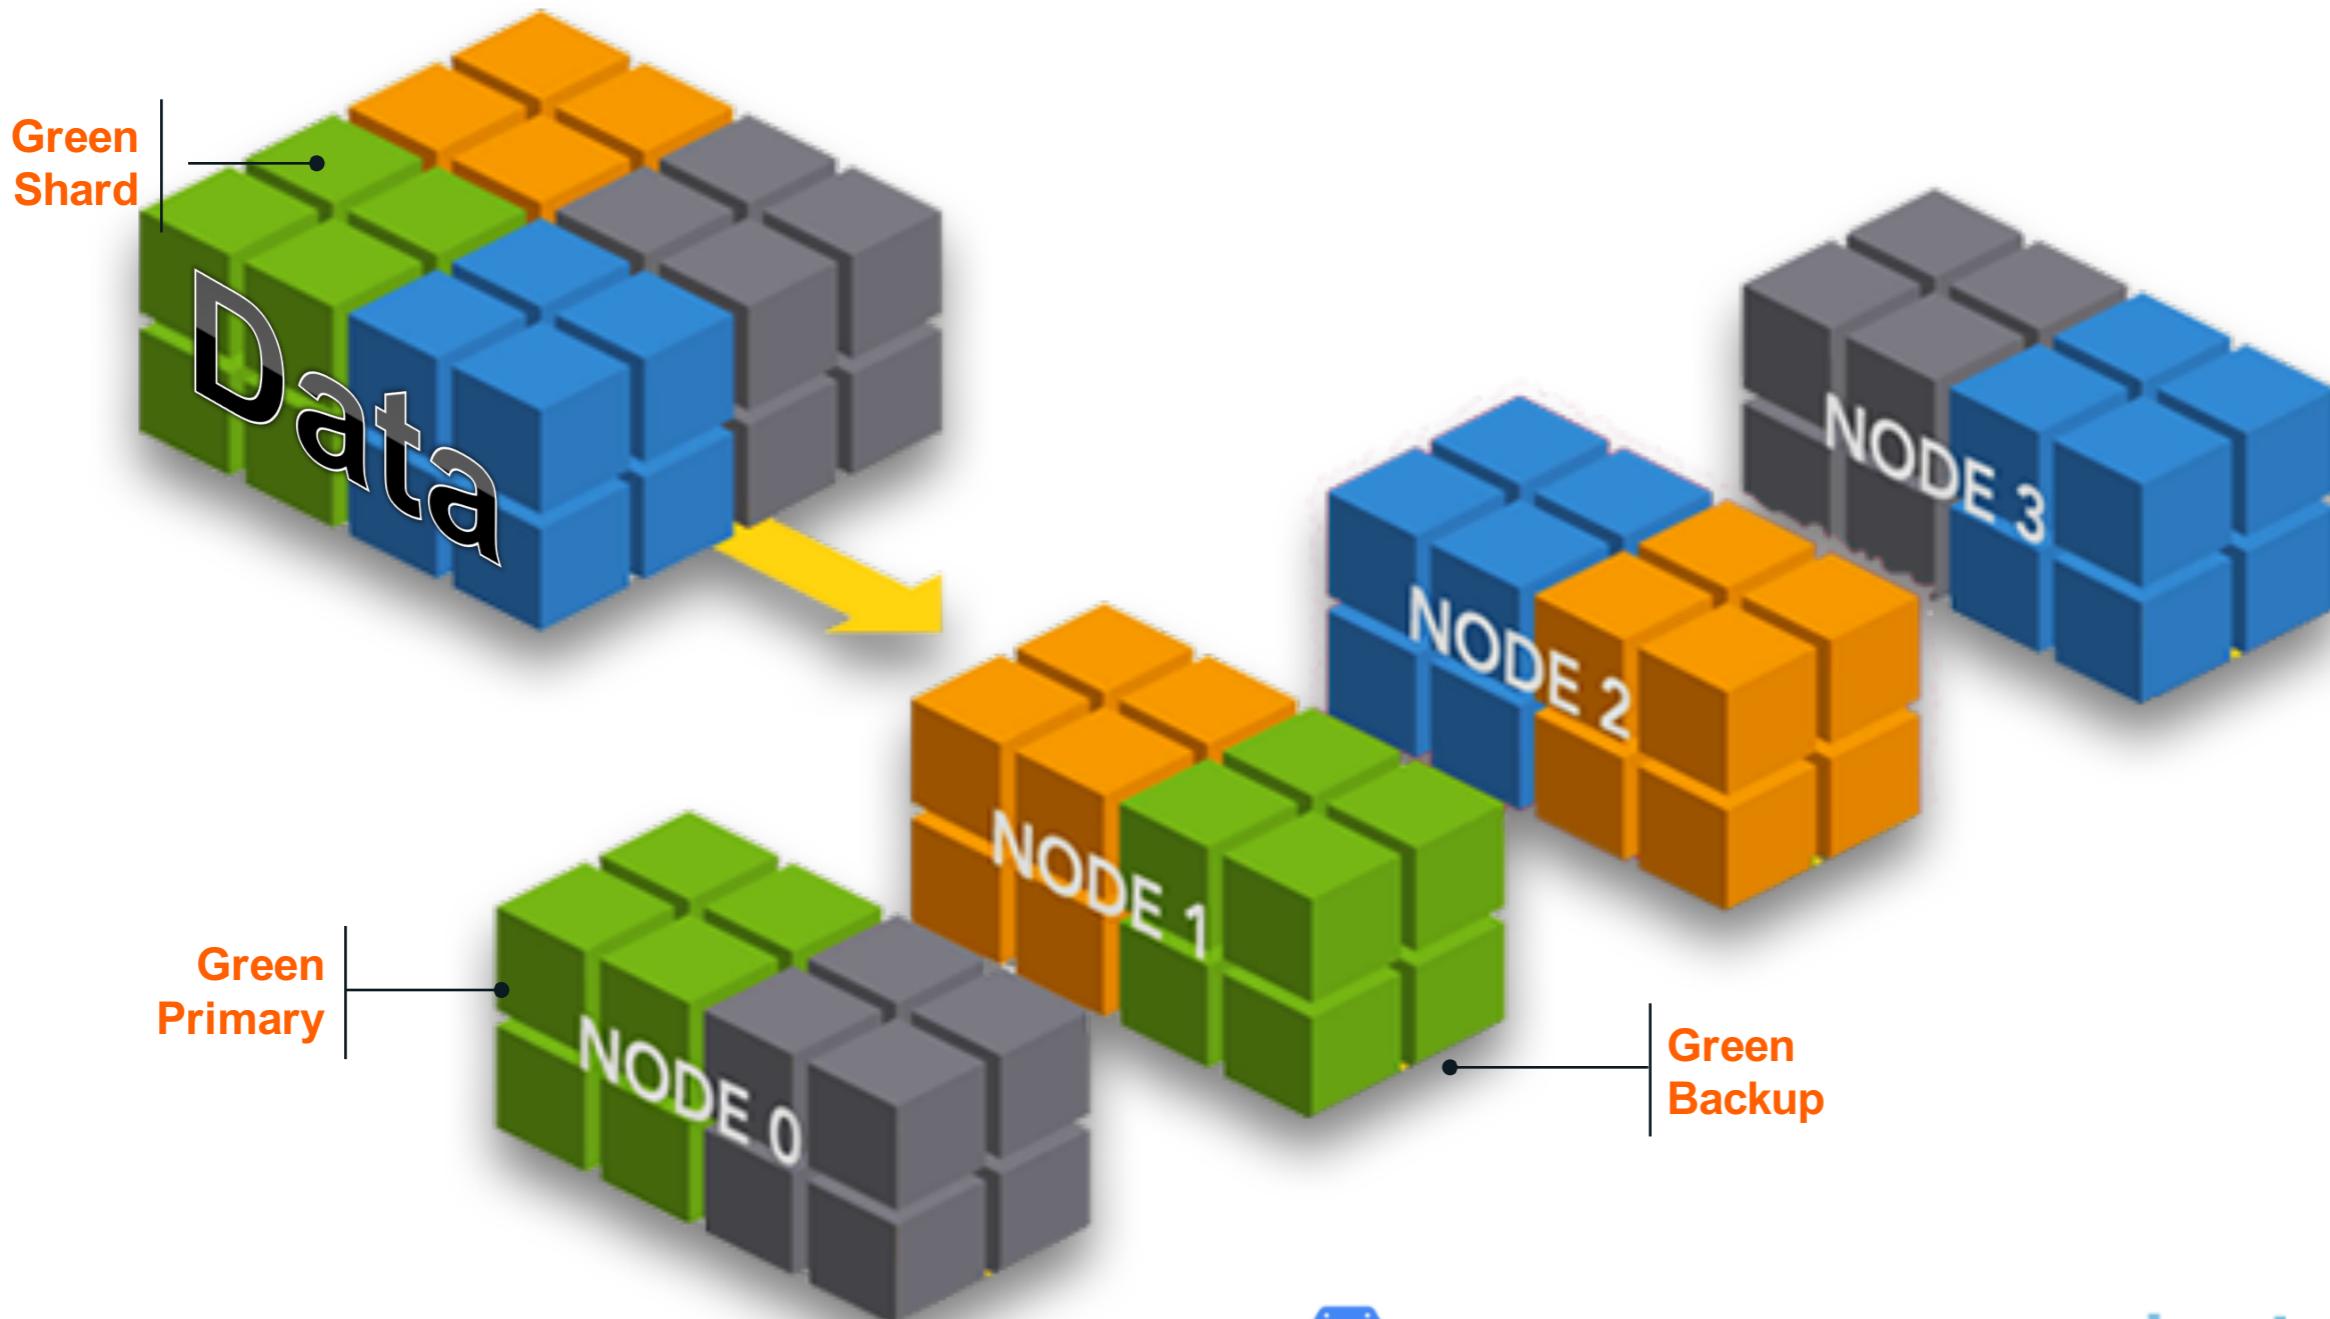


Unified Data Access / Database Caching



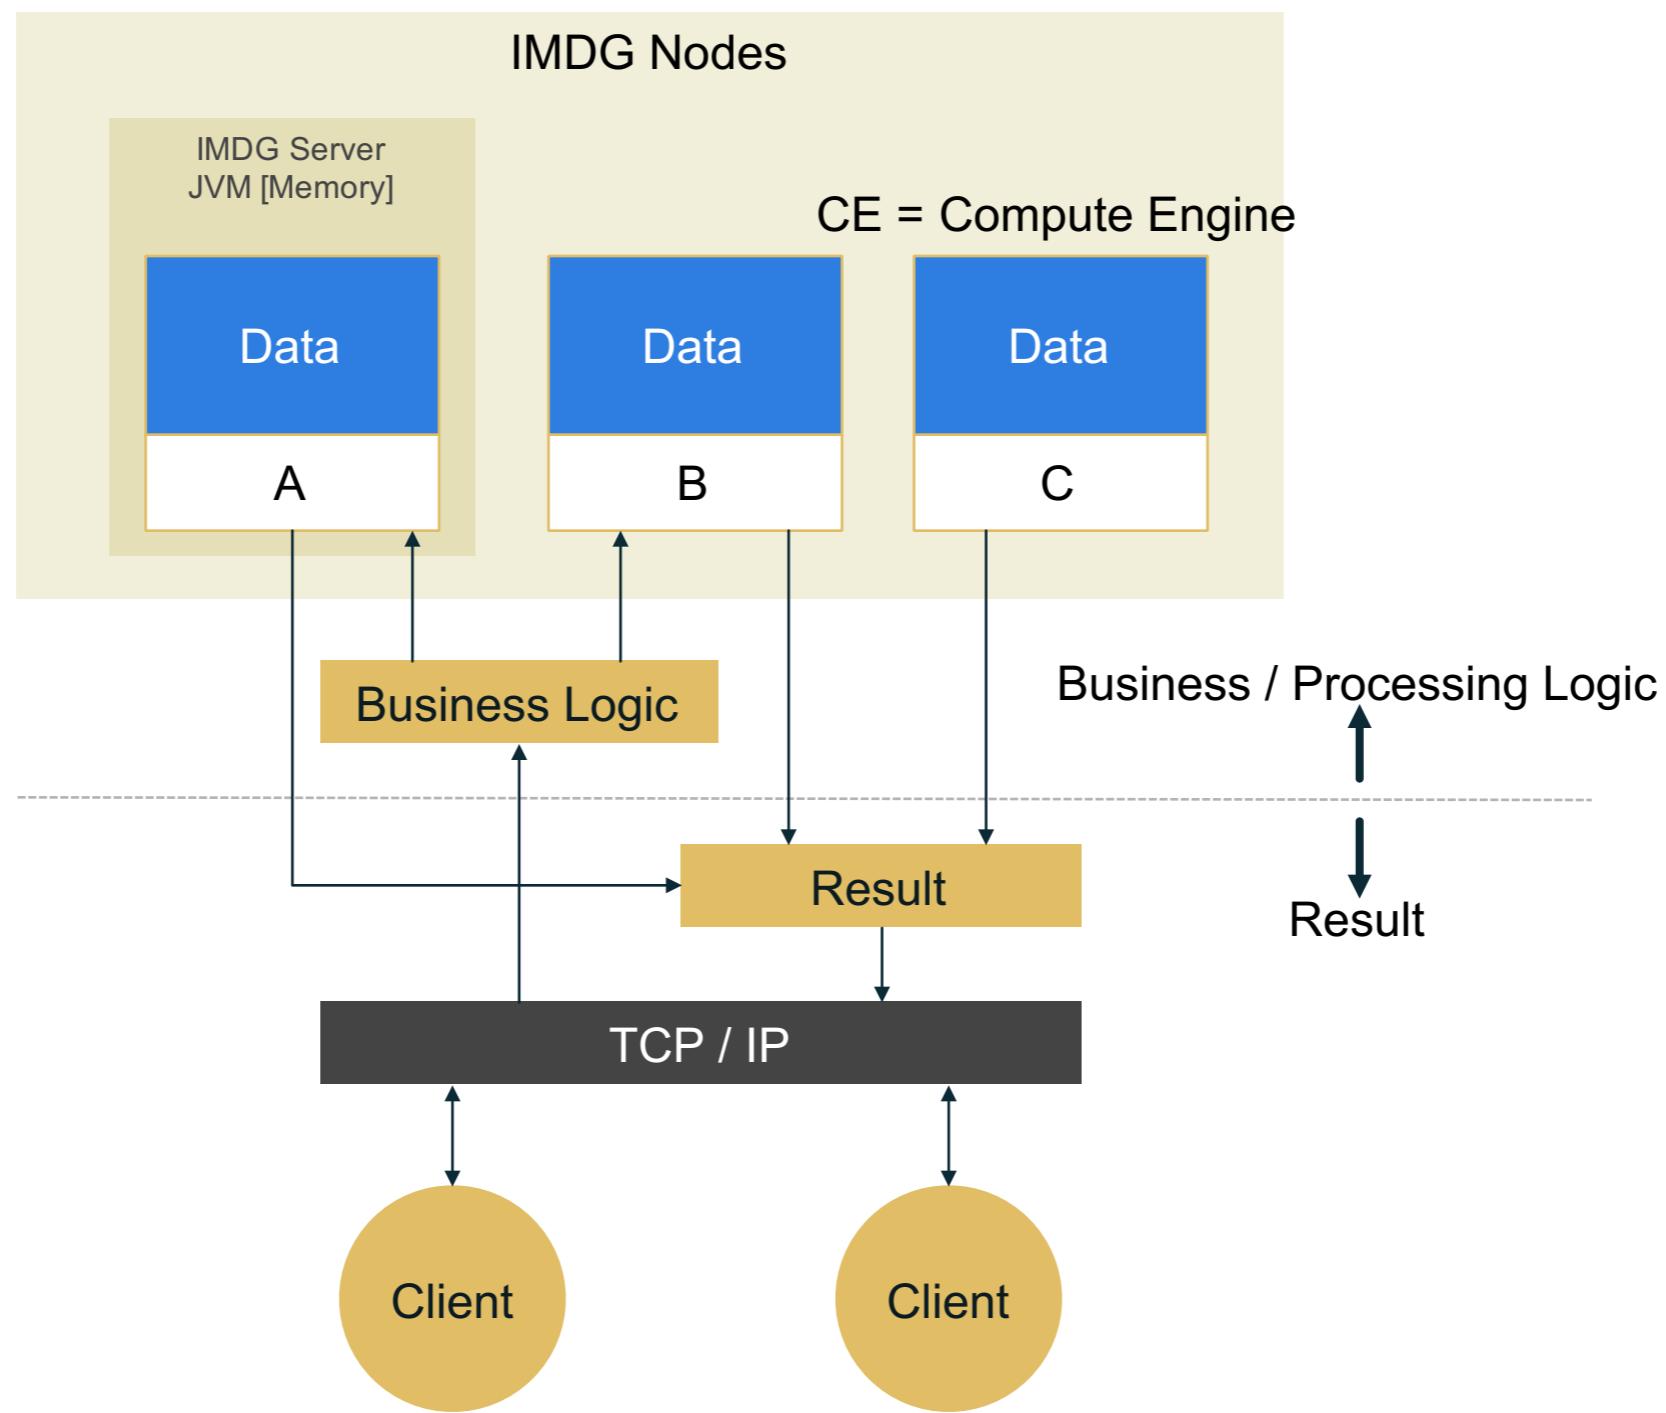
Hazelcast

Elastic Scalability





Distributed Computing





Reactive Clients

The clients (part of an application) enable you to do the operations without being a member of the grid. Clients can be:

- smart: this means that they immediately can send an operation like `map.get(key)` to the member that owns that specific key.
- dumb: it will connect to a random member in the cluster and send requests to this member. This member then needs to send the request to the correct member.

High availability

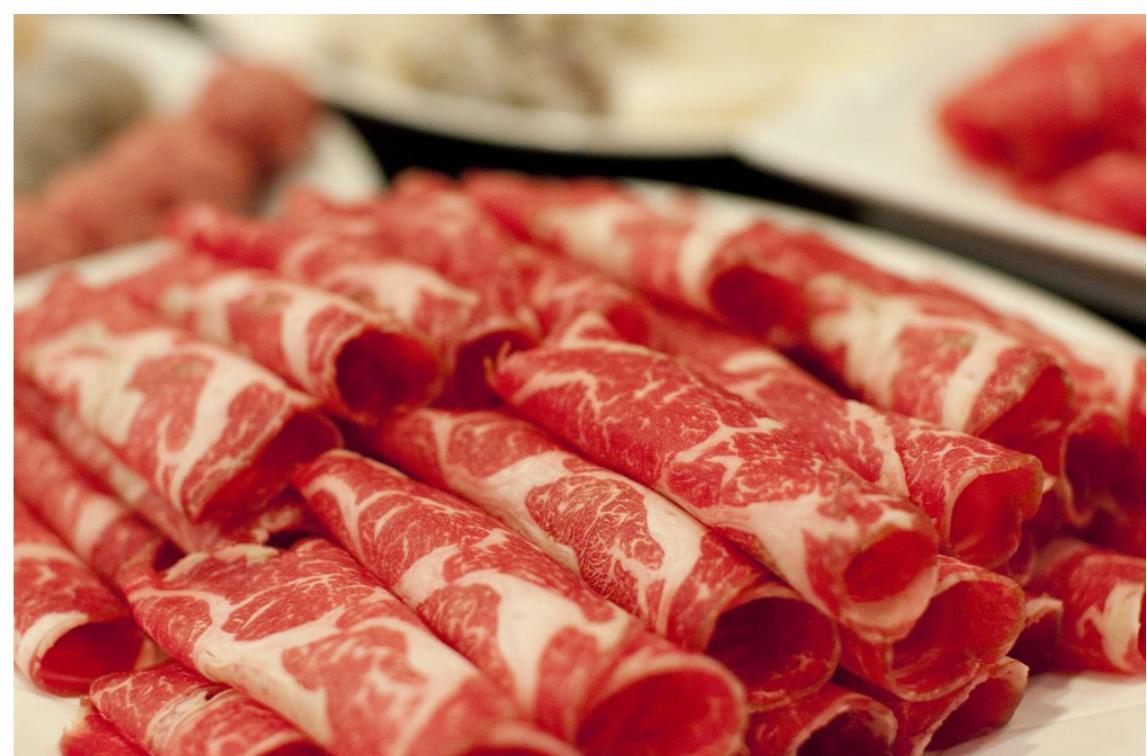
When the connected cluster member dies, client will automatically switch to another live member.



Making In-Memory Reliable, Scalable and Durable

Data Distribution







Data Distribution

Replication - Copying an entire dataset onto multiple servers.
Used for improving speed of access to reference records such as master data.

Partitioning - Splitting up a large monolithic dataset into multiple smaller sets based on data cohesion.

You need to restart your computer. Hold down the Power button for several seconds or press the Restart button.

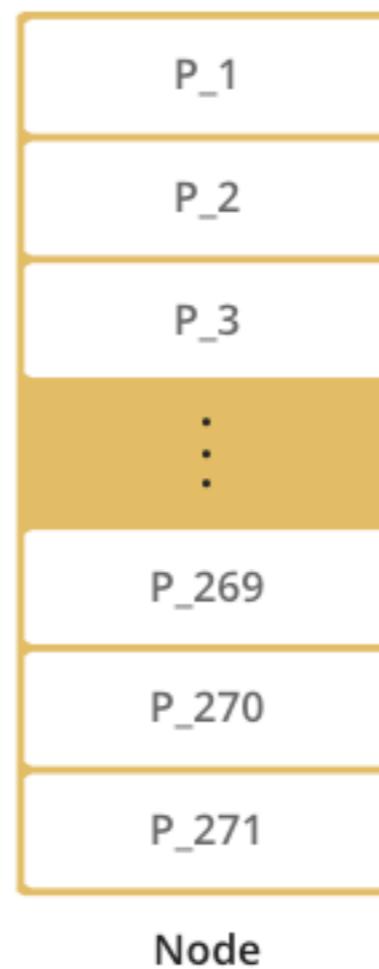
Veuillez redémarrer votre ordinateur. Maintenez la touche de démarrage enfoncée pendant plusieurs secondes ou bien appuyez sur le bouton de réinitialisation.

Sie müssen Ihren Computer neu starten. Halten Sie dazu die Einschalttaste einige Sekunden gedrückt oder drücken Sie die Neustart-Taste.

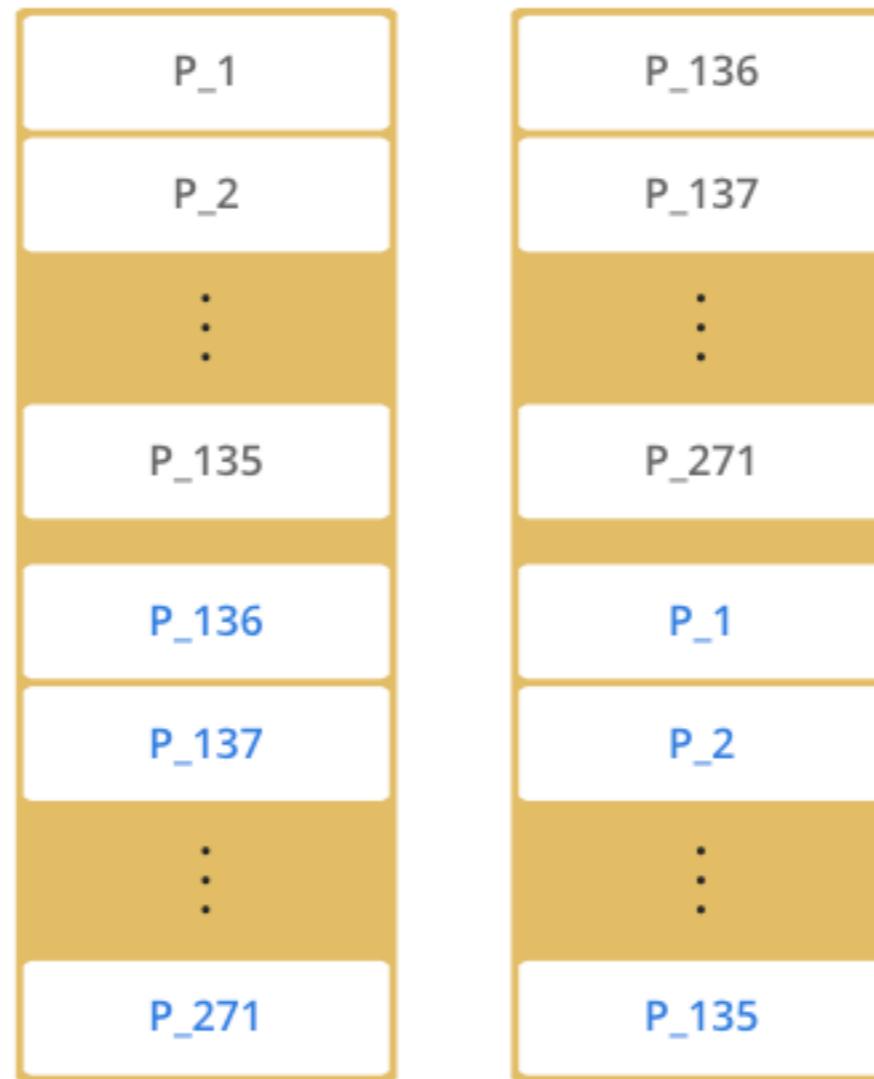
コンピュータを再起動する必要があります。パワー ボタンを数秒間押し続けるか、リセットボタンを押してください。



Data Partitioning

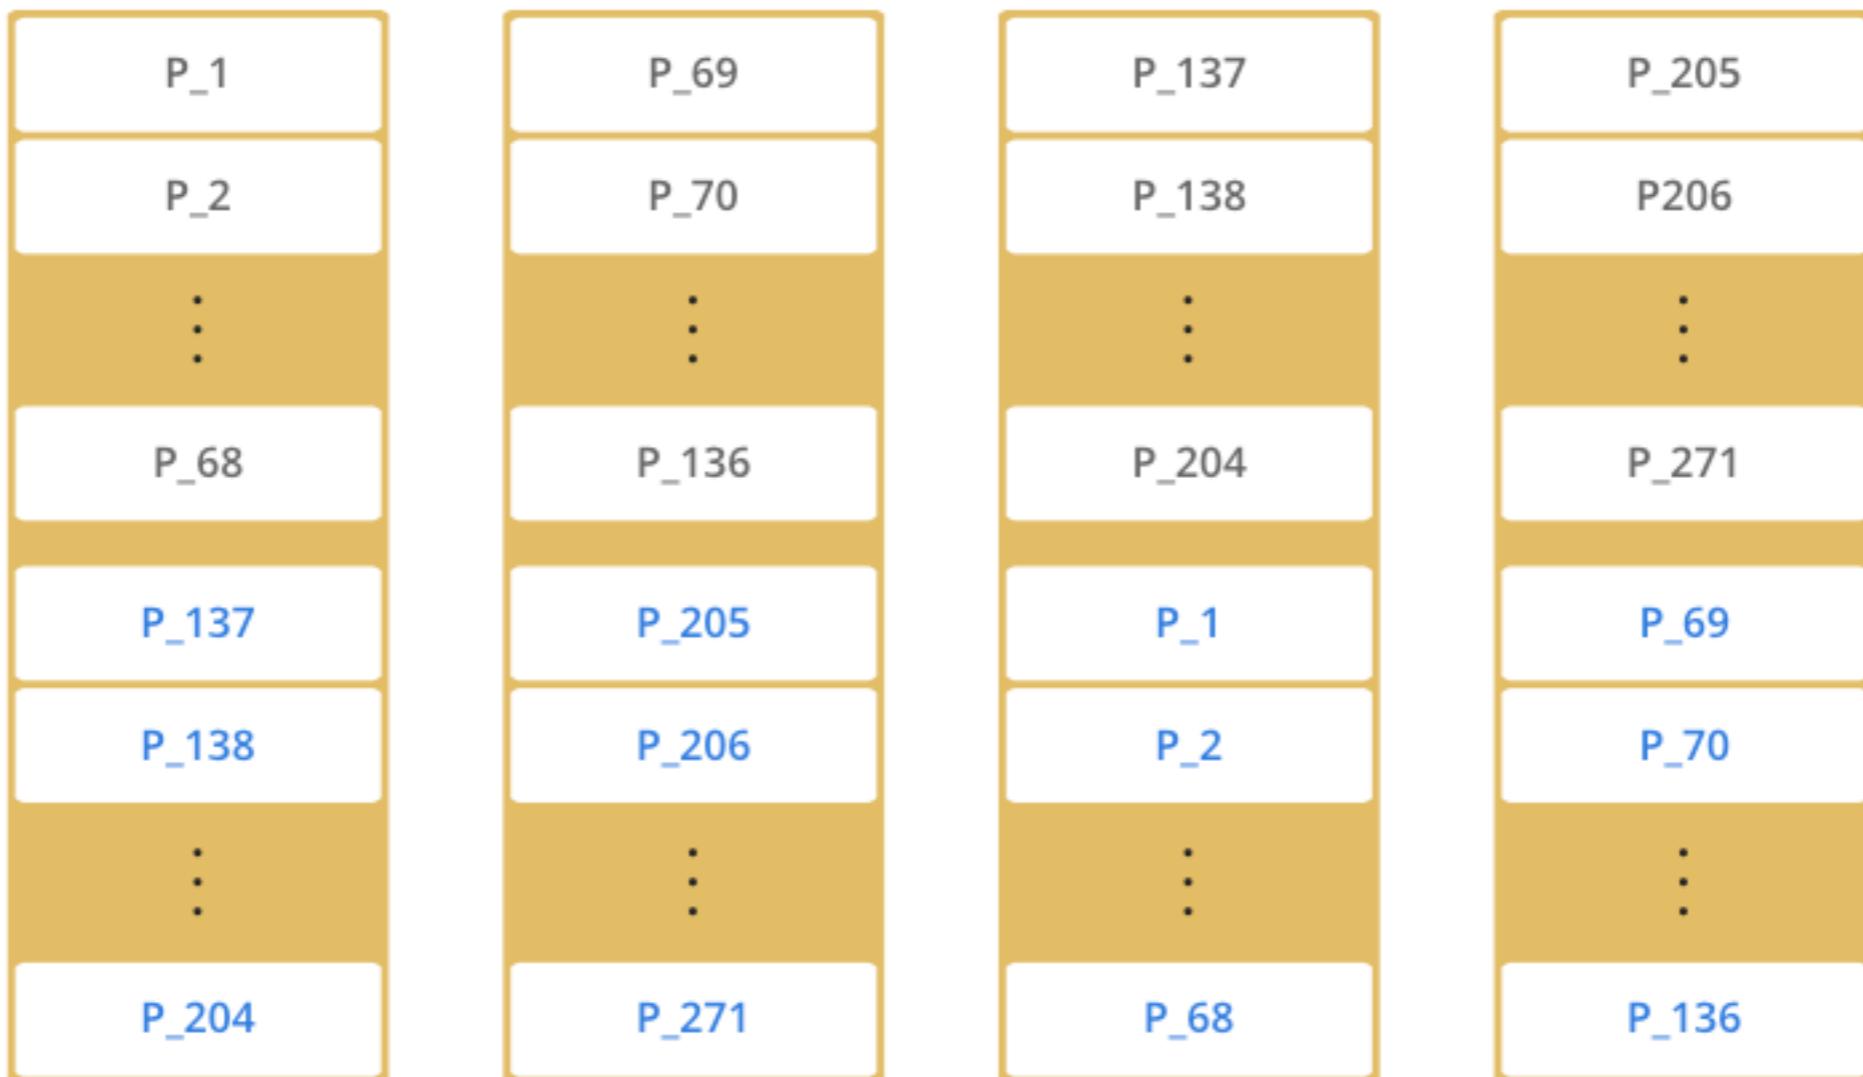


Data Partitioning





Data Partitioning





Rebalance Data On New Node

I DON'T ALWAYS BACKUP THE DATA



BUT WHEN I DO I BACKUP IT IN-MEMORY

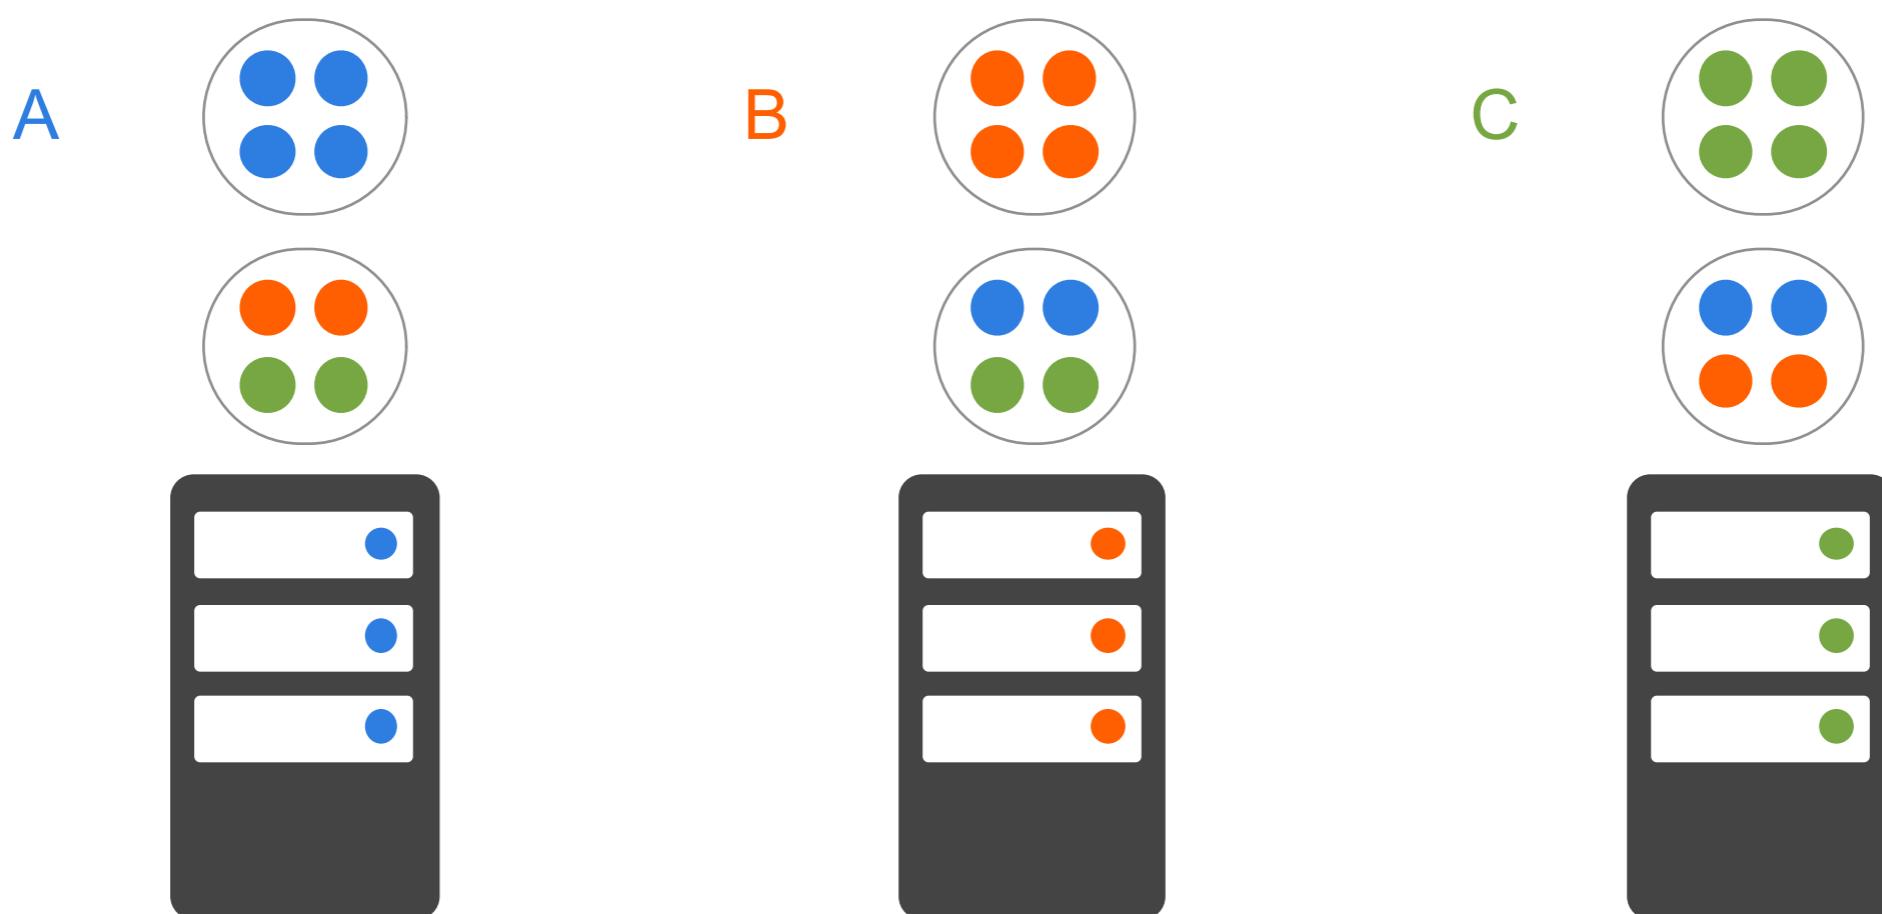
Consistent Hashing

Fixed number of partitions (default 271)

Each key falls into a partition

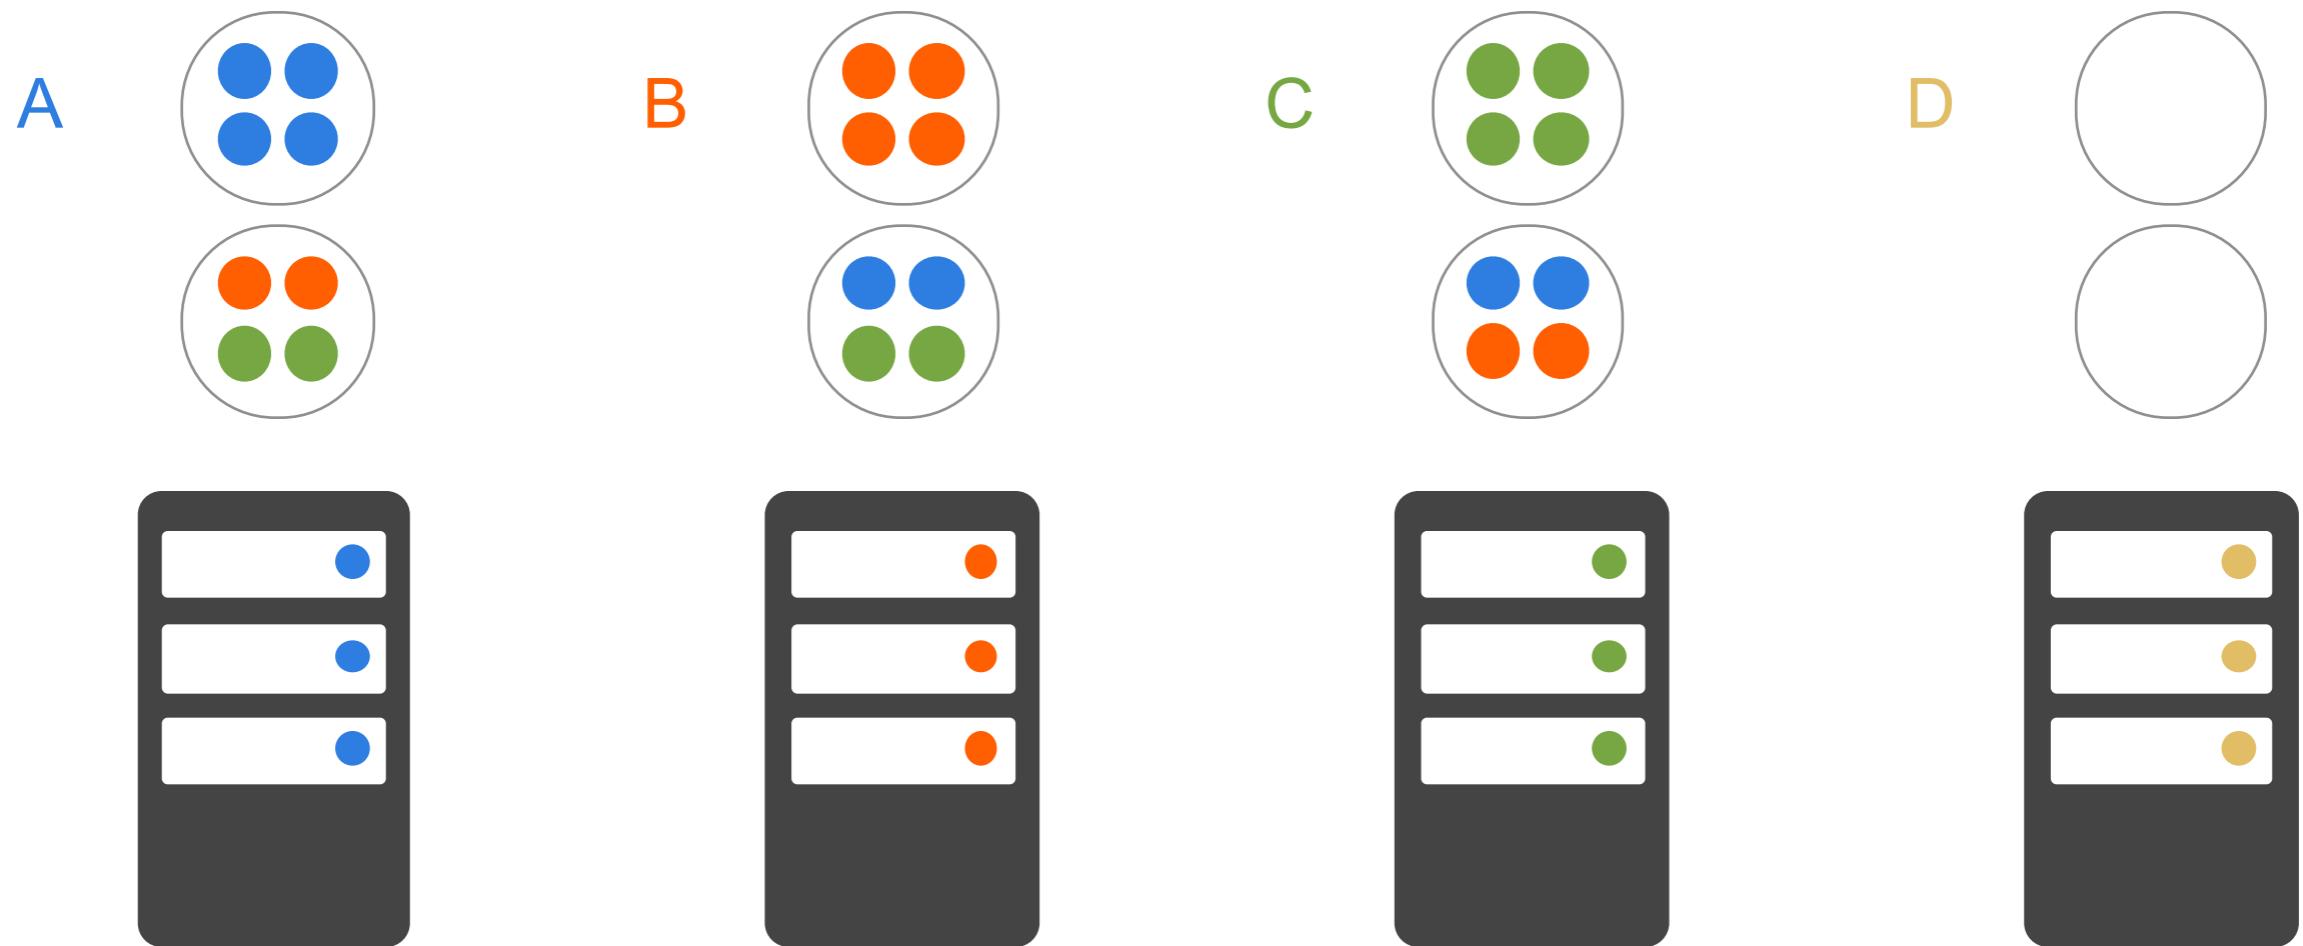
`partitionId = hash(keyData)%PARTITION_COUNT`

Partition ownerships are reassigned upon membership change



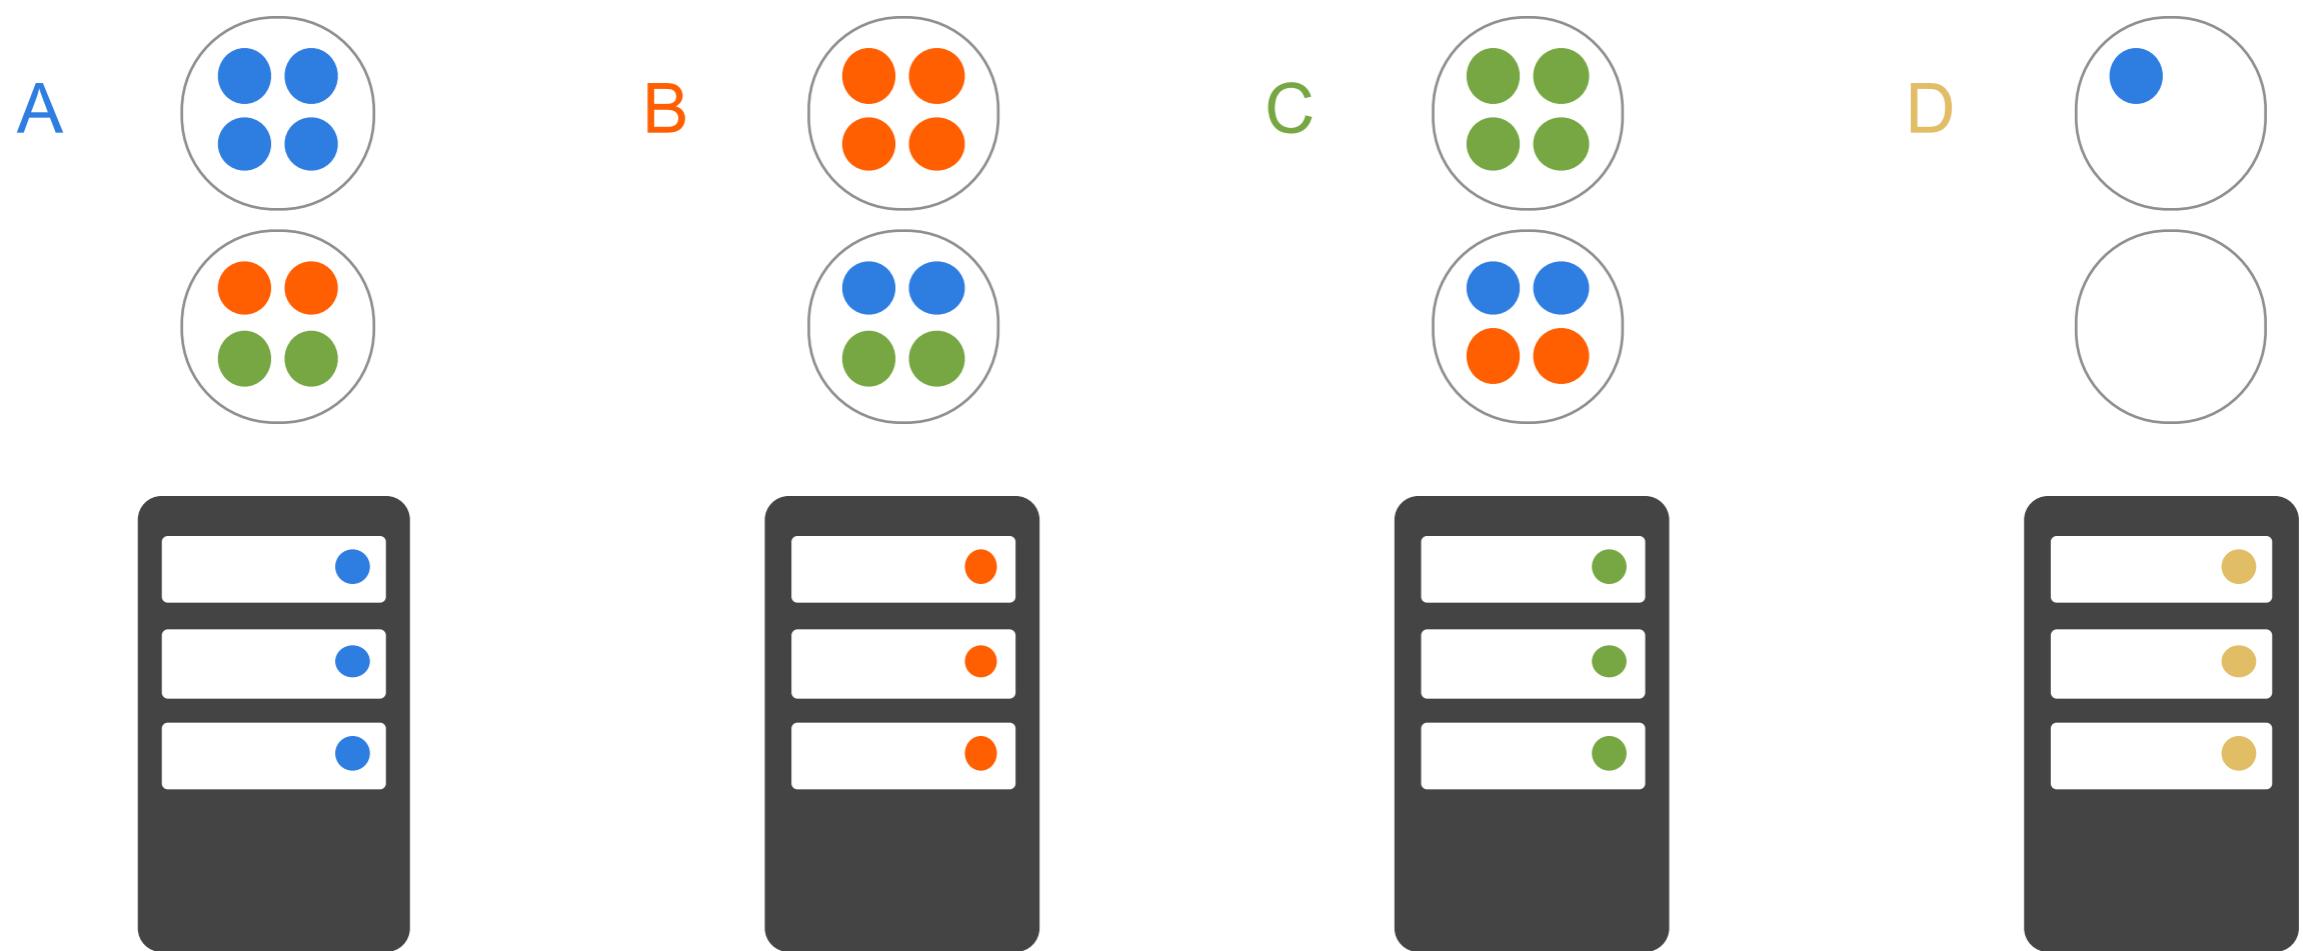


New Node Added



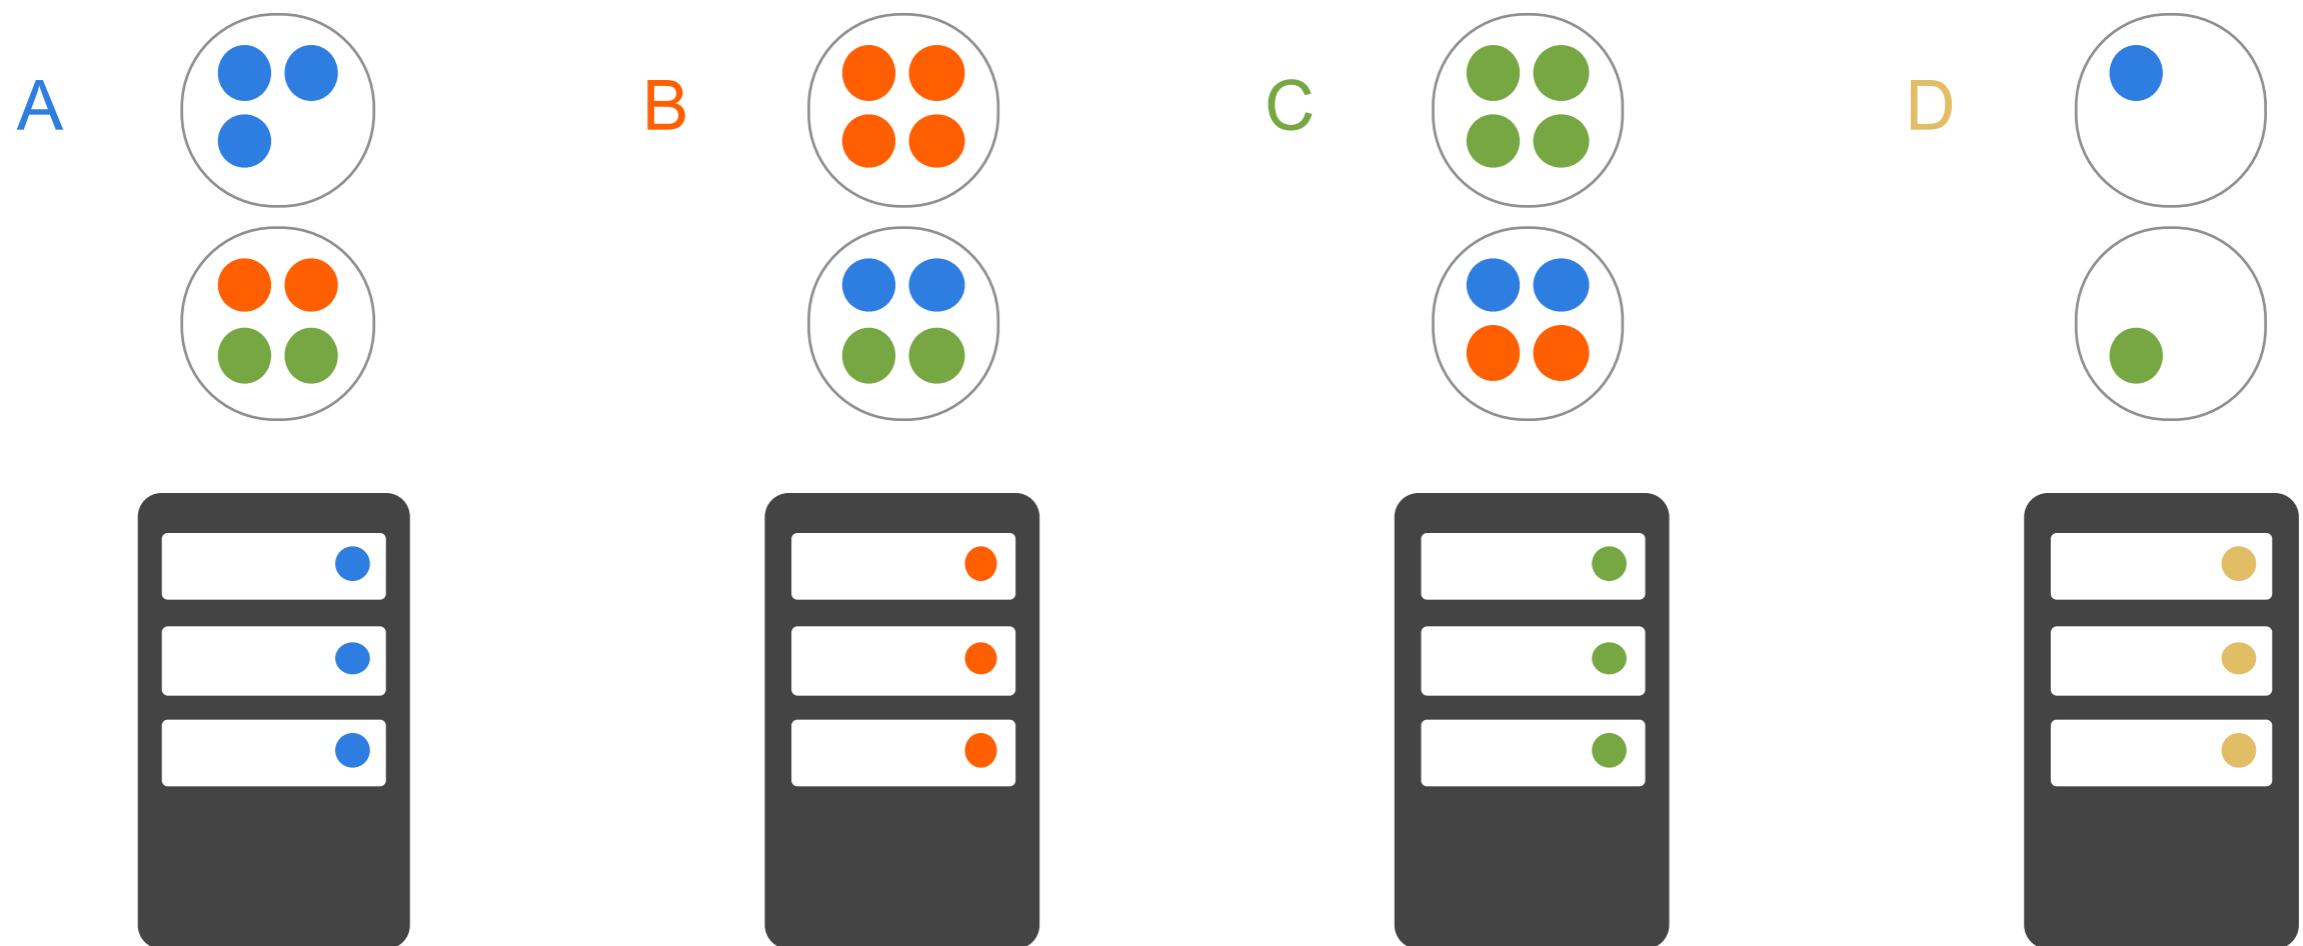


Migration



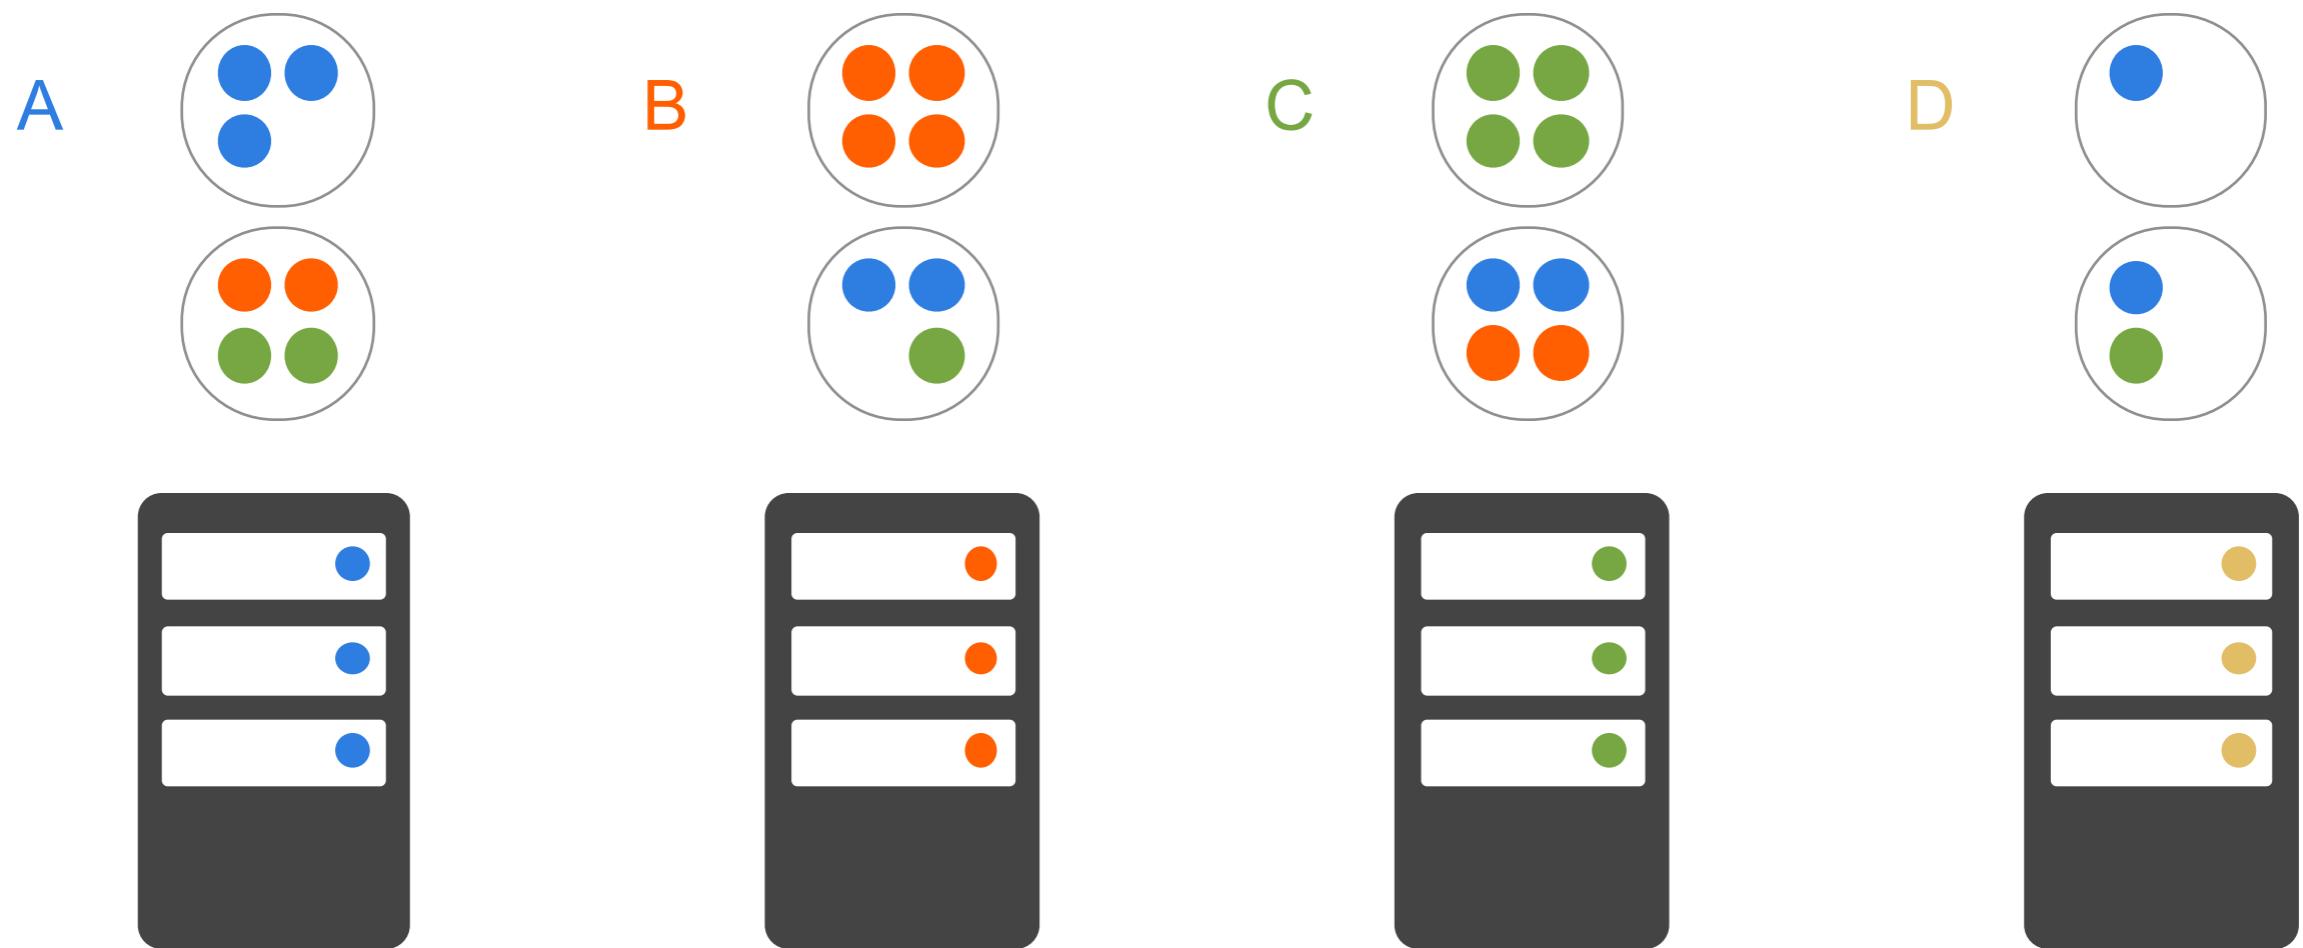


Migration



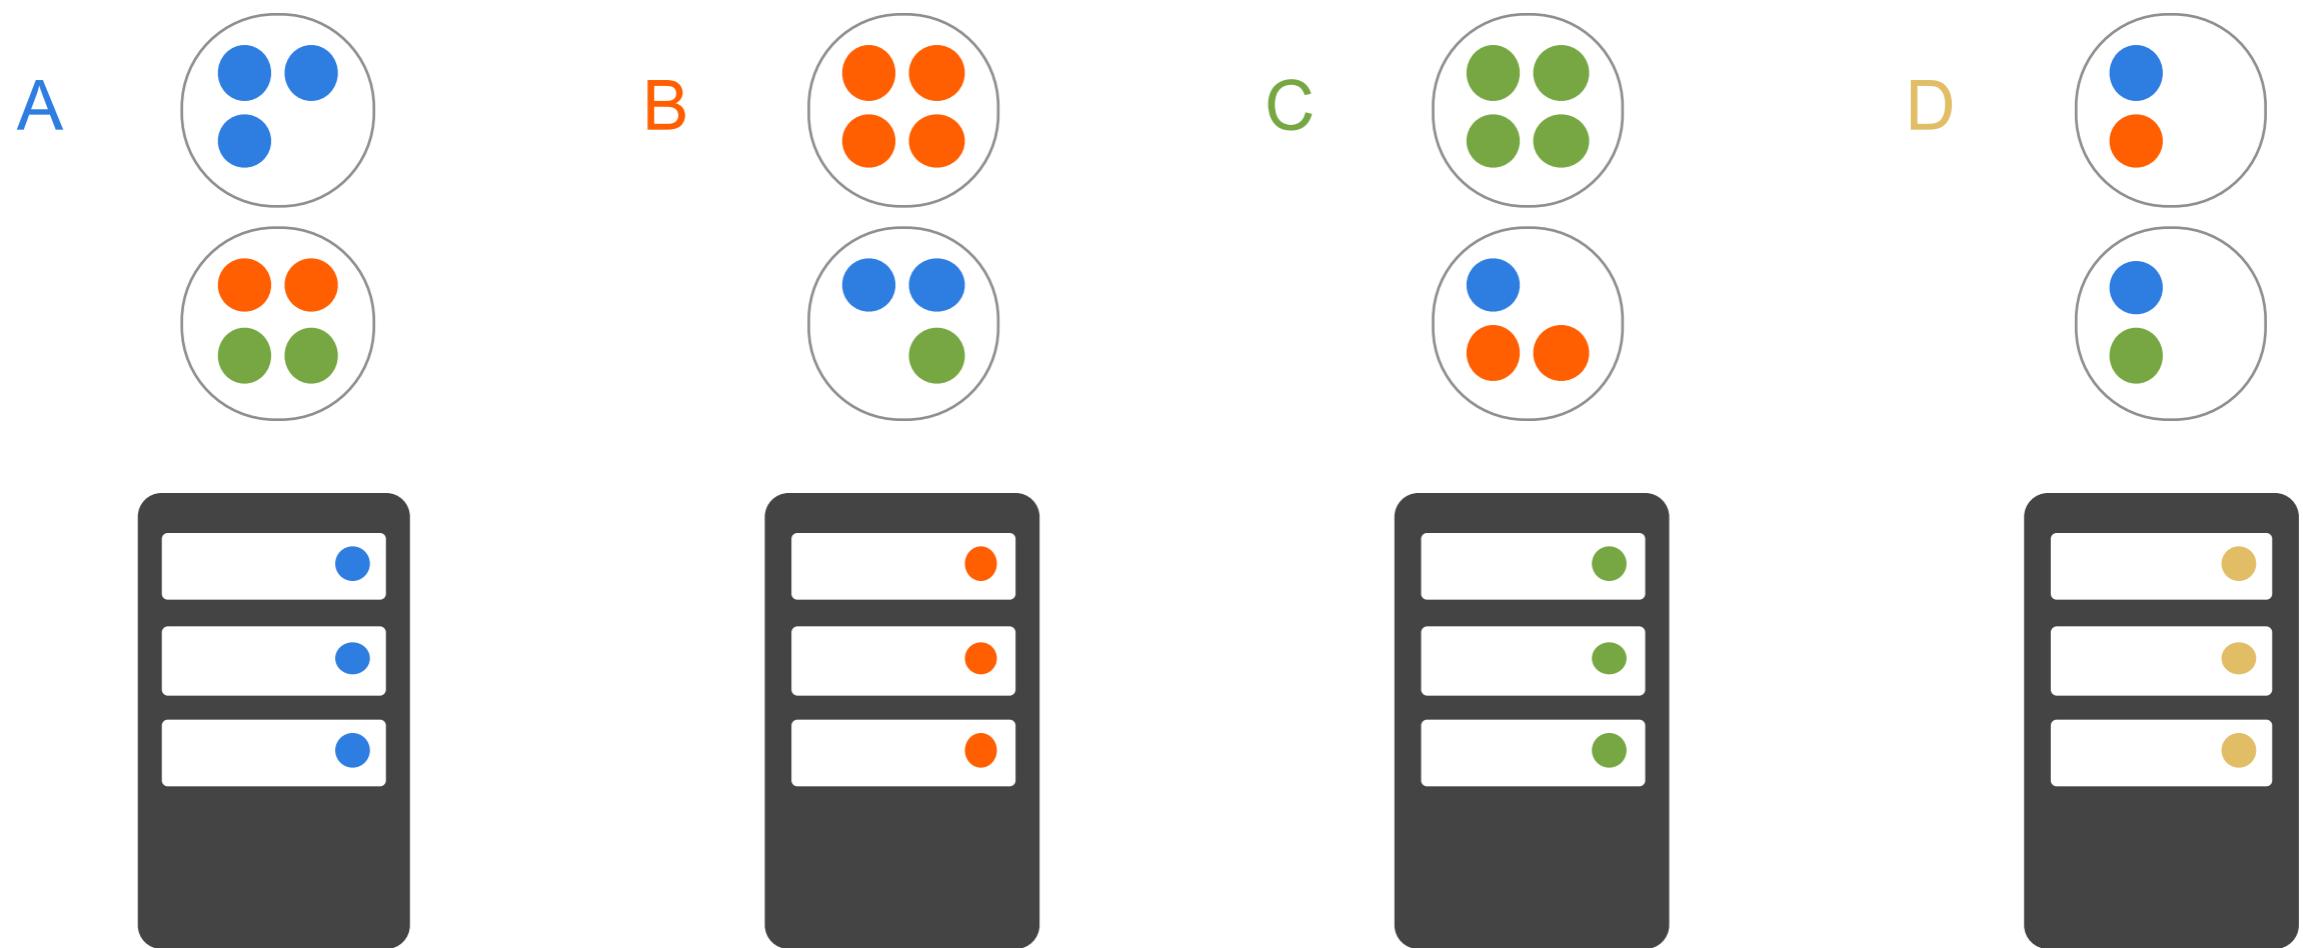


Migration



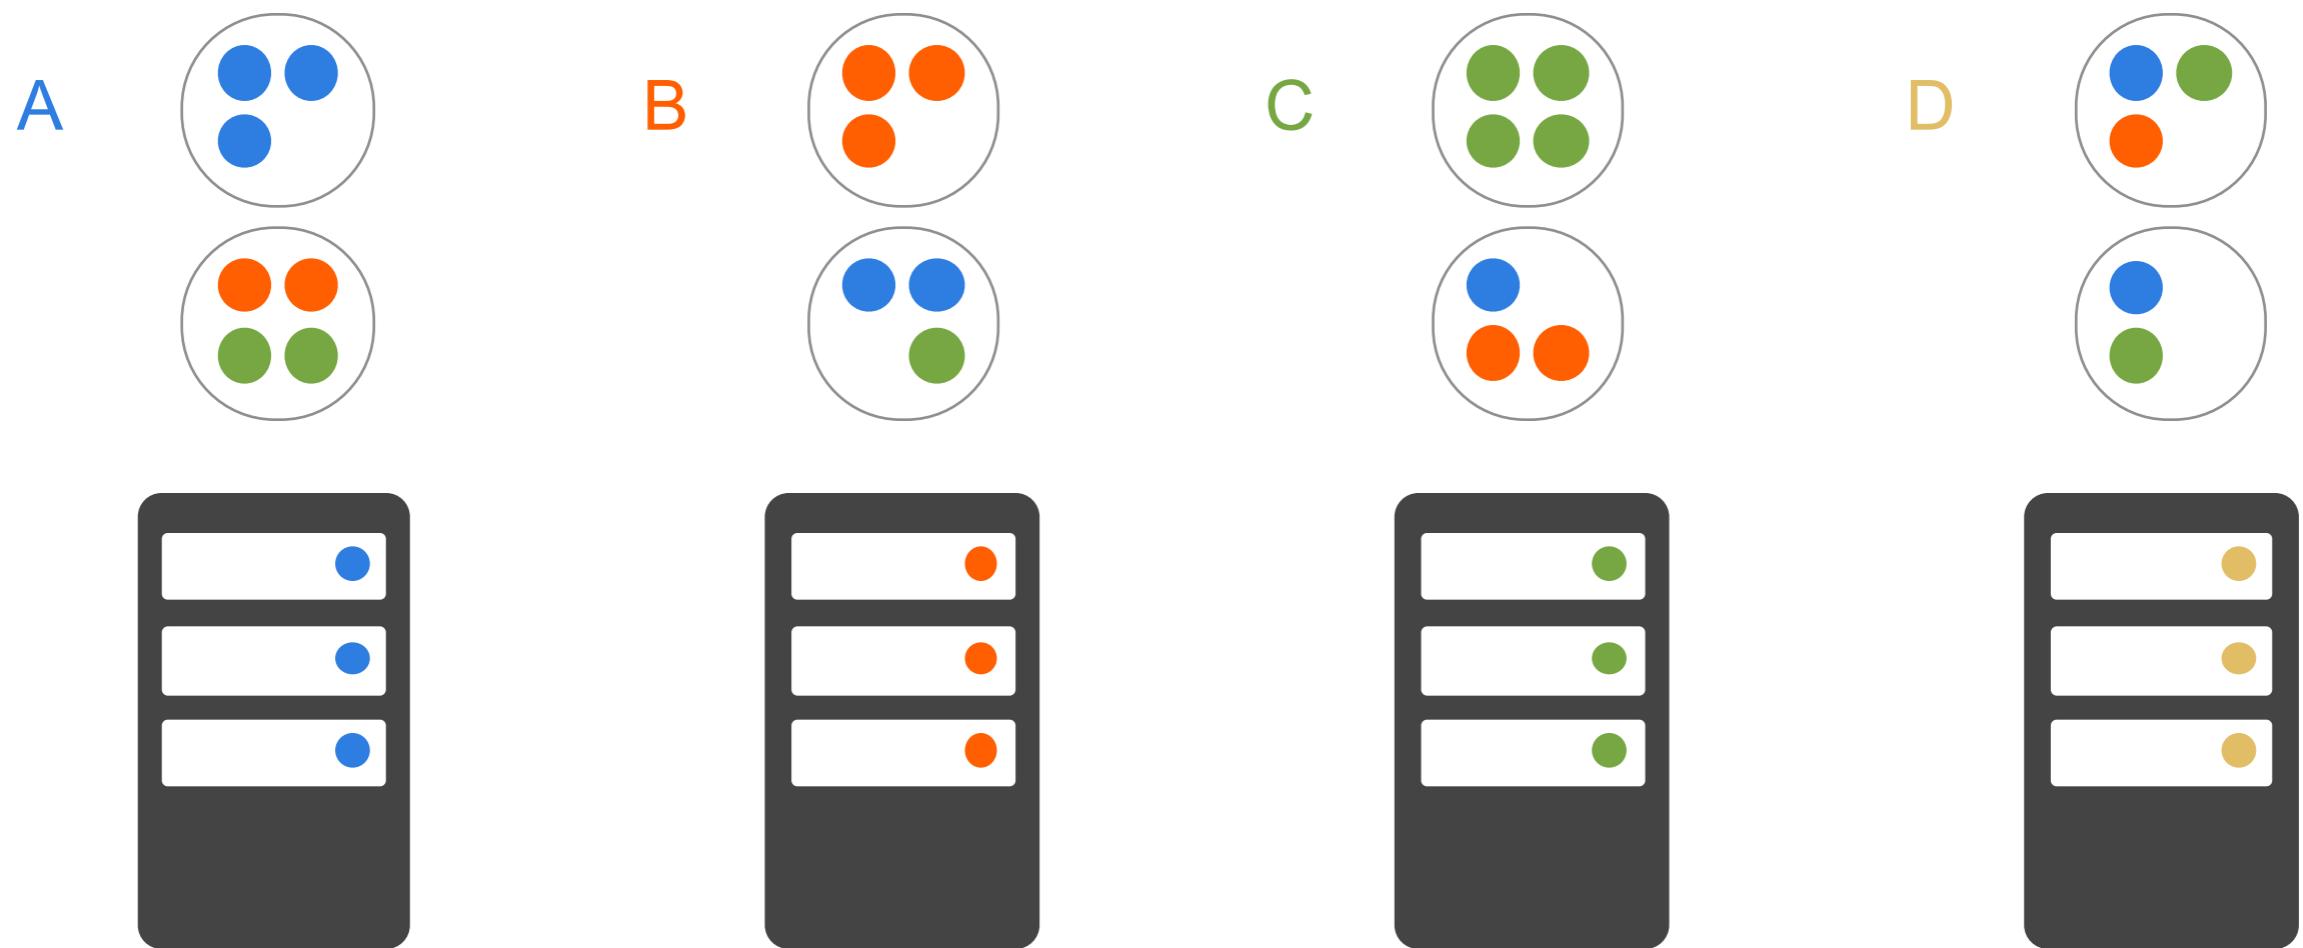


Migration



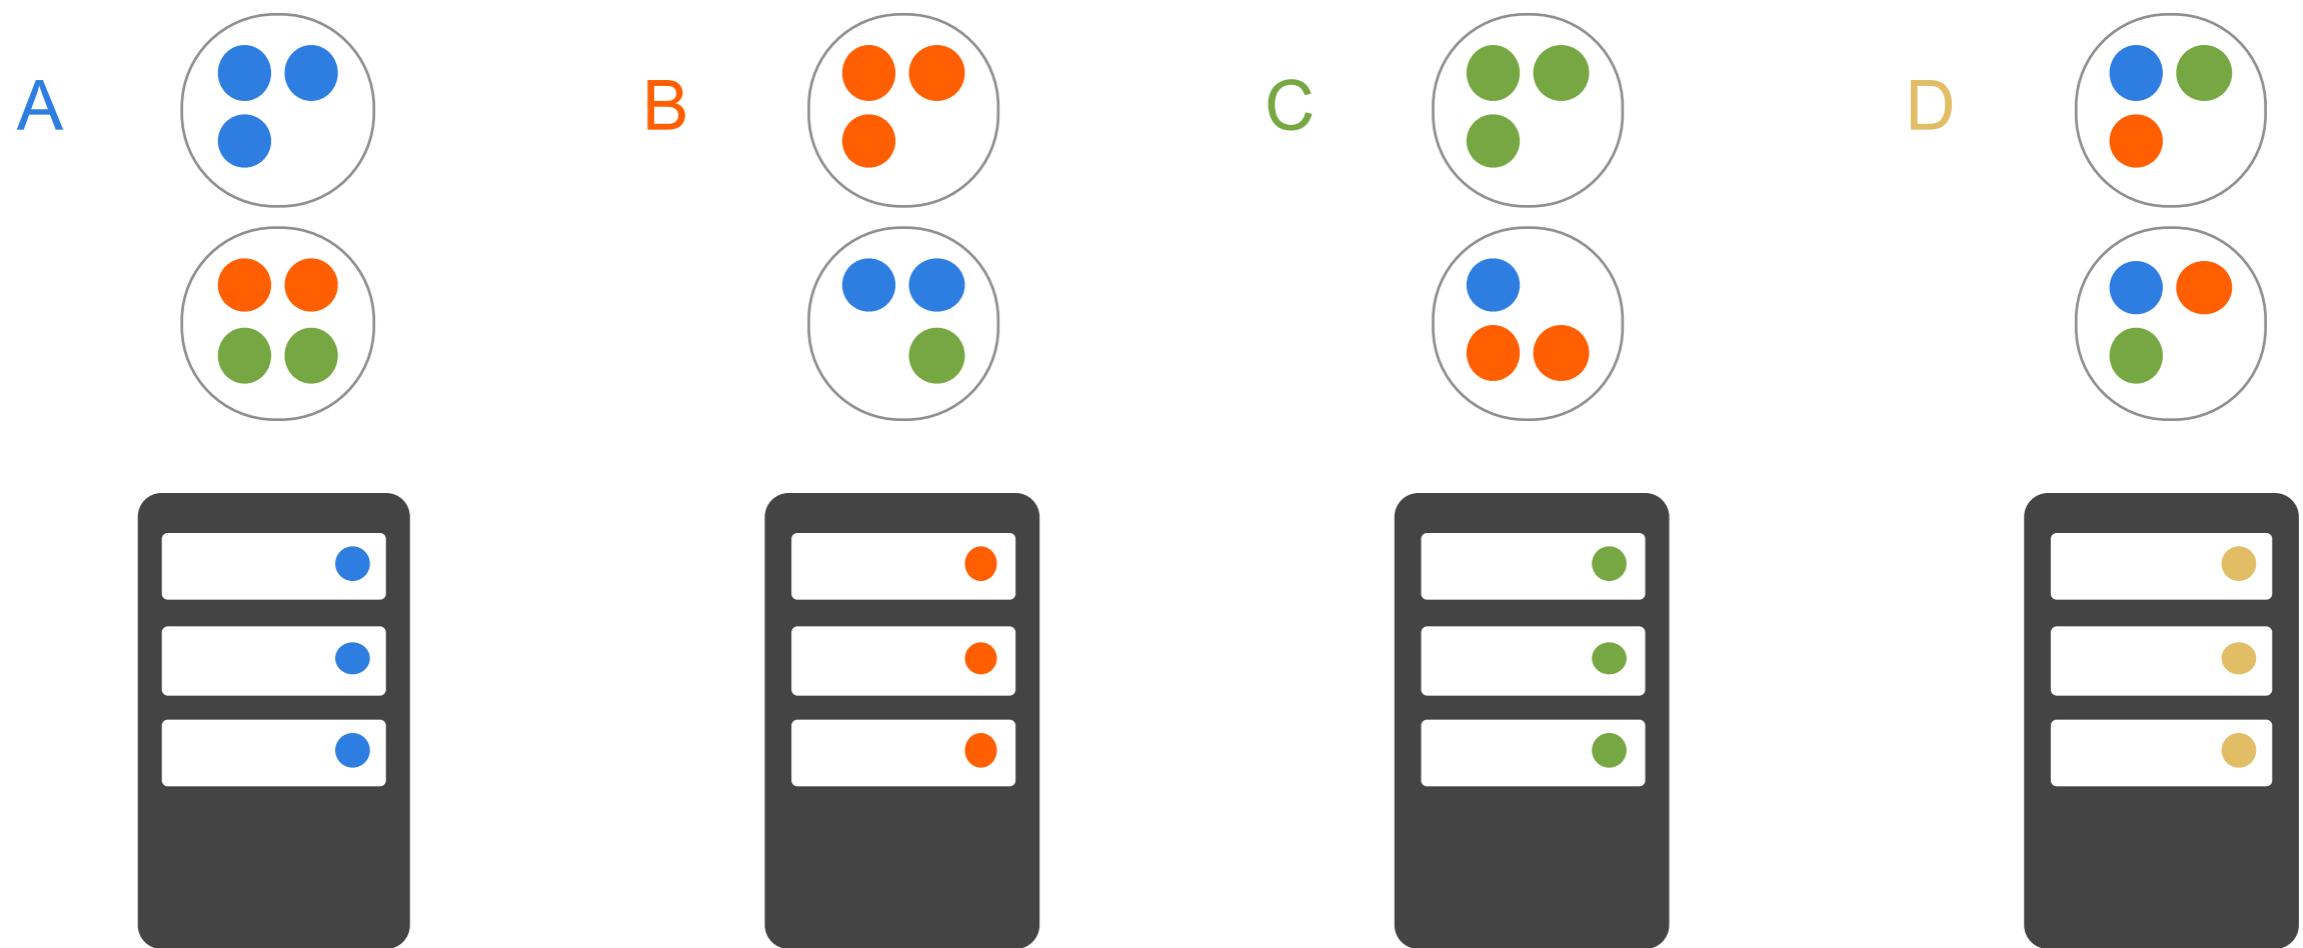


Migration



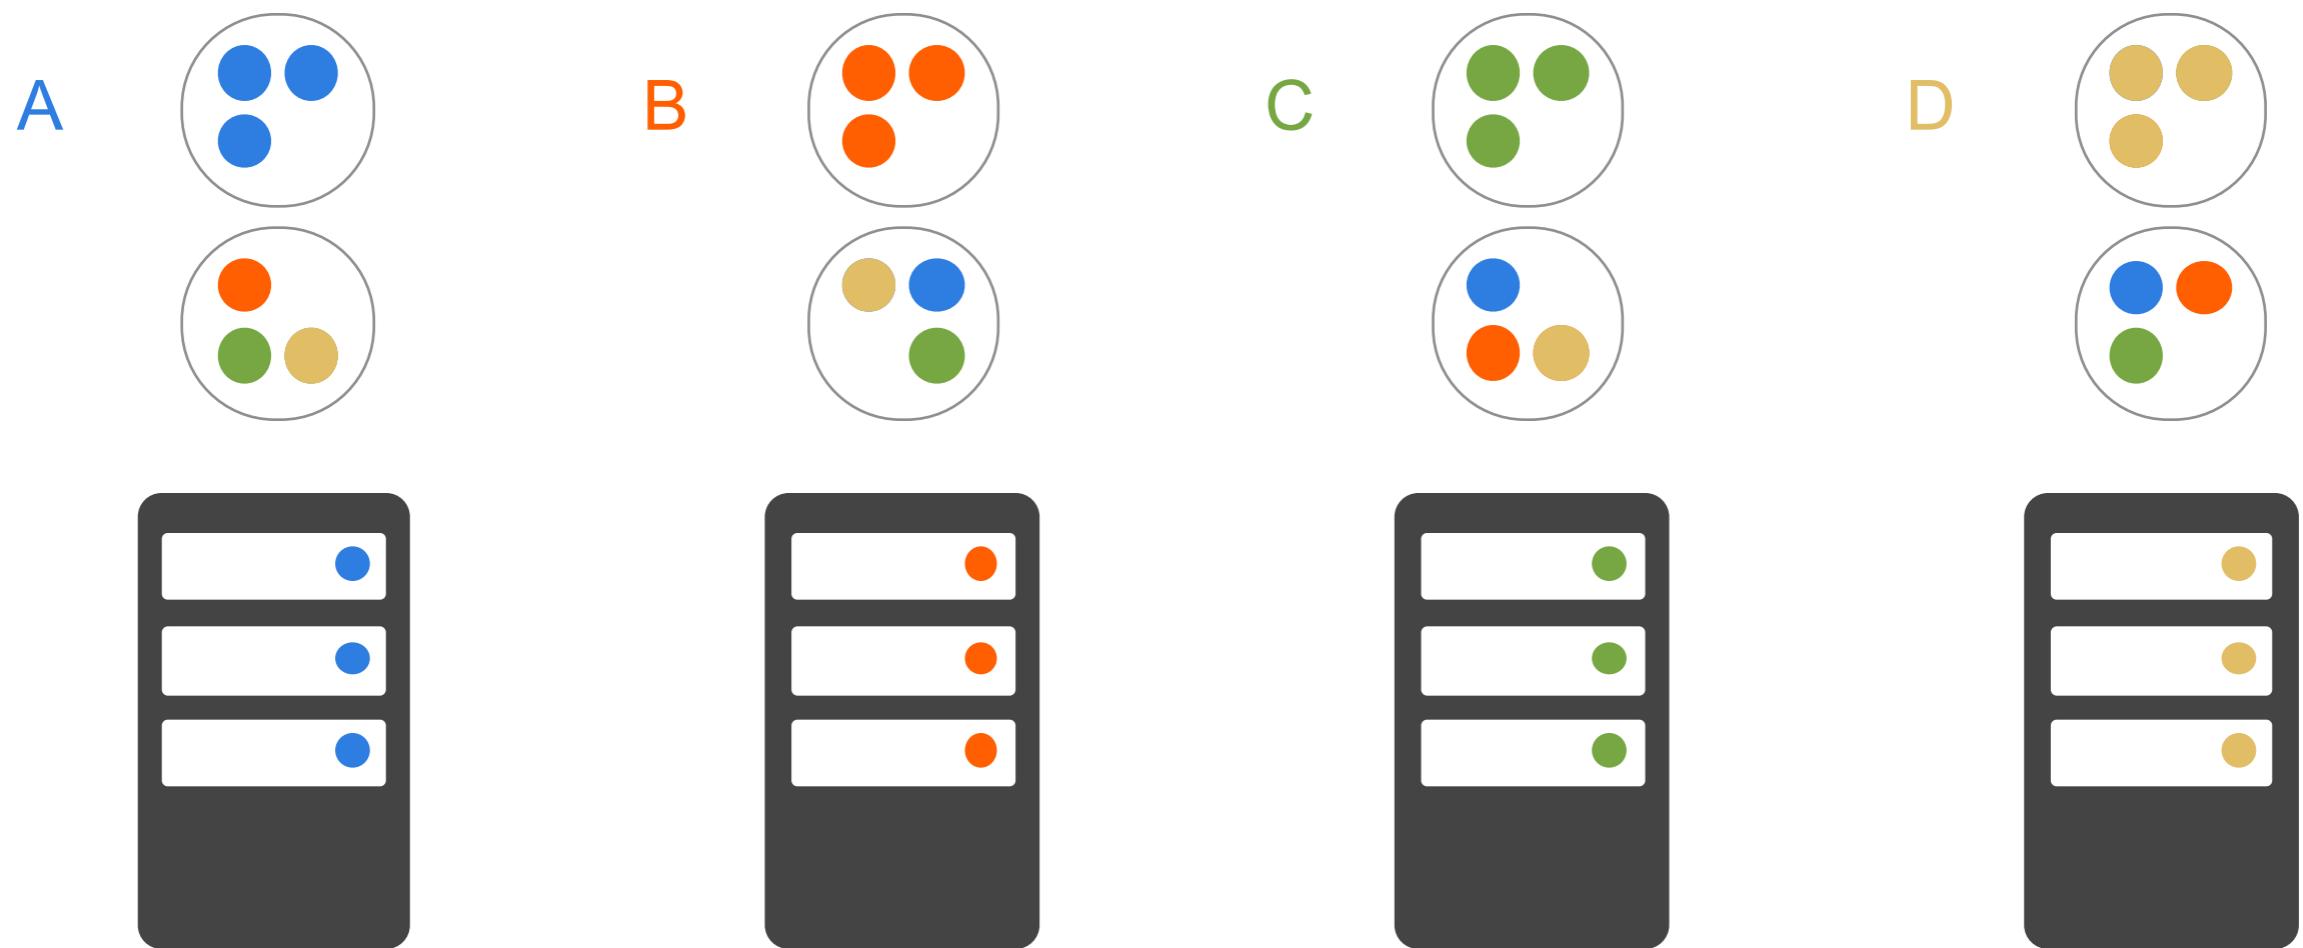


Migration





Migration Complete

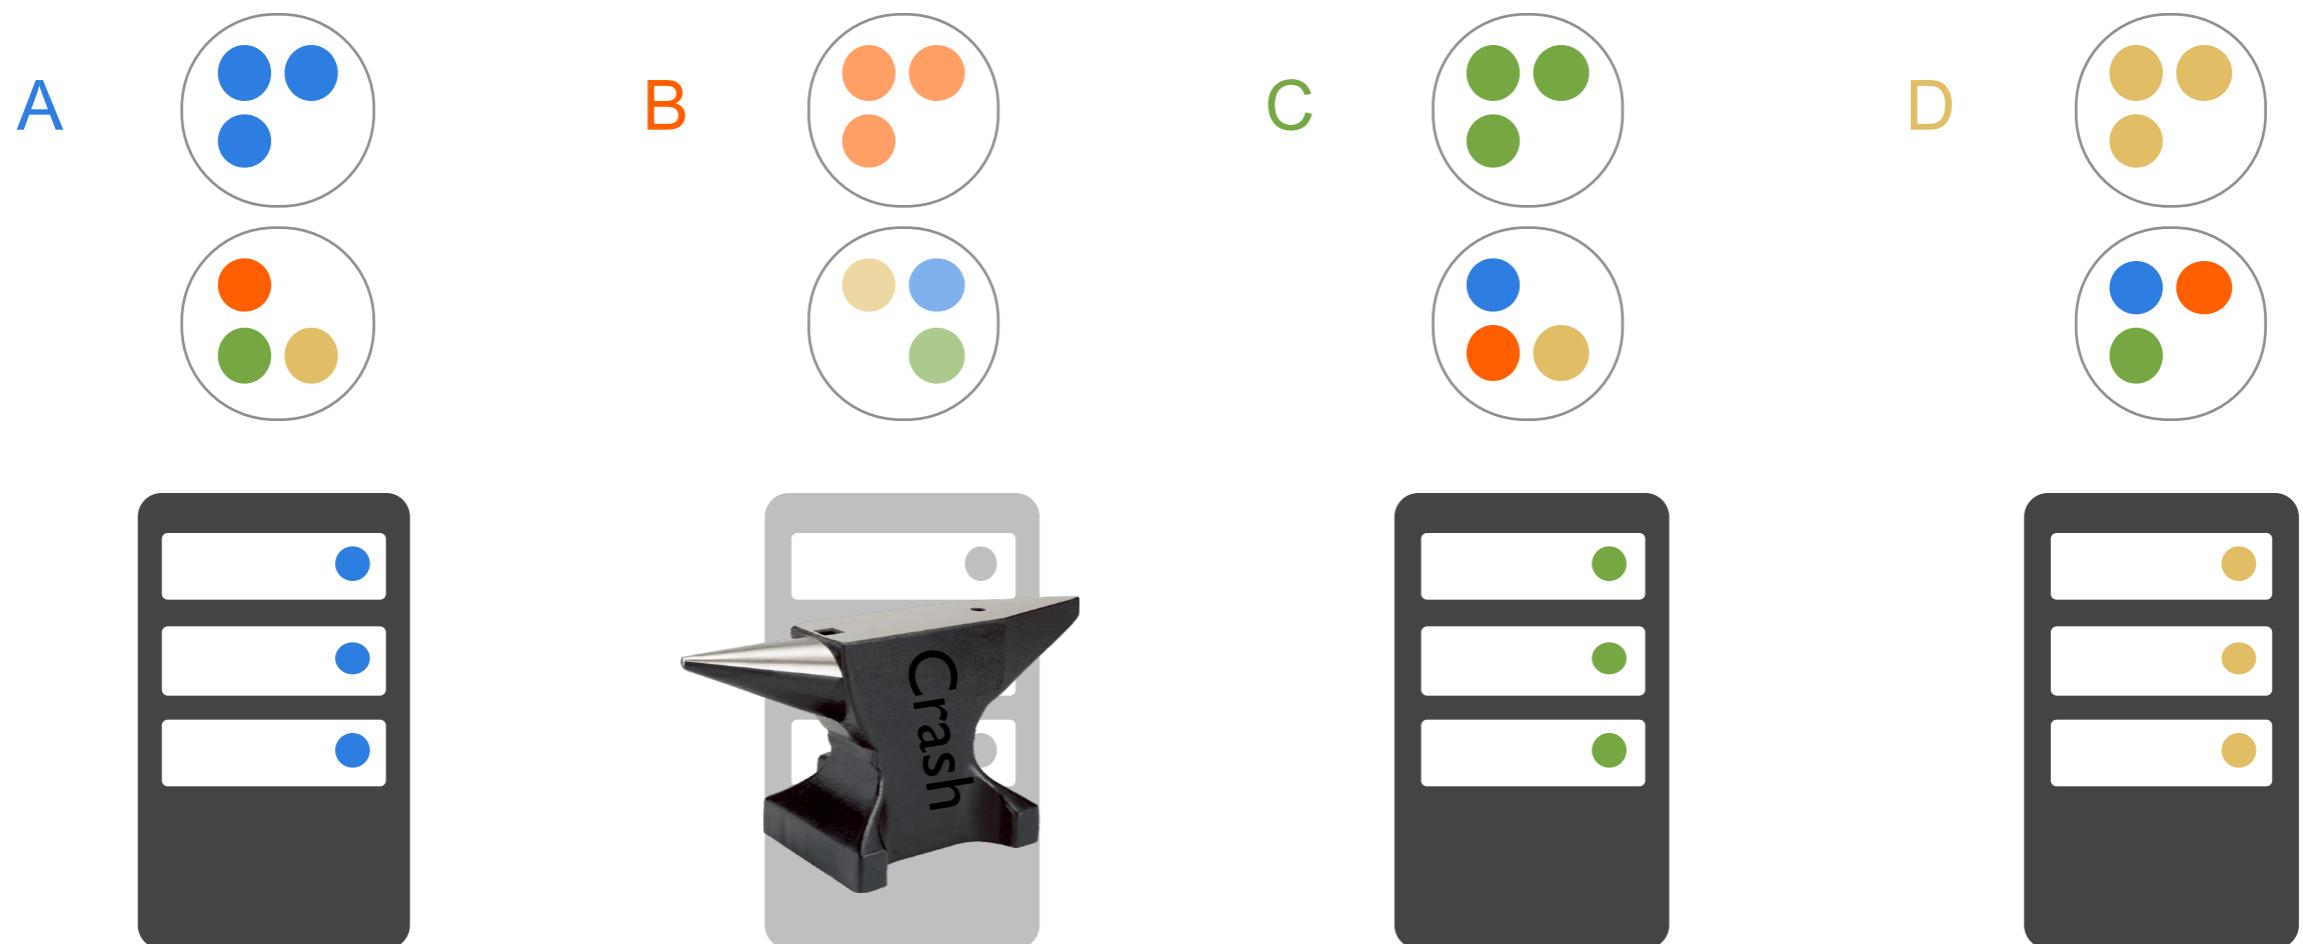




Data Safety when Node Dies

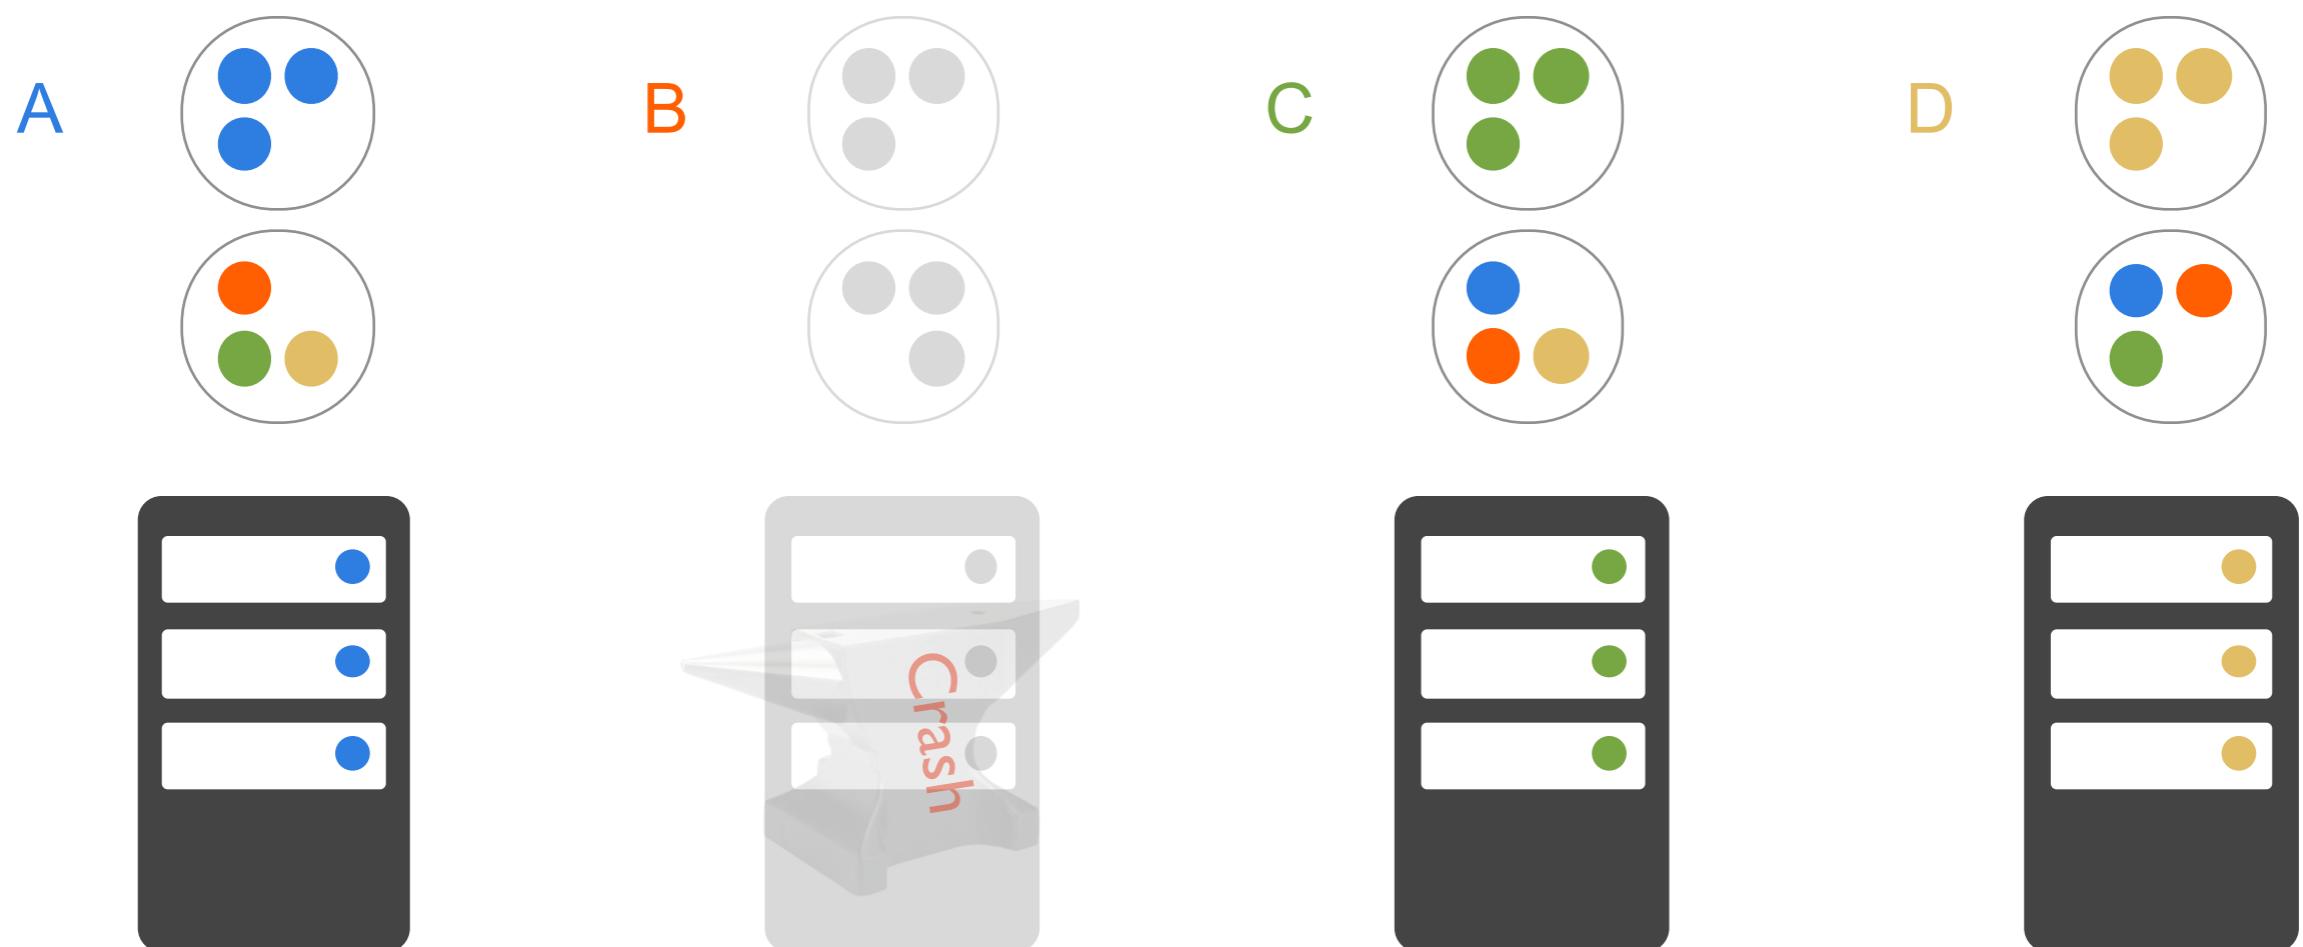


Node Crashes



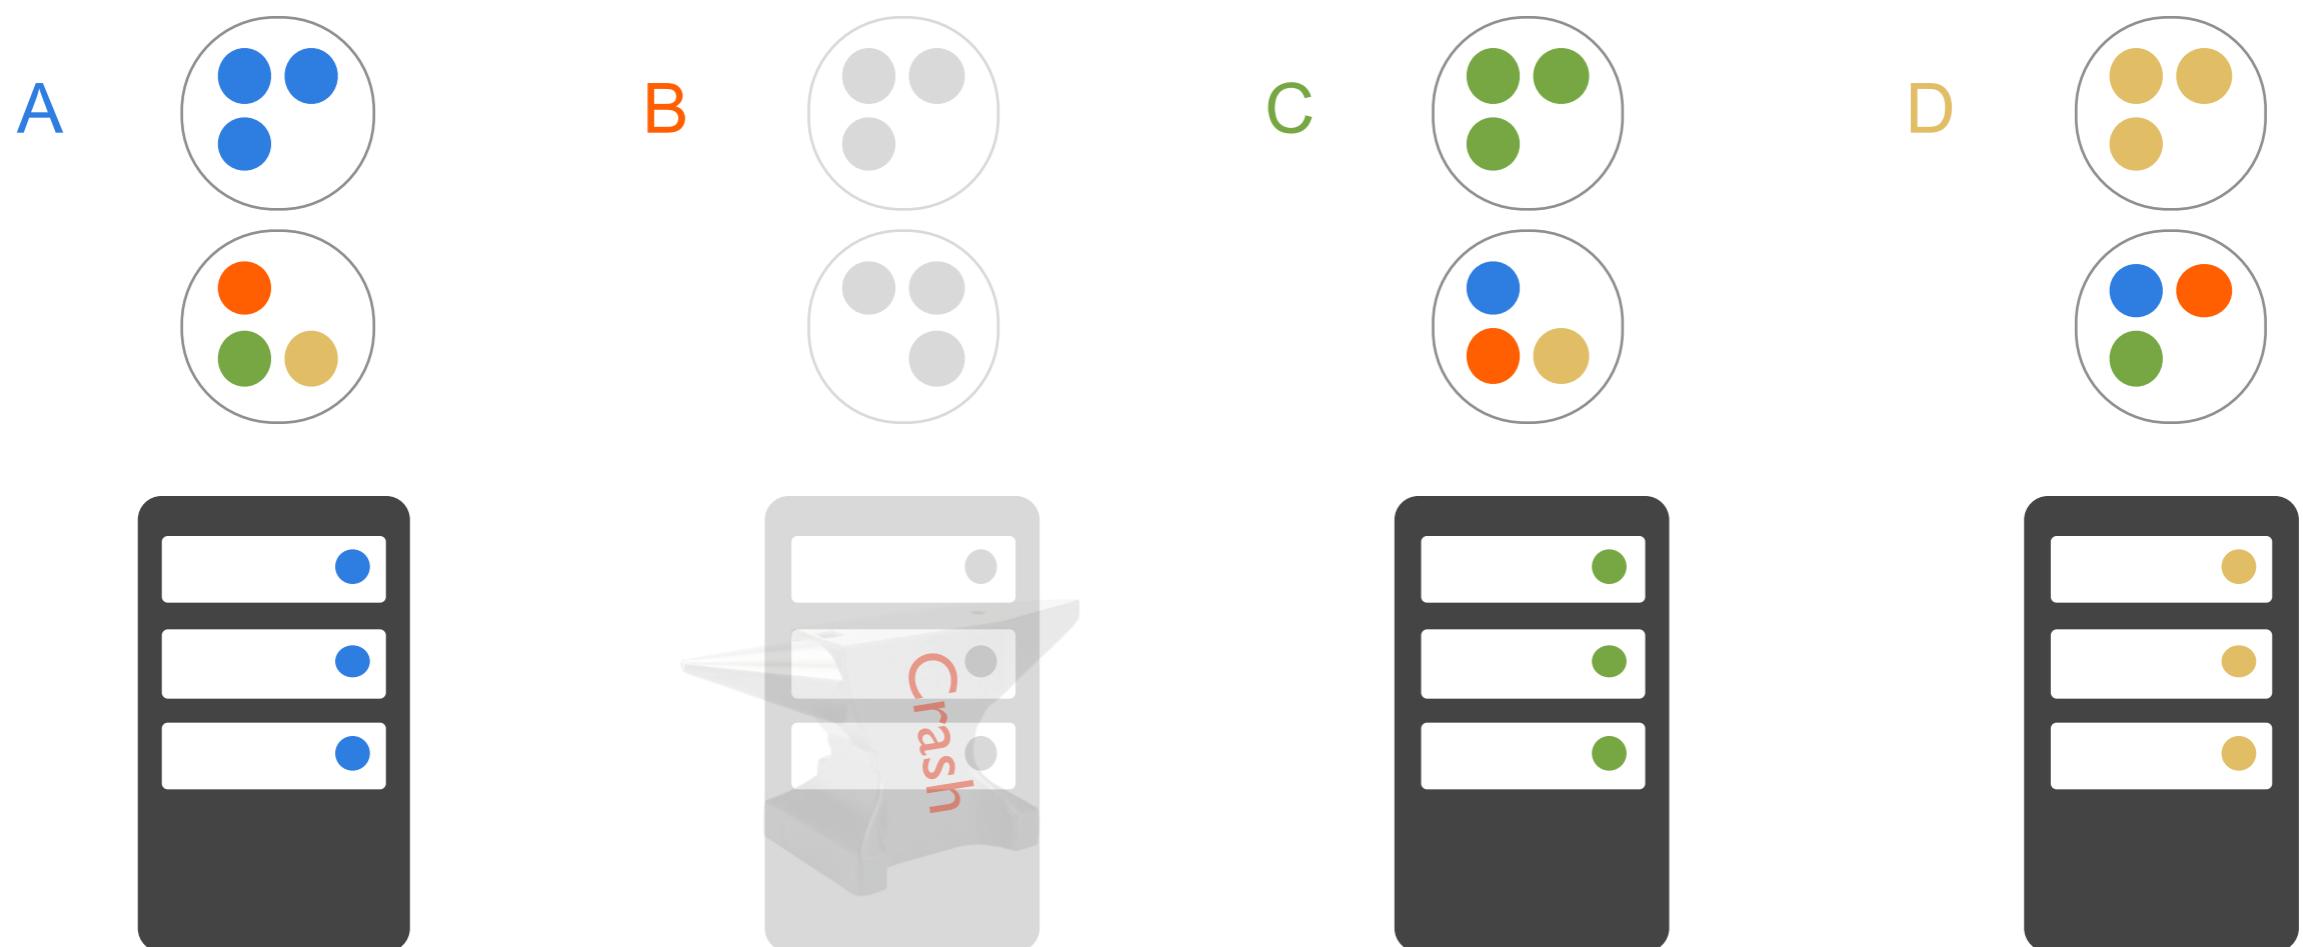


Backups Are Restored



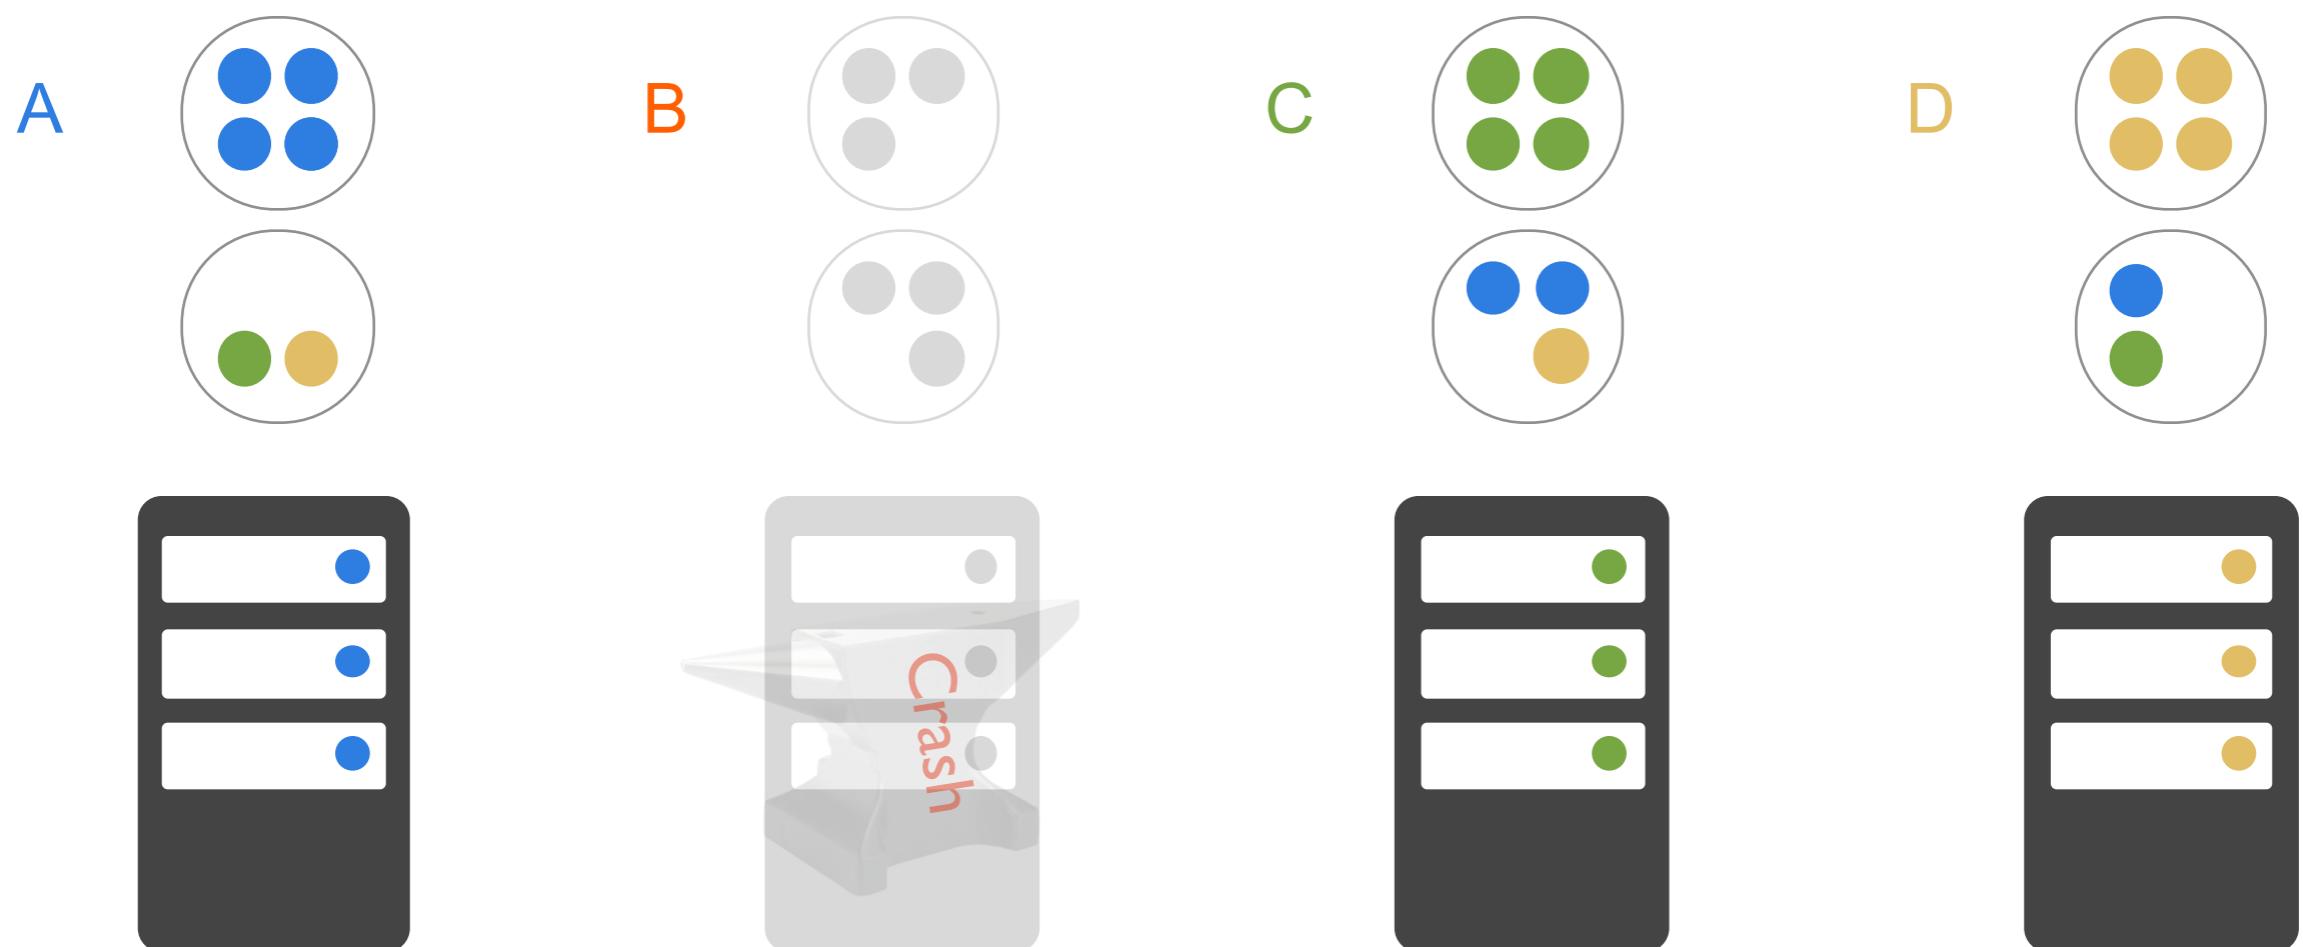


Backups Are Restored



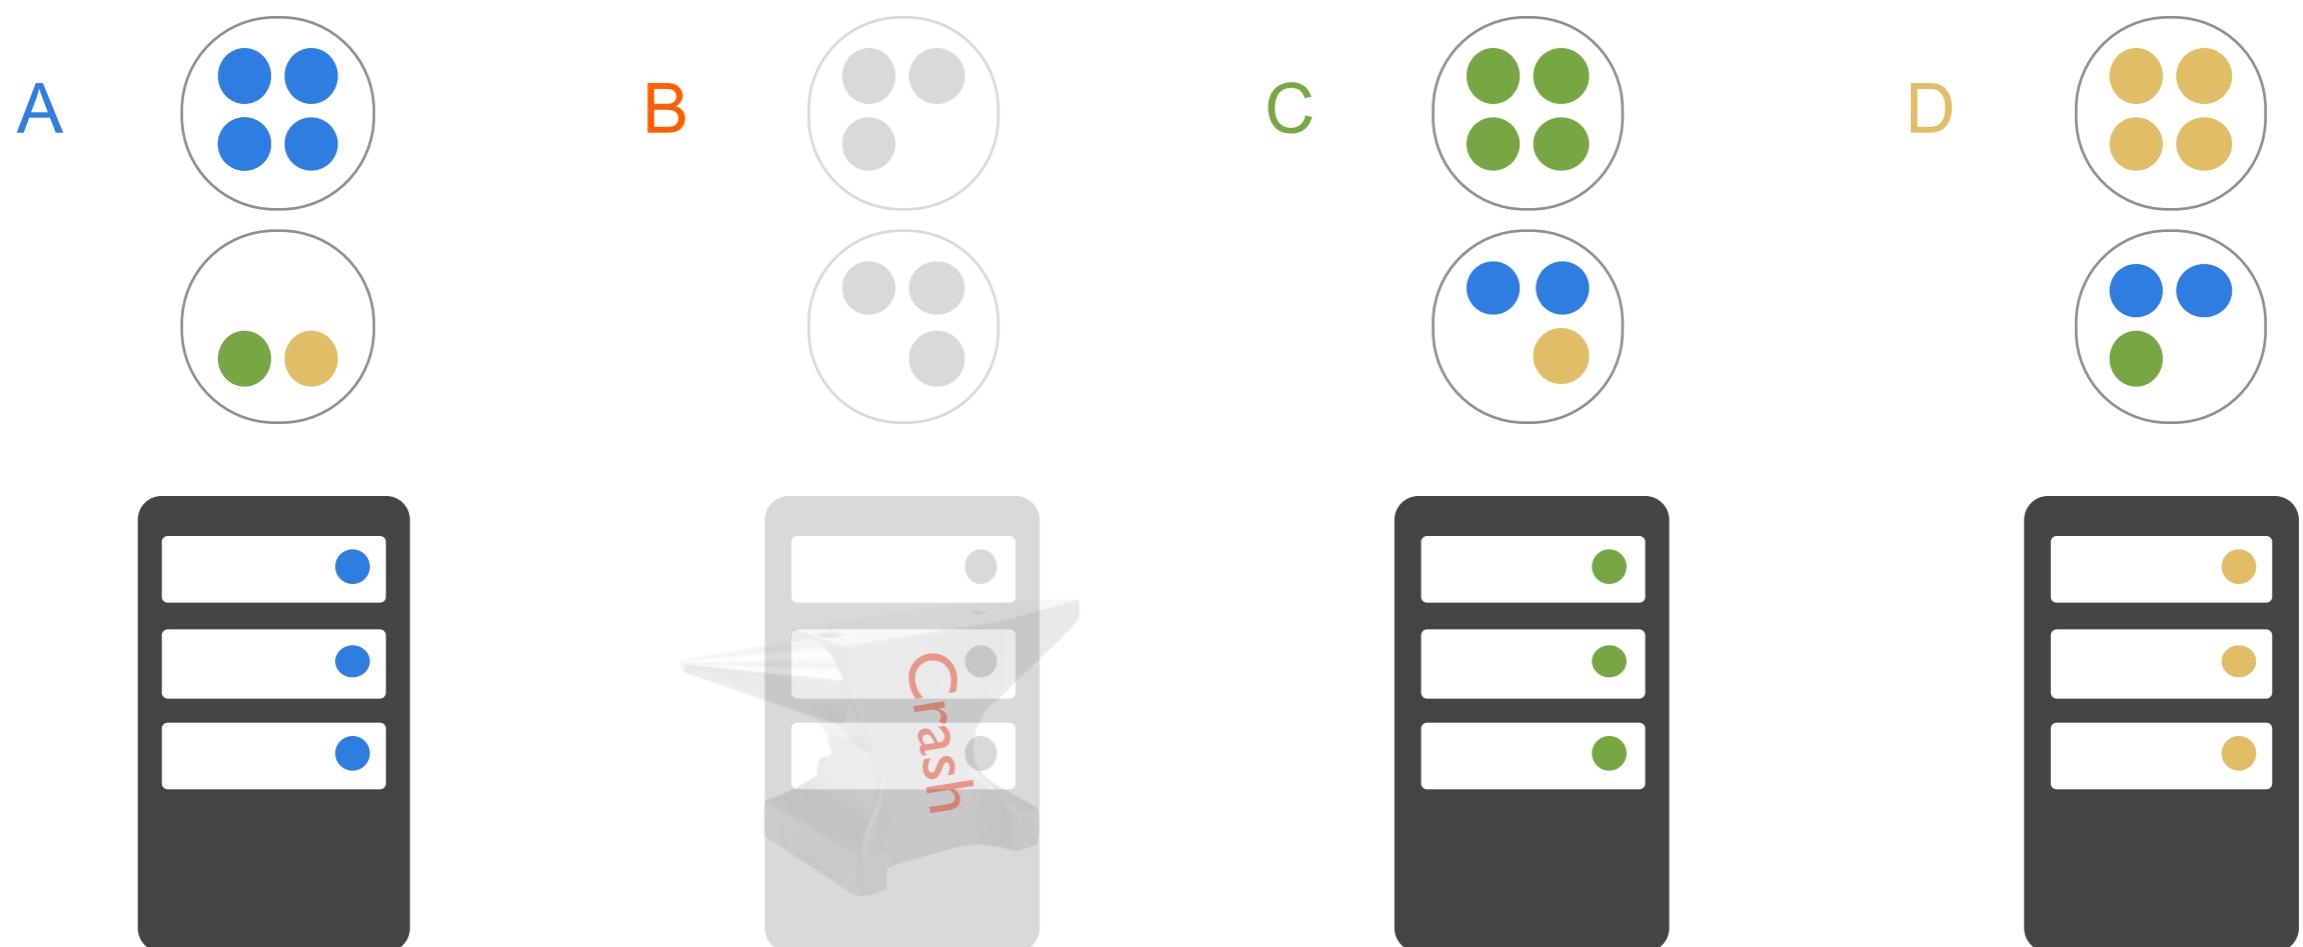


Backups Are Restored



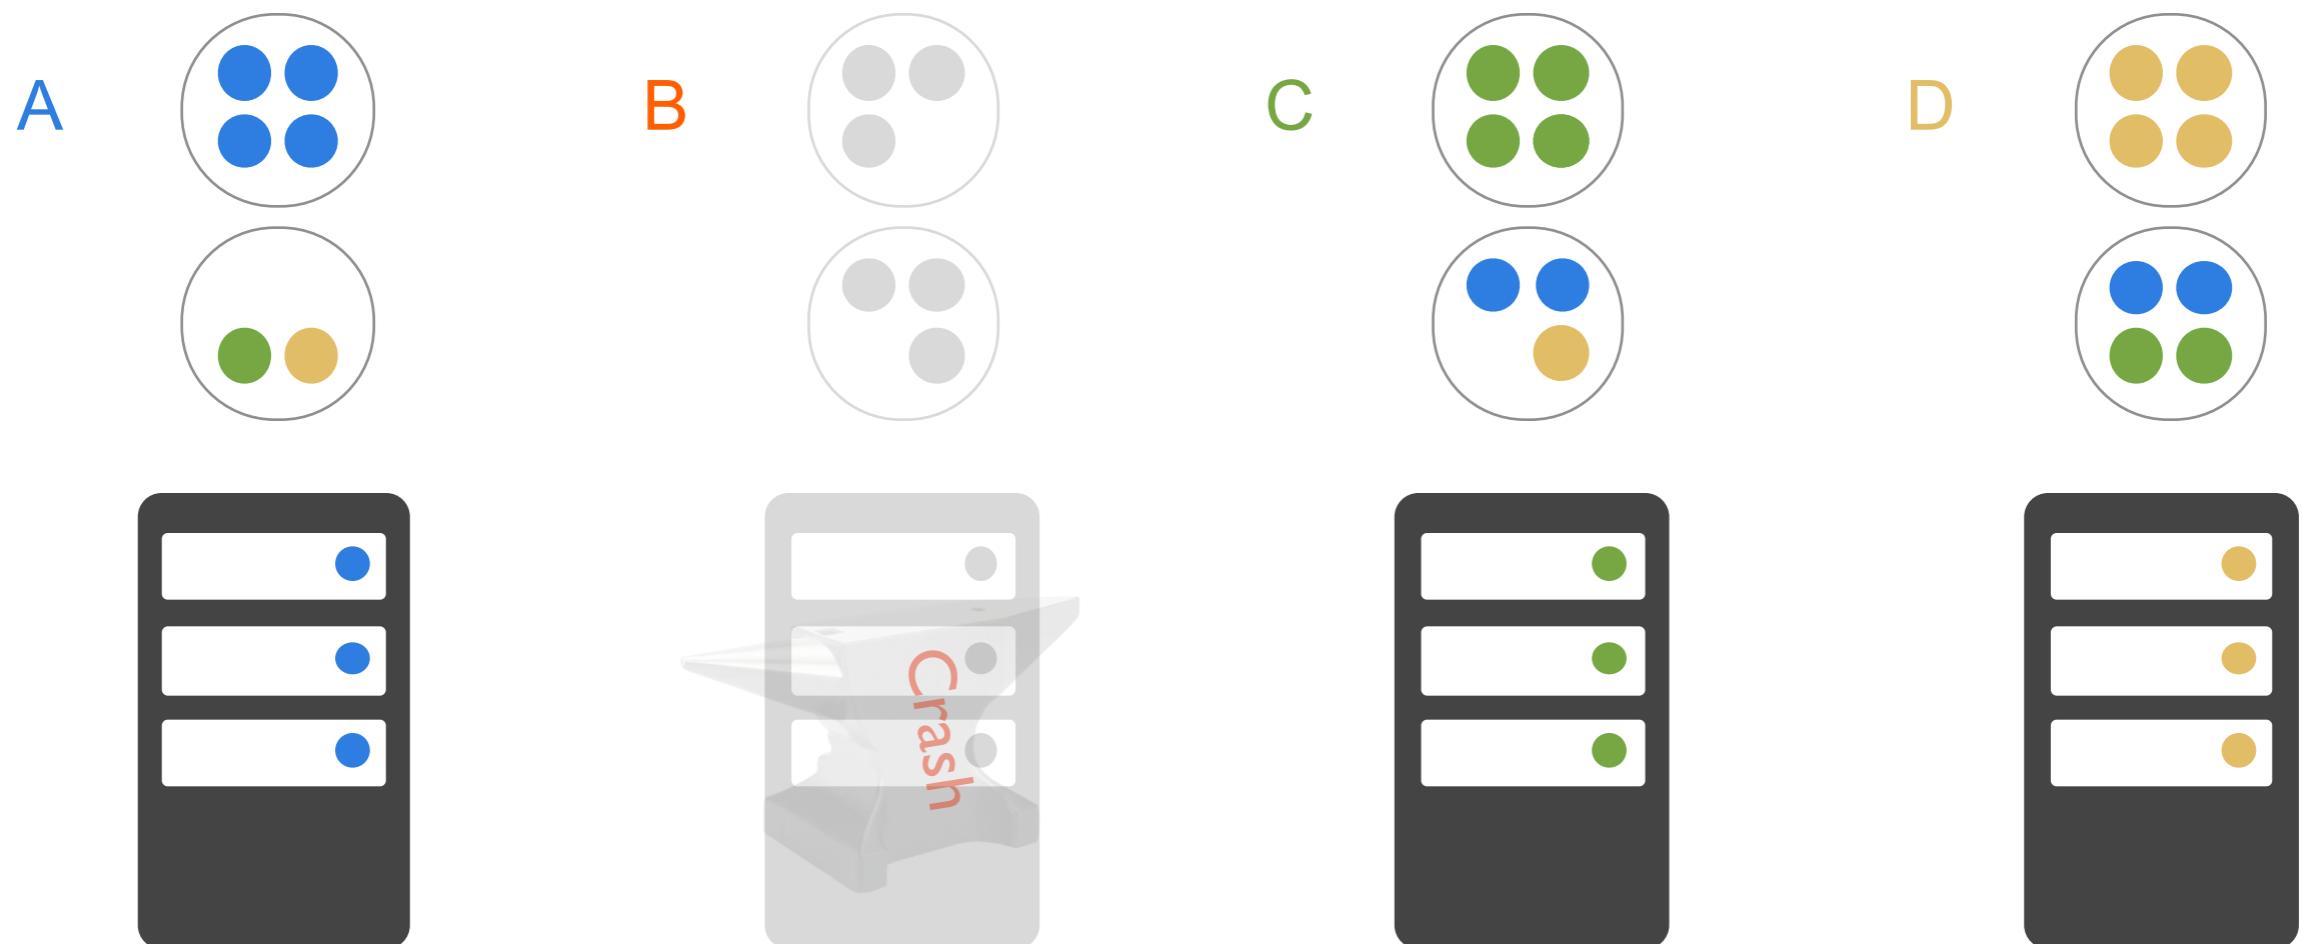


Backups Are Restored



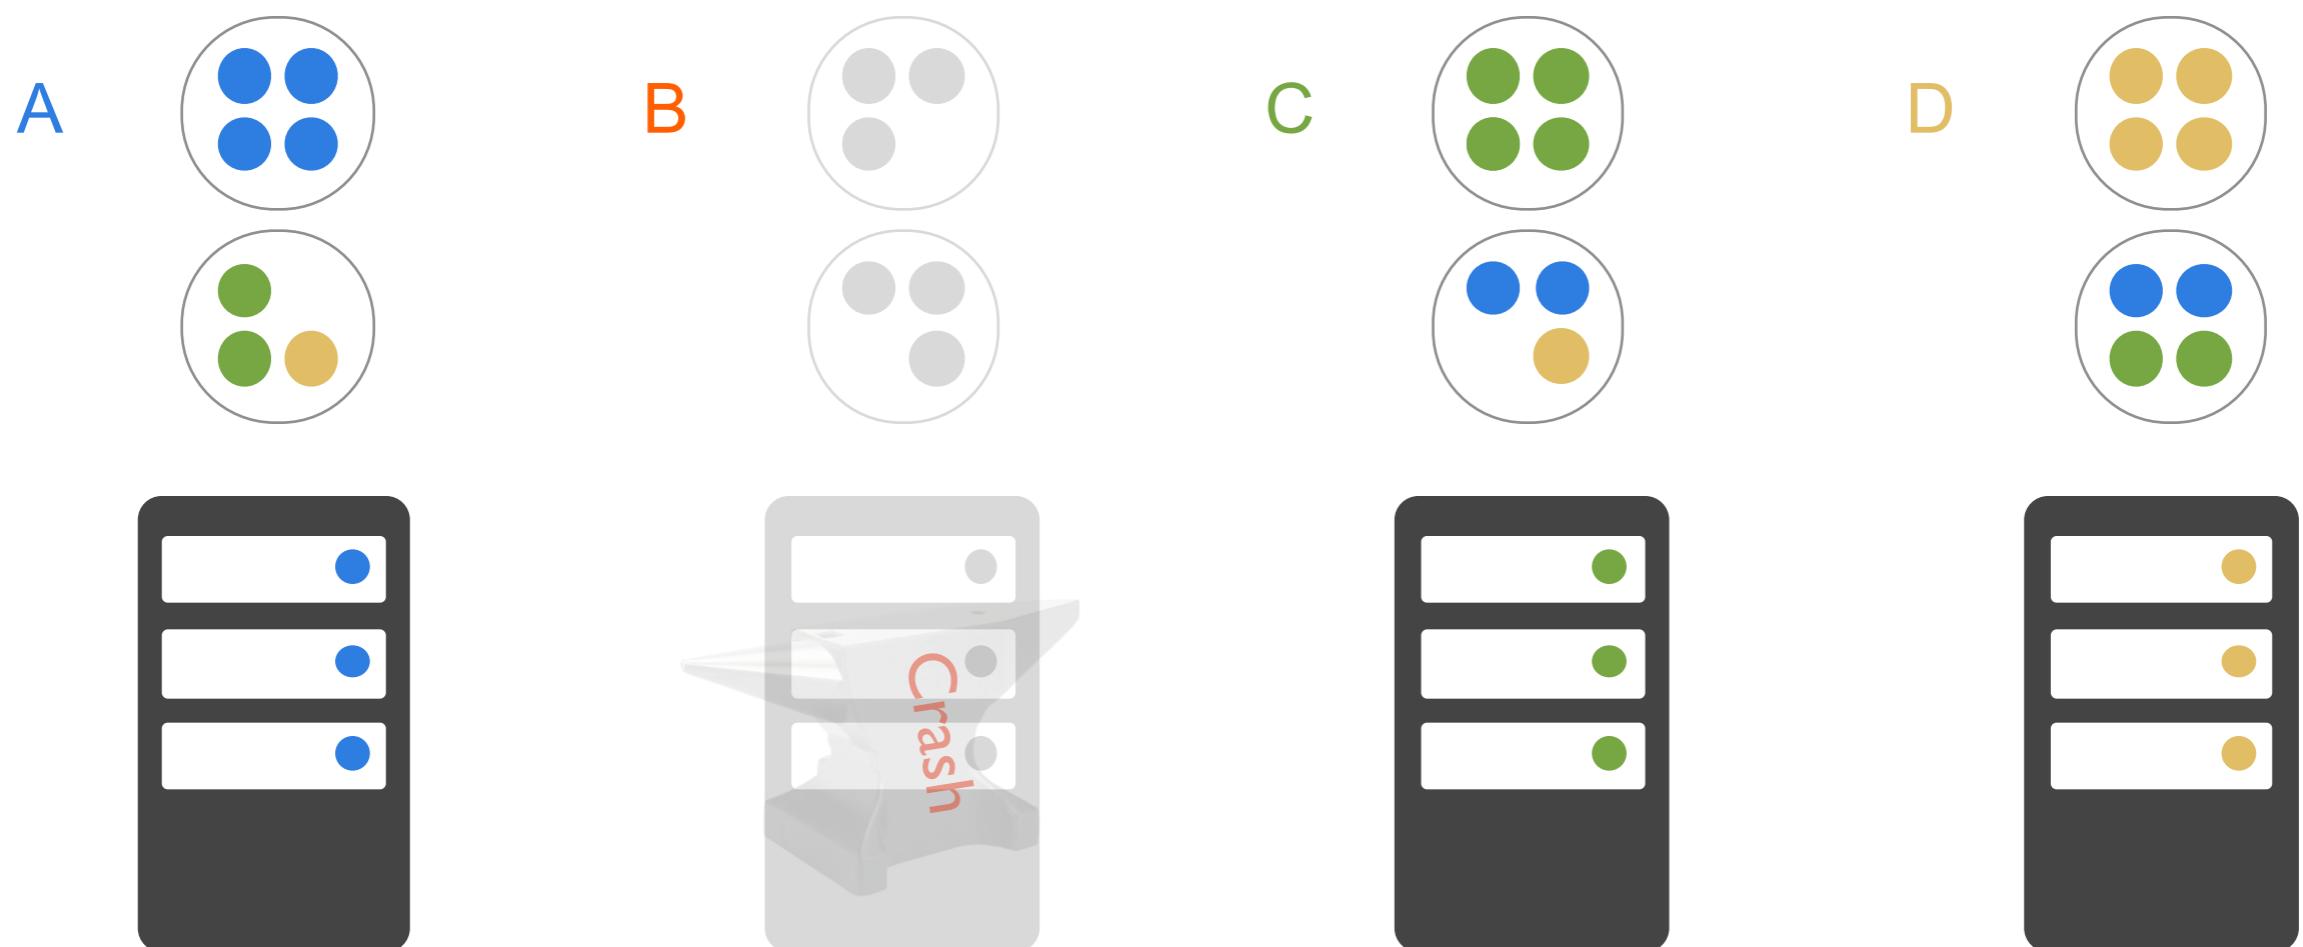


Backups Are Restored



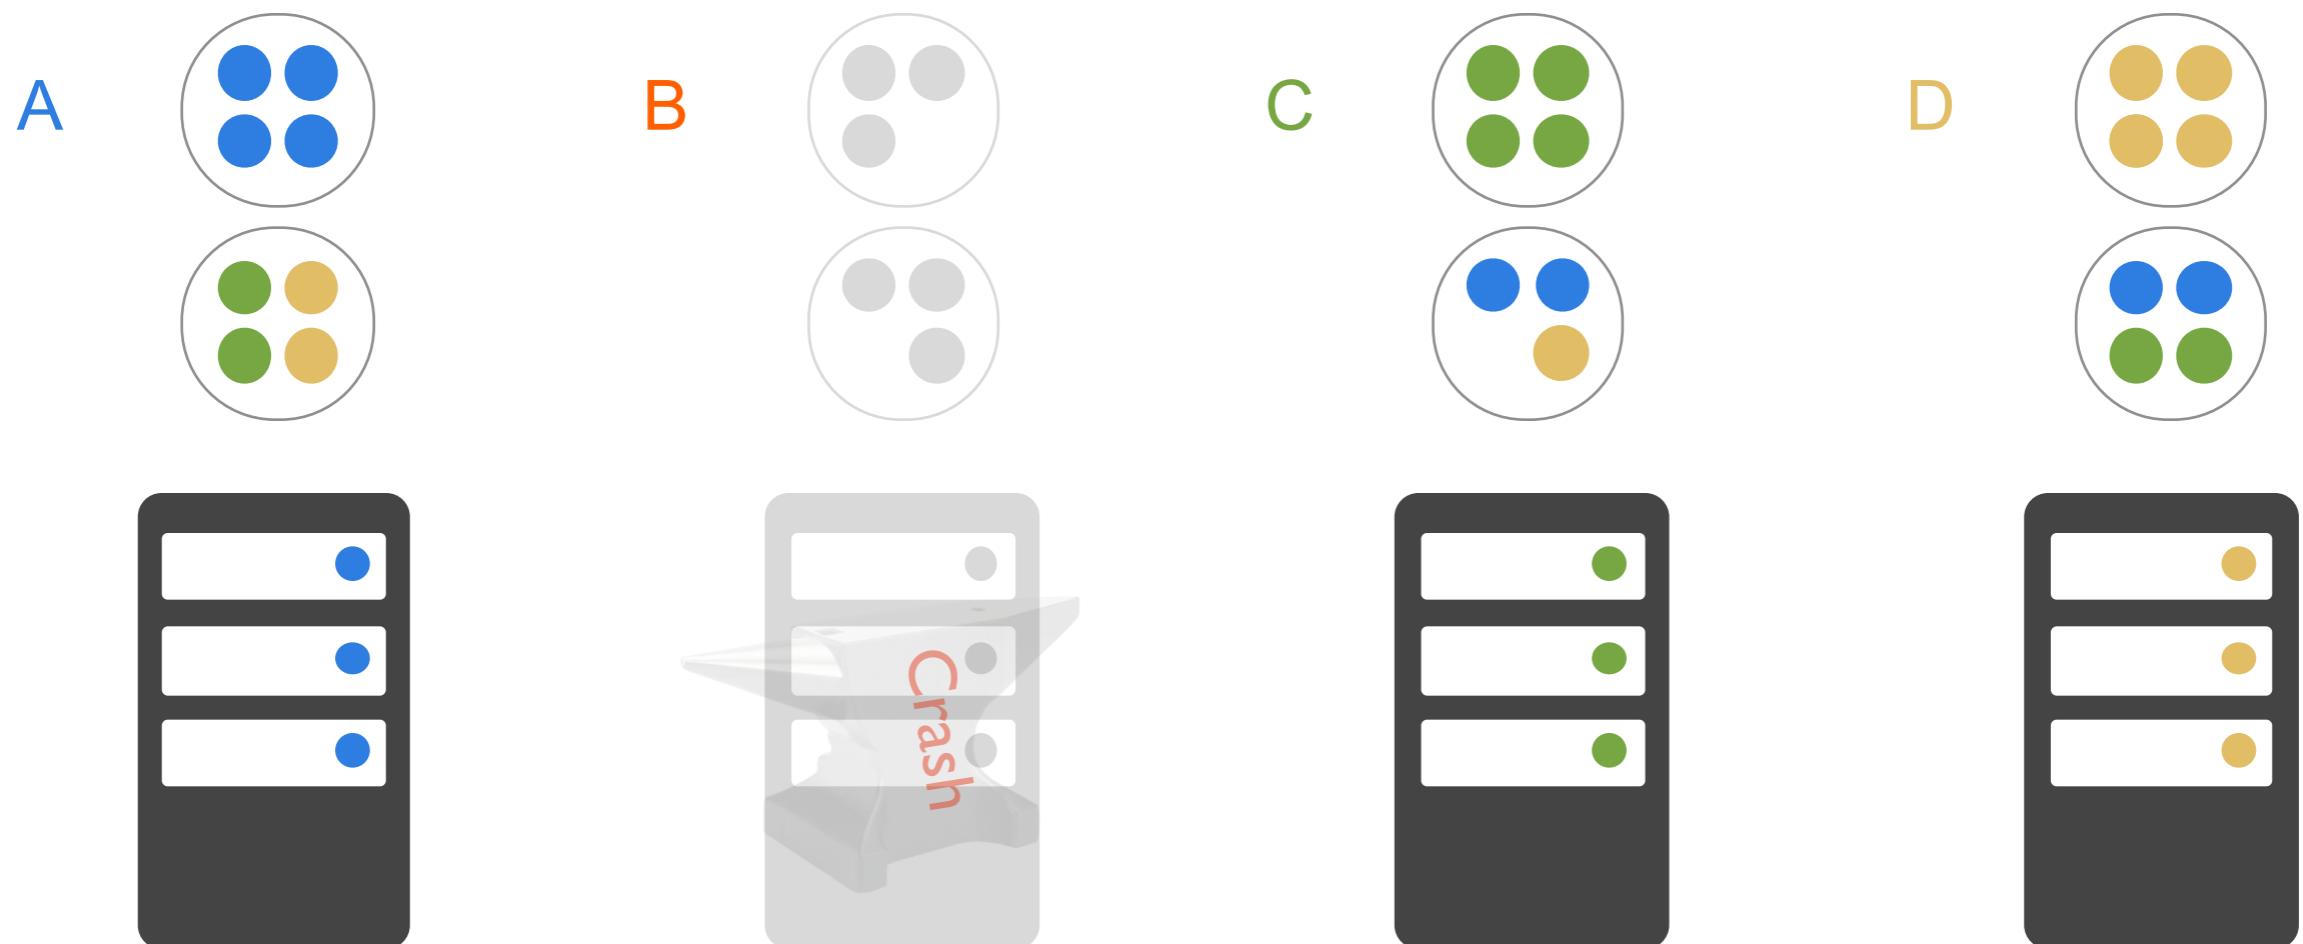


Backups Are Restored



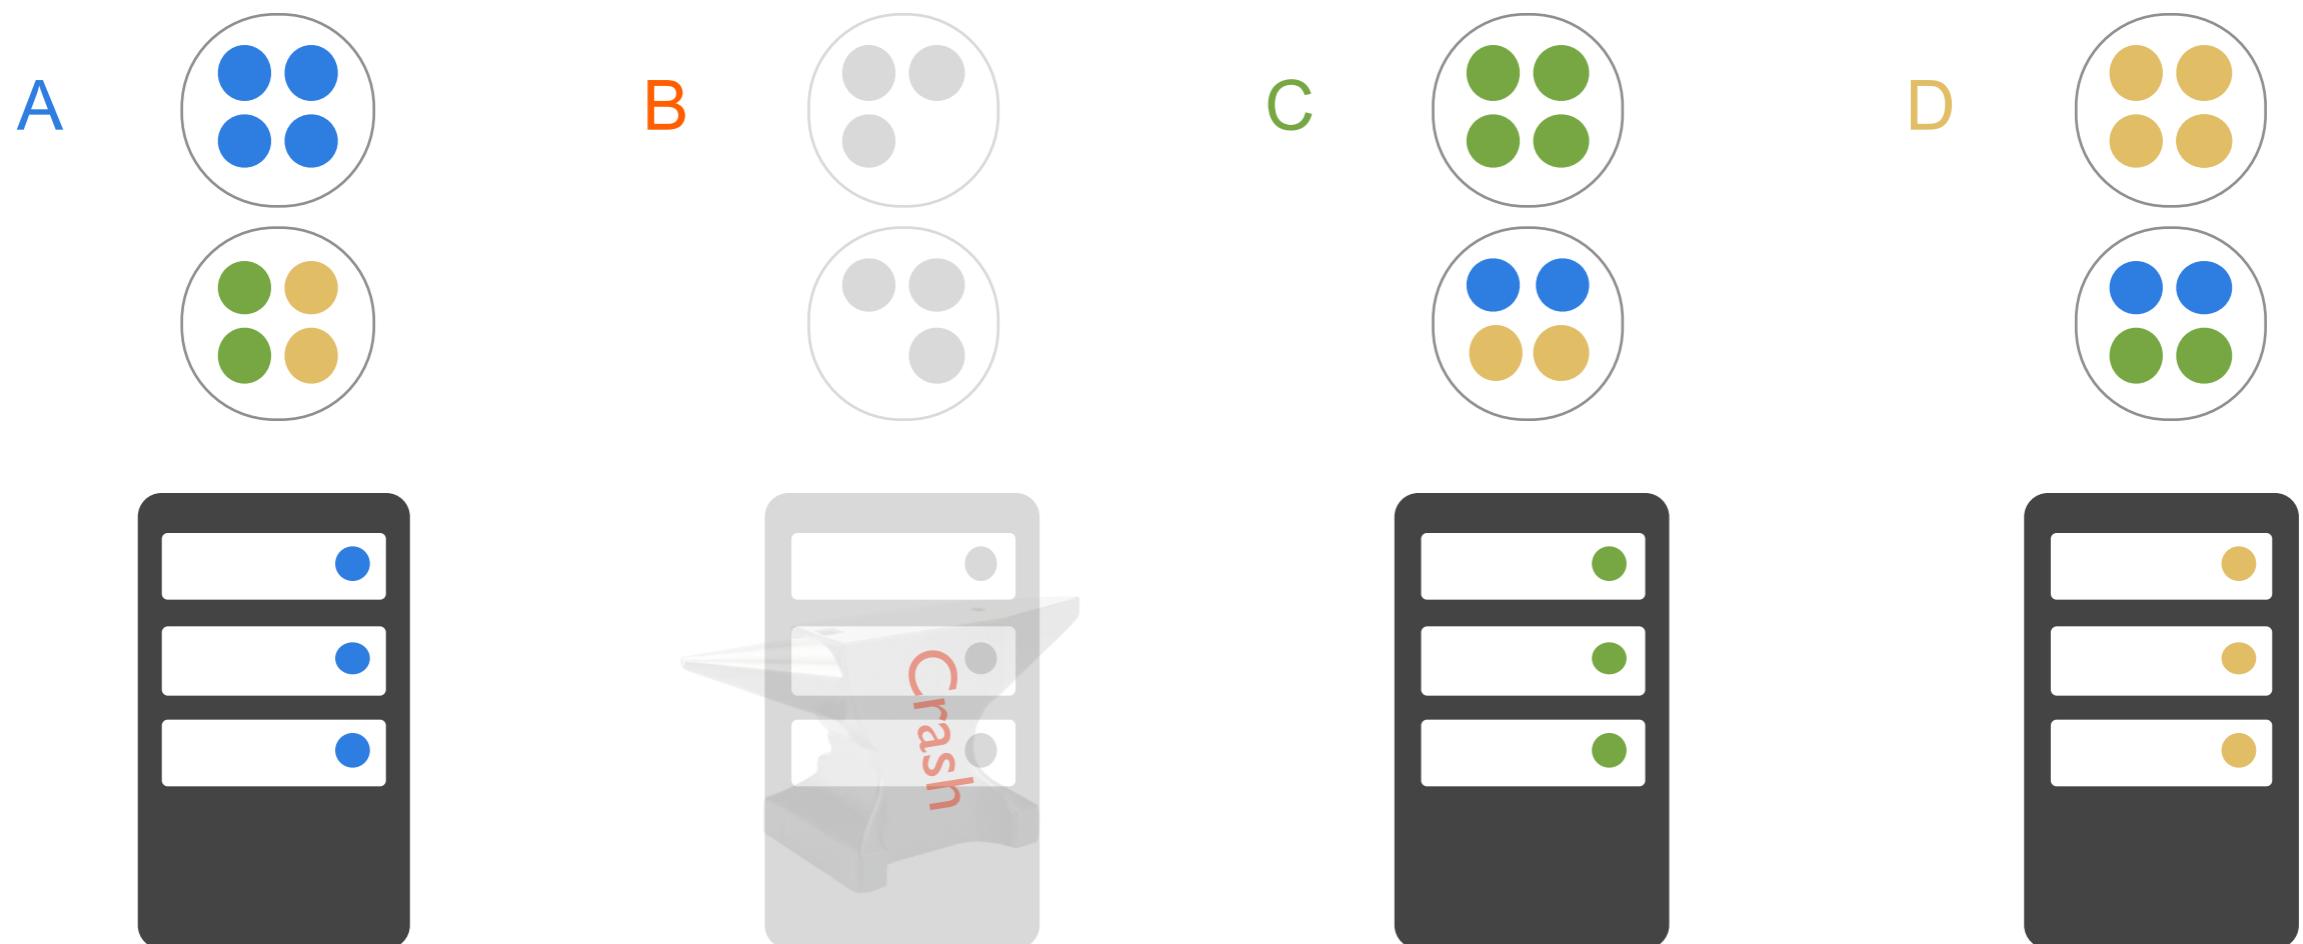


Backups Are Restored





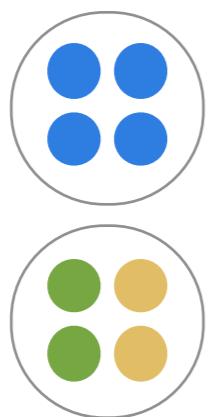
Backups Are Restored



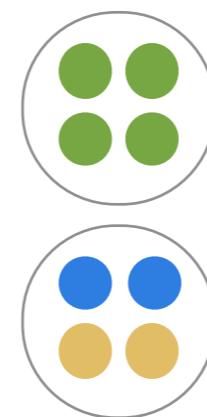


Recovery Is Complete

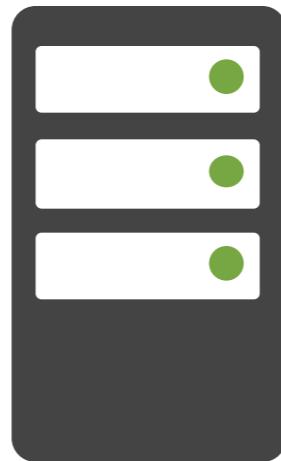
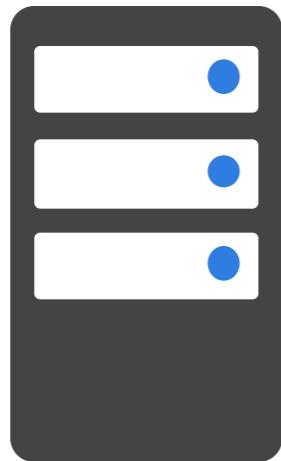
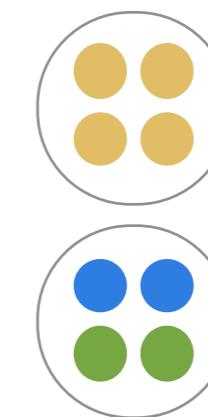
A



C



D





Scalability: Up or Out?



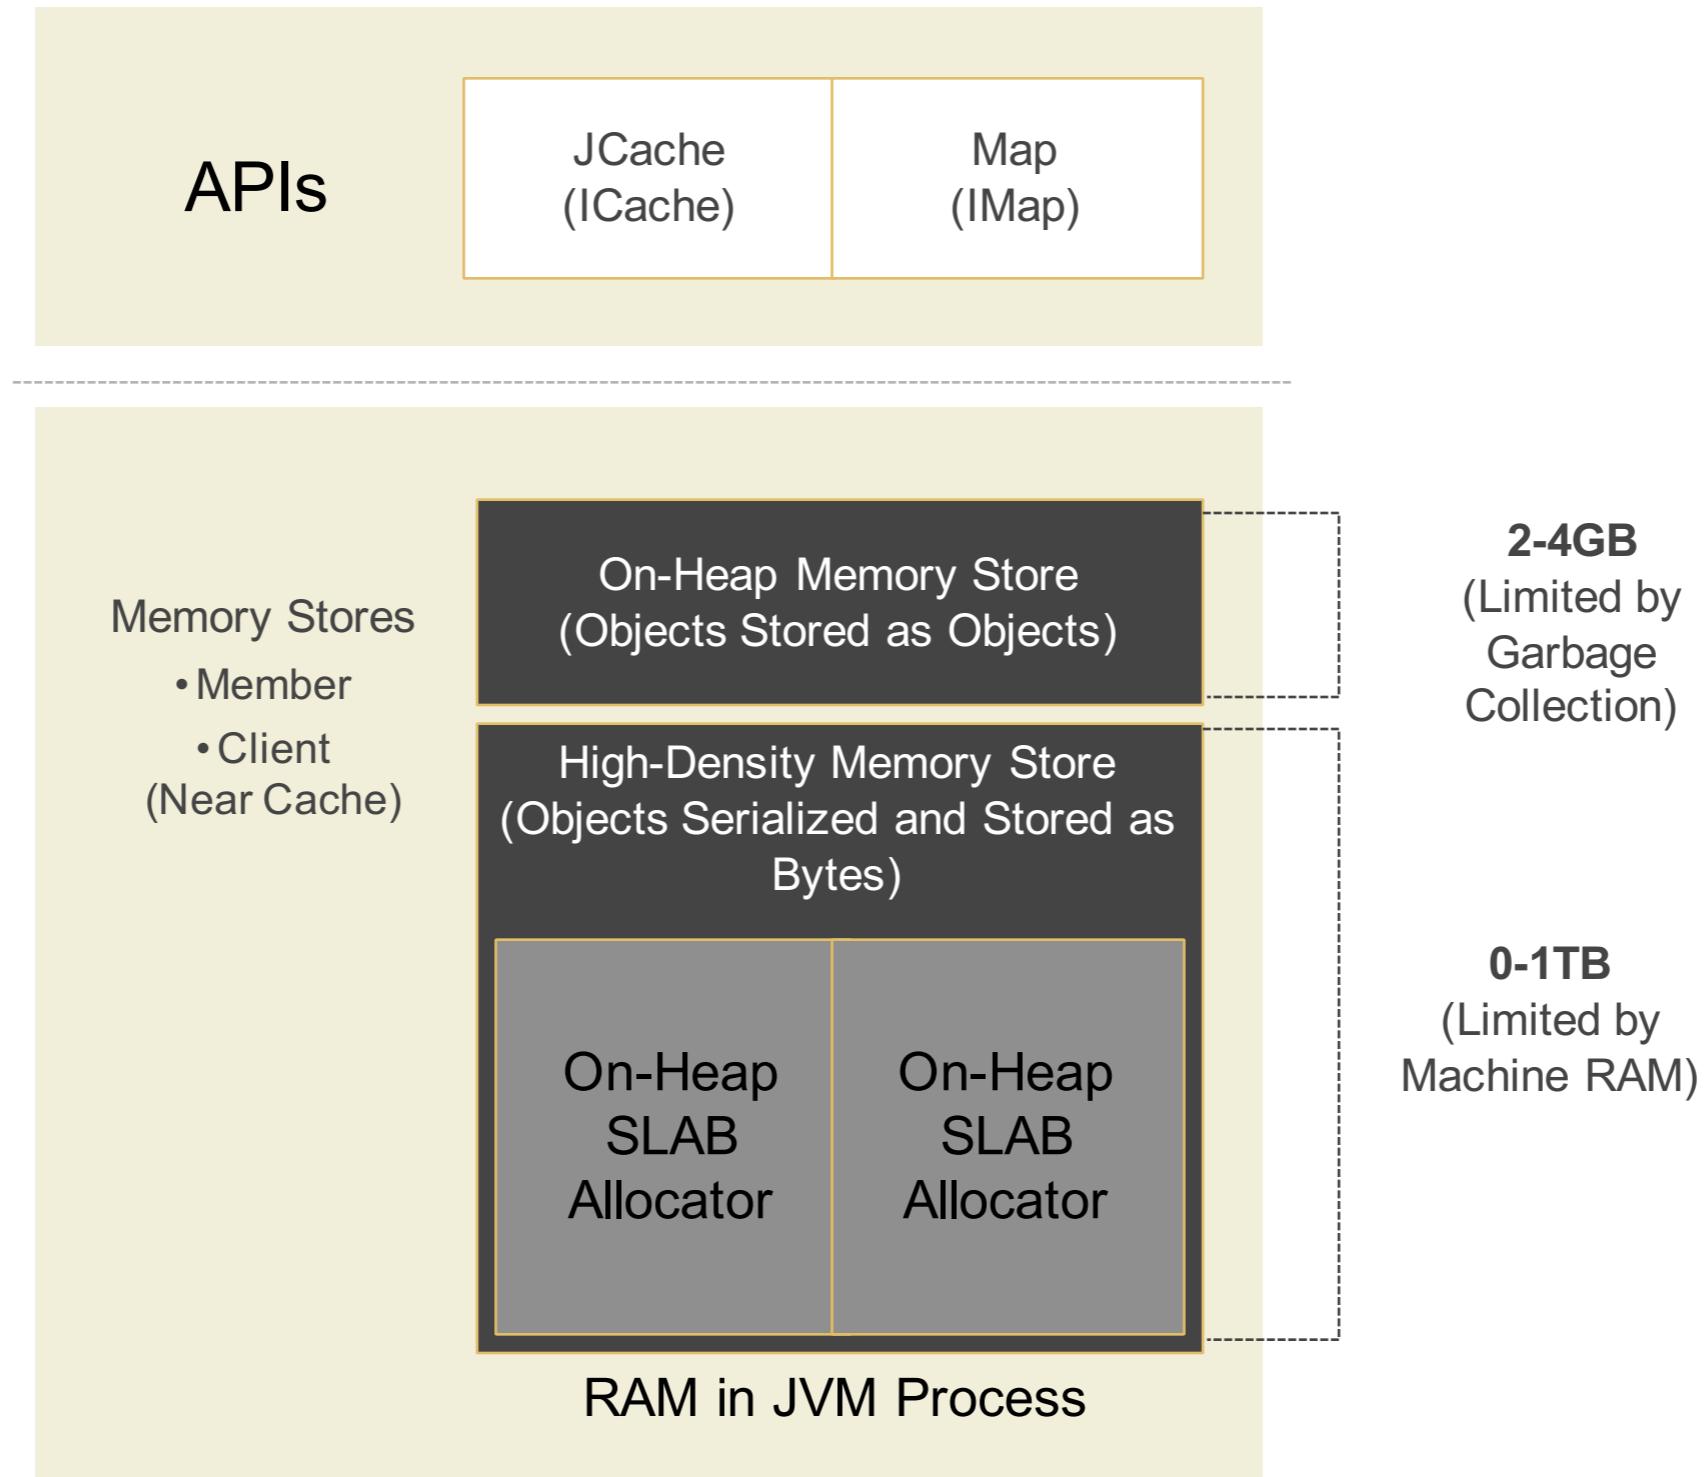




Durability

- Persist to or read from any external data store (SQL database, NoSQL data stores, REST service, EIS, file, etc.)
- Durability of application data (user sessions, shopping carts, master data, configs)
- Write through and write behind
- Hot Restart

High Density Caching



On Heap Vs. High-Density Memory Management

On Heap Memory

0 MB

Native

HD Memory v2

3.3 GB

3.9 GB

Heap Storage

0.6 GB

9 (4900 ms)

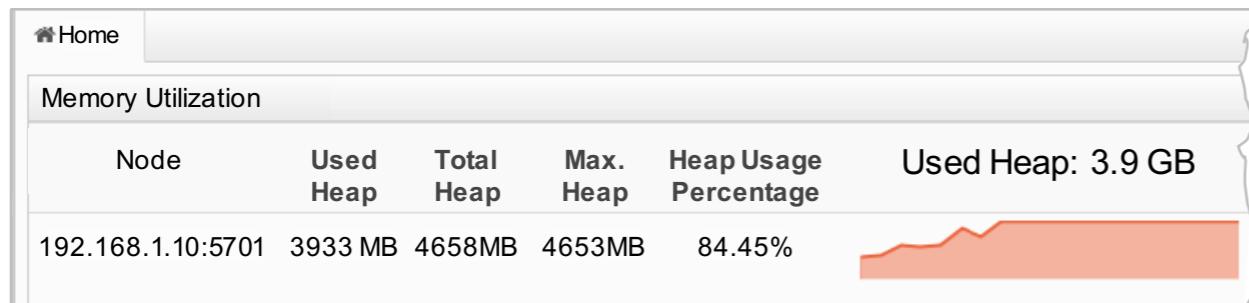
Major GC

0 (0 ms)

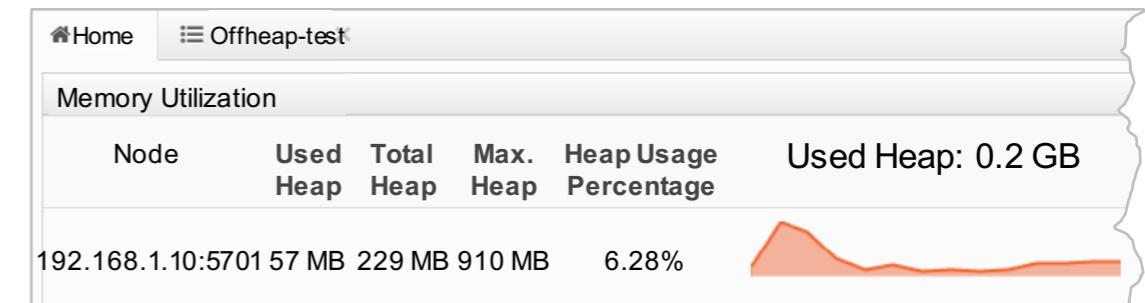
31 (4200 ms)

Minor GC

356 (349 ms)



Example: On Heap Memory



Example: HD Memory v2



Do I really need NoSQL?

NoSQL Landscape

Primary NoSQL technology*



* The results were normalized to exclude non-users

¹ Including Memcache, Riak and a dozen others

NOW YOU KNOW IMDG!





QA

Question, please?

