

Beautiful REST+JSON APIs

DevNexus 2016

Les Hazlewood @lhazlewood

Apache Shiro Project Chair

Expert Group Member, JEE Application Security

CTO, Stormpath, stormpath.com





- User Management API for Developers
- Password security
- Authentication and Authorization
- LDAP & Active Directory Cloud Sync
- Instant-on, scalable, and highly available
- Free for developers

Outline

- APIs, REST & JSON
- REST Fundamentals
- Design

Base URL

Versioning

Resource Format

Return Values

Content Negotiation

References (Linking)

Pagination

Query Parameters

Create/Update

Search

Associations

Errors

IDs

Method Overloading

Resource Expansion

Partial Responses

Caching & Etags

Security

Multi Tenancy

Maintenance

Batch Operations

APIs

- Applications
- Developers
- Pragmatism over Ideology
- Adoption
- Scale

Why REST?

- Scalability
- Generality
- Independence
- Latency (Caching)
- Security
- Encapsulation

Why JSON?

- Ubiquity
- Simplicity
- Readability
- Scalability
- Flexibility

HATEOAS

- **H**ypermedia
- **A**s
- **T**he
- **E**ngine
- **O**f
- **A**pplication
- **S**tate

REST Is Easy

REST Is *&@# \$! Hard

(for providers)

REST *can* be easy

(if you follow some guidelines)

Example Domain: Stormpath

- Applications
- Directories
- Accounts
- Groups
- Associations
- Workflows



Fundamentals

Resources

Nouns, not Verbs

Coarse Grained, not Fine Grained

Architectural style for use-case scalability

What If?

/getAccount

/createDirectory

/updateGroup

/verifyAccountEmailAddress

What If?

/getAccount

/getAllAccounts

/searchAccounts

/createDirectory

/createLdapDirectory

/updateGroup

/updateGroupName

/findGroupsByDirectory

/searchGroupsByName

/verifyAccountEmailAddress

/verifyAccountEmailAddressByToken

...

Smells like bad RPC. DON'T DO THIS.

Keep It Simple

The Answer

Fundamentally two types of resources:

Collection Resource

Instance Resource

Collection Resource

/applications

Instance Resource

/applications/a1b2c3

Behavior

- GET
- PUT
- POST
- DELETE
- HEAD

Behavior

POST, GET, PUT, DELETE

$\neq 1:1$

Create, Read, Update, Delete

Behavior

As you would expect:

GET = Read

DELETE = Delete

HEAD = Headers, no Body

Behavior

Not so obvious:

PUT and POST can *both* be used for
Create *and* Update

PUT for Create

Identifier is known by the client:

```
PUT /applications/clientSpecifiedId
```

```
{  
  ...  
}
```


PUT for Update

Full Replacement

```
PUT /applications/existingId
{
  "name": "Best App Ever",
  "description": "Awesomeness"
}
```

PUT

Idempotent

POST as Create

On a parent resource

```
POST /applications
```

```
{  
  "name": "Best App Ever"  
}
```

Response:

```
201 Created
```

```
Location: https://api.stormpath.com/applications/a1b2c3
```

POST as Update

On instance resource

```
POST /applications/a1b2c3
```

```
{  
  "name": "Best App Ever. Srsly."  
}
```

Response:

200 OK

POST

NOT Idempotent

Media Types

- Format Specification + Parsing Rules
- Request: Accept header
- Response: Content-Type header
- `application/json`
- `application/foo+json`
- `application/ion+json; v=1`
- ...

Design Time!

Base URL

http(s)://foo.io

VS

http://www.foo.com/dev/service/api/rest

http(s)://foo.io

Rest Client
vs
Browser

Versioning

URL

`https://api.stormpath.com/v1`

vs.

Media-Type

`application/json`

`application/foo+json;v=1`

Resource Format

Media Type

Content-Type: application/json

When time allows:

application/foo+json

application/foo2+json;bar=baz

...

Media Type

Don't go overboard!

Media Type \neq Schema!

Most only need 1 *maybe* 2 media types:

- instance resource
- collection resource

`application/ion+json`

`application/ion+json;bar=baz`

...

camelCase

‘JS’ in ‘JSON’ = JavaScript

`myArray.forEach`

Not `myArray.for_each`

`account.givenName`

Not `account.given_name`

Underscores for property/function names are unconventional for JS. Stay consistent.

Date/Time/Timestamp

There's already a standard. Use it: ISO 8601

Example:

```
{  
  ...,  
  "createdAt": "2013-07-10T18:02:24.343Z",  
  ...  
}
```

Use UTC!

createdAt / updatedAt

createdAt / updatedAt

Most people will want this at some point

```
{  
  ... ,  
  "createdAt": "2013-07-10T18:02:24.343Z",  
  "updatedAt": "2014-09-29T07:02:48.761Z"  
}
```

Use UTC!

Response Body

GET obvious

What about POST?

Return the representation in the response when feasible.

Add override (`?_body=false`) for control

Content Negotiation

Header

- Accept header
- Header values comma delimited
- q param determines precedence, defaults to 1, then conventionally by list order

```
GET /applications/a1b2c3
```

```
Accept: application/json,  
text/plain;q=0.8
```

Resource Extension

`/applications/a1b2c3.json`

`/applications/a1b2c3.csv`

...

Conventionally overrides `Accept` header

HREF

- Distributed Hypermedia is paramount!
- **Every accessible Resource has a canonical unique URL**
- Replaces IDs (IDs exist, but are opaque).
- Critical for linking, as we'll soon see

Instance w/ HREF (v1)

GET /accounts/x7y8z9

200 OK

```
{  
  "href": "https://api.stormpath.com/v1/accounts/x7y8z9",  
  "givenName": "Tony",  
  "surname": "Stark",  
  ...  
}
```

Resource References aka 'Linking' (v1)

- Hypermedia is paramount.
 - Linking is fundamental to scalability.
-
- Tricky in JSON
 - XML has it (XLink), JSON doesn't
 - How do we do it?

Instance Reference (v1)

GET /accounts/x7y8z9

200 OK

```
{  
  "href": "https://api.stormpath.com/v1/accounts/x7y8z9",  
  "givenName": "Tony",  
  "surname": "Stark",  
  .../  
  "directory": ????
```

Instance Reference (v1)

GET /accounts/x7y8z9

200 OK

```
{  
  "href": "https://api.stormpath.com/v1/accounts/x7y8z9",  
  "givenName": "Tony",  
  "surname": "Stark",  
  .../  
  "directory": {  
    "href": "https://api.stormpath.com/v1/directories/g4h5i6"  
  }  
}
```

Collection Reference (v1)

GET /accounts/x7y8z9

200 OK

```
{  
  "href": "https://api.stormpath.com/v1/accounts/x7y8z9",  
  "givenName": "Tony",  
  "surname": "Stark",  
  .../  
  "groups": {  
    "href": "https://api.stormpath.com/v1/accounts/x7y8z9/groups"  
  }  
}
```

Linking v2 (recommended)

Instance HREF (v2)

GET /accounts/x7y8z9

200 OK

```
{  
  "meta": {  
    "href": "https://api.stormpath.com/v1/accounts/x7y8z9",  
  },  
  "givenName": "Tony",  
  "surname": "Stark",  
  ...  
}
```

Instance Reference (v2)

GET /accounts/x7y8z9

200 OK

```
{
  "meta": { ... },
  "givenName": "Tony",
  "surname": "Stark",
  ...,
  "directory": {
    "meta": {
      "href": https://api.stormpath.com/v1/directories/g4h5i6
    }
  }
}
```

Collection Reference (v2)

GET /accounts/x7y8z9

200 OK

```
{
  "meta": { ... },
  "givenName": "Tony",
  "surname": "Stark",
  ...,
  "groups": {
    "meta": {
      "href": "https://api.stormpath.com/v1/accounts/x7y8z9/groups",
      "rel": ["collection"]
    }
  }
}
```

Reference Expansion

(aka Entity Expansion, Link Expansion)

Account and its Directory?

```
GET /accounts/x7y8z9?expand=directory
```

```
200 OK
```

```
{  
  "meta": {...},  
  "givenName": "Tony",  
  "surname": "Stark",  
  .../  
  "directory": {  
    "meta": { ... },  
    "name": "Avengers",  
    "description": "Hollywood's plan for more $",  
    "createdAt": "2012-07-01T14:22:18.029Z",  
    ...  
  }  
}
```

Partial Representations

GET

/accounts/x7y8z9?fields=givenName,surname,
directory(name)

Collections!

Collections

- A first class resource ‘citizen’
- Own href / metadata
- Own properties
- Different from all other collections

GET /accounts/x7y8z9/groups

200 OK

```
{
  "meta": { ... },
  "offset": 0,
  "limit": 25,
  "size": 289,
  "first": { "meta":{"href": ".../accounts/x7y8z9/groups?offset=0"} },
  "previous": null,
  "next": { "meta":{"href": ".../accounts/x7y8z9/groups?offset=25"} },
  "last": { "meta":{"href": "..."} },
  "items": [
    {
      "meta": { "href": "...", ... }
    },
    ...
  ]
}
```

Pagination

Collection Resource supports query params:

- Offset
- Limit

.../applications?offset=50&limit=25

GET /accounts/x7y8z9/groups

200 OK

```
{
  "meta": { ... },
  "offset": 0,
  "limit": 25,
  "first": { "meta":{"href": ".../accounts/x7y8z9/groups?offset=0"} },
  "previous": null,
  "next": { "meta":{"href": ".../accounts/x7y8z9/groups?offset=25"} },
  "last": { "meta":{"href": "..."} },
  "items": [
    {
      "meta": { "href": "...", ...}
    },
    {
      "meta": { "href": "...", ...}
    },
    ...
  ]
}
```

Sorting

```
GET .../accounts?  
    orderBy=surname,givenName%20desc
```


Search

“Find all accounts with a
‘company.com’ email address that
can login to a particular
application”

```
GET /applications/x7y8z9/accounts?email=*company.com
```

```
200 OK
```

```
{
  "meta": { ... },
  "offset": 0,
  "limit": 25,
  "first": { "meta":{
    "href":"/applications/x7y8z9/accounts?email=*company.com&offset=0"}
  },
  "previous": null,
  "next": { "meta":{
    "href": "/applications/x7y8z9/accounts?email=*company.com&offset=25"}
  },
  "last": { "meta":{"href": "..."}},
  "items": [
    { "meta": { "href": "...", ...} },
    { "meta": { "href": "...", ...} },
    ...
  ]
}
```

Search cont'd

- Filter search

.../accounts?q=some+value

- Attribute Search

.../accounts?**surname**=Joe&**email**=*company.com

Search cont'd

- Starts with

`?email=joe*`

- Ends with

`?email=*company.com`

- Contains

`?email=*foo*`

Search cont'd

- Range queries

“all accounts created between September 1st and the 15th”

```
.../accounts?createdAt=[2014-09-01,2014-09-15]
```

Creating And Updating

Remember HATEOS?

“Let me go read the docs to figure out
how to POST”

NO! This is not HATEOS.

How do browsers work?

Forms

Forms: Create / Update

```
{
  "foo": "bar",
  "baz": "boo",
  ...

  "createAccount": {
    "meta": {
      "href": ".../v1/users",
      "type": "form"
      "method": "POST",
    },
    "items": [
      { "type": "text",
        "name": "login",
        "required": "true",
        "label": "Username or Email"
      },
      { "type": "password",
        "name": "password",
        "required": "true",
        "label": "Password"
      }
    ]
  }
}
```

Forms: Search / Query

```
{
  ...

  "findAccounts": {
    "meta": { "href": ".../v1/users", "type": "form", "method": "GET" },
    "items": [
      { "type": "text",
        "name": "username",
        "label": "Username"
        "orderable": true
      },
      { "type": "email",
        "name": "email",
        "label": "Email"
        "orderable": true
      },
      { "type": "text",
        "name": "givenName",
        "label": "First Name"
        "orderable": true
      },
      { "type": "text",
        "name": "surname",
        "label": "Last Name"
        "orderable": true
      },
    ]
  }
}
@lhazlewood
```

Many To Many

Group to Account

- A group can have many accounts
- An account can be in many groups
- Each mapping is a resource:

GroupMembership

GET /groupMemberships/231k3j2j3

200 OK

```
{
  "meta": {"href": ".../groupMemberships/231k3j2j3"},
  "account": {
    "meta": {"href": "..."}
  },
  "group": {
    "meta": {"href": "..."}
  },
  ...
}
```

GET /accounts/x7y8z9

200 OK

```
{
  "meta": {"href": ".../accounts/x7y8z9"},
  "givenName": "Tony",
  "surname": "Stark",
  ...,
  "groups": {
    "meta": {"href": ".../accounts/x7y8z9/groups"}
  },
  "groupMemberships": {
    "meta": {"href": ".../groupMemberships?accountId=x7y8z9"}
  }
}
```

Async or Long-Lived Operations

POST /emails

```
{  
  "from": me@somewhere.com,  
  "subject": "Hi!"  
  "body": "..."  
}
```

204 Accepted

Location: /emails/23Sd932sS1

```
{  
  "status": "queued",  
  ...  
}
```

GET /emails/23Sd932sS1

Expires: 2014-09-29T18:00:00.000Z

```
{  
  "status": "sent",  
  ...  
}
```

Batch Operations

- Each batch is represented as a resource
- Batches are likely to be a collection
- Batches are likely to have a status
- Batch deletes easier than create/update

Batch Delete

“Delete all company.com accounts”

```
DELETE /accounts?  
      email=*@company.com
```

Batch Create / Update

Already have a Collection concept. Use it.

Batch Create or Update

POST /accounts

```
{  
  "items": [  
    { ... account 1 ... },  
    { ... account 2 ... },  
    ...  
  ]  
}
```

Batch Operations: The ‘Catch’

HTTP Caching is bypassed entirely 😞

204 Accepted

Location: /batches/a1b2c3

```
{
  "status": "processing", //overall status
  "size": "n",
  "limit": 25,
  ...,
  "items": {
    { response 1 (w/ individual status) ...},
    { response 2 (w/ individual status) ...},
    ...
  }
}
```

Errors

- As descriptive as possible
- As much information as possible
- Developers are your customers

POST /directories

409 Conflict

```
{
  "status": 409,
  "code": 40924,
  "property": "name",
  "message": "A Directory named 'Avengers' already exists.",
  "developerMessage": "A directory named 'Avengers' already exists. If you have a stale local cache, please expire it now.",
  "moreInfo": "https://www.stormpath.com/docs/api/errors/40924"
}
```


Security

Avoid sessions when possible

Authenticate every request if necessary

Stateless

Authorize based on resource content, NOT URL!

Use Existing Protocol:

Oauth 1.0a, Oauth2, Basic over SSL only

Custom Authentication Scheme:

Only if you provide client code / SDK

Only if you really, *really* know what you're doing

Use API Keys and/or JWTs instead of Username/Passwords

401 vs 403

- 401 “Unauthorized” *really* means Unauthenticated

“You need valid credentials for me to respond to this request”

- 403 “Forbidden” *really* means Unauthorized

“Sorry, you’re not allowed!”

HTTP Authentication Schemes

- Server response to issue challenge:

WWW-Authenticate: *<scheme name>*
realm="Application Name"

- Client request to submit credentials:

Authorization: *<scheme name> <data>*

API Keys

- Entropy
- Password Reset
- Independence
- Scope
- Speed
- Limited Exposure
- Traceability

IDs

- IDs should be opaque
- Should be globally unique
- Avoid sequential numbers (contention, fusing)
- Good candidates: UUIDs, 'Url64'

HTTP Method Overrides

POST /accounts/x7y8z9?_method=DELETE

Caching & Concurrency Control

Server (initial response):

ETag: "686897696a7c876b7e"

Client (later request):

If-None-Match: "686897696a7c876b7e"

Server (later response):

304 Not Modified

Maintenance

Use HTTP Redirects

Create abstraction layer / endpoints when migrating

Use well defined custom Media Types

IETF RFC?

ionwg.github.io/draft-ion.html



- Free for developers
- Eliminate months of development
- Automatic security best practices
- Single Sign On for your apps
- API Authentication & Key Management
- Token Authentication for SPAs / Mobile
- Authorization

Libraries and integrations:
<https://docs.stormpath.com>